
EGA: Manifold-Preserving Adapter Training for Out-of-Distribution Retrieval in Vector Search

Anonymous Author(s)

Affiliation

Address

email

Abstract

Adapting pre-trained vision-language models for vector retrieval typically relies on global contrastive objectives, which optimize semantic alignment across all training categories simultaneously. We identify a fundamental limitation of this approach: when applied through lightweight adapters, global objectives produce dense gradient fields that systematically distort the latent manifold of unseen categories, causing out-of-distribution (OOD) retrieval performance to collapse below the frozen baseline. To address this, we propose **Euclidean Geodesic Alignment (EGA)**, a residual adapter trained with a small-margin triplet loss on the class-aware neighborhood graph of pre-trained embeddings. By operating locally rather than globally, EGA converges to a sparse-gradient regime in which 96.5% of training triplets produce zero gradient at steady state, leaving the broader embedding space—including regions occupied by unseen categories—structurally intact. Extensive evaluations on CIFAR-10, CIFAR-100, FGVC-Aircraft, and Food-101 demonstrate that EGA consistently outperforms global contrastive methods including ICon and SRL across both in-distribution and OOD settings. On strictly unseen fine-grained categories, competing adapters degrade below the frozen CLIP baseline in label precision, whereas EGA improves it by up to 18.5%.

1 Introduction

Vector similarity search is the core mechanism for retrieving high dimensional representations in modern applications. While foundational vision language models such as CLIP [21] offer rich zero shot semantic capabilities, their latent manifolds often exhibit a systemic disconnect between semantic alignment and geometric proximity. This geometric illusion means that true semantic neighbors frequently overlap with noise in the Euclidean space, rendering standard approximate nearest neighbor search inefficient [24].

To address these metric inconsistencies, recent approaches apply global contrastive objectives or structural regularizers [1, 4] to recalibrate the representation space. However, these methods rely on batch wide alignment and uniformity pressures that generate dense gradient fields during training. When applied through lightweight adapters for parameter efficient tuning, this global optimization aggressively reconfigures the entire latent manifold. Consequently, the delicate geometric structures of unseen fine grained categories are systematically distorted, causing out of distribution retrieval performance to collapse below that of the frozen baseline.

To overcome this fundamental limitation, we propose Euclidean Geodesic Alignment (EGA), a parameter efficient manifold smoothing approach. EGA introduces a residual adapter equipped with zero initialization and L2 normalization constraints to refine the pre trained semantic space. Instead of forcing global uniformity, EGA optimizes local geometric anchors via a small margin triplet loss.

36 This localized objective compresses class specific clusters and resolves metric misalignment without
37 subjecting the broader pre trained semantic relations to destructive global gradients.

38 Extensive empirical evaluations across CIFAR-10, CIFAR-100, FGVC-Aircraft, and Food-101
39 validate the superiority of our method. EGA converges to a sparse gradient regime where 96.5 percent
40 of training triplets produce zero gradient at steady state, effectively preserving the structural integrity
41 of unseen categories. As a result, EGA consistently outperforms state of the art global contrastive
42 methods in both in distribution and out of distribution settings. For instance, on strictly unseen fine
43 grained domains, competing methods degrade below the frozen baseline, whereas EGA significantly
44 elevates label precision by up to 18.5 percent and achieves parity with multi probe search baselines
45 using only single cell lookups.

46 1.1 Contributions

47 This paper makes the following contributions:

- 48 • We identify the systemic disconnect between semantic alignment and geometric proximity
49 in foundational latent manifolds, revealing that distance overlap between true neighbors and
50 noise renders Euclidean indexing unreliable.
- 51 • We propose Euclidean Geodesic Alignment (EGA), a parameter efficient adapter leveraging
52 residual structures and local geometric anchors to smooth manifold irregularities without
53 destroying pre trained semantic relations.
- 54 • We demonstrate through rigorous empirical evaluation that EGA significantly accelerates
55 approximate nearest neighbor search and exhibits superior out of distribution generalization
56 on complex fine grained domains, outperforming contemporary global contrastive objectives.

57 2 Related Work

58 **Vision Language Representations and Semantic Manifolds** Vision language foundational models,
59 such as CLIP [21], have established a dominant paradigm for constructing joint semantic spaces
60 through large scale contrastive learning. Despite their remarkable zero shot capabilities, recent studies
61 reveal that these continuous representations suffer from inherent geometric failures [24]. Specifically,
62 the modality gap and cross modal misalignment lead to feature spaces that are overly sparse and
63 structurally fractured [22, 5]. Furthermore, attempts to align complex compositional concepts
64 often struggle in standard Euclidean spaces, highlighting the systemic disconnect between semantic
65 meaning and spatial proximity [20]. This geometric illusion is exacerbated by the representation
66 degradation problem prevalent in contrastive objectives. Extensive research indicates that latent
67 embeddings tend to occupy narrow conic regions, a phenomenon known as anisotropy, which renders
68 standard similarity metrics ineffective [36]. When models undergo continual learning or adapt
69 to imbalanced fine grained categories, the feature space becomes increasingly distorted [18, 31].
70 Consequently, the dense gradient updates required by conventional multi view alignment often
71 destruct specific latent structures rather than preserving them [30]. To mitigate these issues, recent
72 adaptation strategies impose strict geometric constraints, such as orthogonality regularization, to
73 balance new knowledge acquisition with the stability of the original representation [23]. While these
74 methods attempt to preserve the global topological structure, they still rely on global optimization
75 pressures that inadvertently alter the continuous manifold. Unlike these approaches, EGA employs a
76 lightweight residual adapter with local geometric anchors to smooth specific class clusters, thereby
77 correcting the Euclidean distance metric without subjecting the broader pre trained semantic relations
78 to destructive global gradients.

79 **Global Contrastive Objectives and Structural Regularization** Recent advancements in repre-
80 sentation learning often rely on global contrastive objectives to optimize the latent space [27, 17]. A
81 common paradigm involves jointly enforcing alignment for positive pairs and uniformity for the over-
82 all feature distribution, a strategy applied across various domains including graph encoders [28] and
83 multimodal recommendation systems [37]. While pushing embeddings toward a uniform distribution
84 enhances robustness to background noise [29], it inherently assumes that all latent regions require the
85 same degree of global scattering. To further refine complex distributions, structural regularization
86 techniques introduce auxiliary constraints. For instance, methods like I Con formulate a unifying

framework to systematically align multi view representations [1], while geometric constraints are leveraged to improve imbalanced regression tasks [4]. Other approaches attempt to decouple the learning objectives [16] or apply information theoretic regularizers [35] to control the density of the learned representations. Furthermore, region based distillation techniques have been proposed to enhance fine grained understanding [13]. However, these global contrastive objectives and structural regularizers typically generate dense gradient fields throughout the training process. By continuously applying optimization pressure across the entire batch, they inadvertently disrupt the delicate pre established geometric structures of unseen fine grained categories. Unlike these global approaches that aggressively reconfigure the entire latent manifold through dense updates, EGA utilizes local geometric anchors to perform sparse and targeted metric smoothing, thereby preserving the structural integrity of out of distribution semantics.

Parameter Efficient Manifold Adaptation Parameter efficient adaptation emerged as a standard to prevent catastrophic forgetting and reduce computational overhead when transferring large foundation models [8, 9]. This paradigm has extensively shifted towards vision language models, dominating recent literature due to the prohibitive cost of full scale fine tuning [19]. Methods such as Tip Adapter demonstrate that lightweight tuning can effectively specialize representations for downstream tasks without altering the frozen backbone [34]. Furthermore, extensive benchmarking confirms that parameter efficient techniques can rival full fine tuning while preserving the generalizable representations of the foundation model [25]. To further protect the pre trained manifold during initial training steps, zero initialization strategies have proven highly effective. For instance, conditional control mechanisms in image generation initialize adapter weights to zero, ensuring the model behaves exactly like the frozen base model before any optimization occurs [33]. While these existing methods successfully maintain global knowledge and prevent catastrophic forgetting, they are not designed to resolve local metric inconsistencies for vector search. Unlike these general adaptation strategies, EGA integrates a zero initialized residual adapter with a localized geometric objective, explicitly smoothing class specific clusters to correct Euclidean distances without compromising the parameter efficiency or stability of the pre trained manifold.

Vector Similarity Search and Metric Consistency Vector similarity search is the foundational engine for retrieving high dimensional representations at scale. The systems community has developed highly optimized indexing algorithms to handle billion scale datasets. Foundational works such as FAISS leverage hardware acceleration for fast similarity search [14], while graph based indices and hybrid structures like DiskANN and SPANN achieve exceptional speed by efficiently managing memory and disk hierarchies [26, 2]. Recent advancements continue to refine these topologies, exploring crossing sparse proximity graphs [32] and analyzing the theoretical constructions of navigable graphs for high dimensional nearest neighbor search [3]. Furthermore, system level studies optimize retrieval latency and stability by aligning best first search algorithms with solid state drives [6] and enabling direct insertions in large scale indices [7]. To reduce the computational cost of exact distance calculations during the retrieval phase, approximation techniques such as low rank matrix factorization have also been introduced to accelerate inner product evaluations [12]. However, extensive theoretical analyses reveal that popular approximate nearest neighbor implementations still face severe worst case performance limitations [11]. While these sophisticated indexing and approximation systems continuously push the boundaries of retrieval throughput, they operate under the strict assumption that the underlying Euclidean space perfectly reflects semantic similarity. If the foundational manifold suffers from geometric distortion, these highly optimized engines will simply retrieve semantically irrelevant neighbors at an accelerated rate. Unlike hardware and system level optimizations that treat the embedding space as an immutable input, EGA proactively recalibrates the metric manifold prior to index construction, thereby enabling standard search backends to achieve superior semantic recall without requiring complex algorithmic reengineering.

3 Methodology

3.1 Problem Formulation

Let $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$ denote a frozen vision-language encoder (e.g., CLIP ViT-B/32) that maps an image $x \in \mathcal{X}$ to a unit-normalized embedding $\mathbf{z} = \phi(x) \in \mathbb{S}^{d-1}$. We are given a set of labeled training images $\mathcal{D}_{\text{train}} = \{(x_i, y_i)\}_{i=1}^N$, where $y_i \in \mathcal{Y}_{\text{seen}}$ belongs to a set of *seen* classes. At inference time,

the retrieval system must operate over a database that also contains images from a disjoint set of *unseen* classes $\mathcal{Y}_{\text{unseen}}$, with $\mathcal{Y}_{\text{seen}} \cap \mathcal{Y}_{\text{unseen}} = \emptyset$.

Our goal is to learn a lightweight adapter $f_\theta : \mathbb{S}^{d-1} \rightarrow \mathbb{S}^{d-1}$ such that the transformed embeddings $\hat{\mathbf{z}} = f_\theta(\mathbf{z})$ support more accurate approximate nearest neighbor (ANN) retrieval on *both* seen and unseen classes, without access to any unseen-class labels during training.

Evaluation metrics. We evaluate retrieval quality along two complementary dimensions.

Label Precision@K measures the semantic quality of retrieved results—whether the K nearest neighbors in the transformed space share the same class label as the query:

$$\text{LP@K} = \frac{1}{|\mathcal{Q}|} \sum_{q \in \mathcal{Q}} \frac{1}{K} \sum_{k=1}^K \mathbf{1}[y_{\hat{n}_k(q)} = y_q], \quad (1)$$

where $\mathcal{Q}$ is the query set and $\hat{n}_k(q)$ denotes the k -th nearest neighbor of q in the transformed embedding space.

ANNS Recall@K measures the geometric fidelity of the ANN index—the fraction of true K nearest neighbors (under exact search) that are recovered by the approximate index at a given search budget n_{probe} :

$$\text{AR@K} = \frac{1}{|\mathcal{Q}|} \sum_{q \in \mathcal{Q}} \frac{|\mathcal{N}_K^{\text{exact}}(q) \cap \mathcal{N}_K^{\text{approx}}(q)|}{K}, \quad (2)$$

where $\mathcal{N}_K^{\text{exact}}(q)$ and $\mathcal{N}_K^{\text{approx}}(q)$ denote the exact and approximate top- K neighbor sets, respectively.

The two metrics capture distinct failure modes: an adapter may improve LP@K by tightening within-class clusters (semantic improvement) while simultaneously improving AR@K by making the geometry more amenable to IVF-based indexing (geometric improvement). As we show in Section 4, global contrastive objectives achieve high AR@K on seen classes but collapse on *both* metrics for unseen classes, whereas EGA improves both simultaneously under the OOD setting.

3.2 Manifold-Preserving Adapter Design

EGA introduces a lightweight adapter f_θ composed of stacked residual blocks with a final refinement layer. The architecture is deliberately constrained by three design principles, each of which serves a distinct role in preserving the pretrained manifold.

Residual connection with zero initialization. Each EGA block applies a residual update of the form

$$\mathbf{z}' = \mathbf{z} + g_\theta(\mathbf{z}), \quad (3)$$

where g_θ is a learned transformation initialized so that $g_\theta(\mathbf{z}) \approx \mathbf{0}$ at the start of training. Concretely, the final linear layer of g_θ is initialized to zero, making f_θ an identity mapping at initialization. This guarantees that training begins from the pretrained geometry rather than a random perturbation, and that the adapter can only *refine* rather than replace the CLIP embedding. As confirmed by our ablation (Table 1), removing the residual connection causes accuracy to collapse from 0.70 to 0.01, because the lightweight MLP cannot reconstruct complex semantic relationships from seen-class data alone.

Feature gating. Inside each residual block, the transformed signal passes through a feature-gating module:

$$\text{Gate}(\mathbf{h}) = \mathbf{h} \odot \sigma(W_2 \text{GELU}(W_1 \mathbf{h})), \quad (4)$$

where $W_1 \in \mathbb{R}^{(d/4) \times d}$, $W_2 \in \mathbb{R}^{d \times (d/4)}$, σ is the sigmoid function, and $\odot$ denotes elementwise multiplication. This structure, analogous to Squeeze-and-Excitation networks [10], learns a per-dimension importance weight that amplifies retrieval-relevant dimensions and suppresses noisy ones. The bottleneck ratio of 1/4 is chosen to prevent the gating network from overfitting on seen-class data, as validated in Table 2.

178 **ℓ_2 normalization.** The final output of f_θ is projected back onto the unit hypersphere $\mathbb{S}^{d-1}$ via ℓ_2
 179 normalization:

$$\hat{\mathbf{z}} = \frac{f_\theta(\mathbf{z})}{\|f_\theta(\mathbf{z})\|_2}. \quad (5)$$

180 This constraint ensures that the transformed embeddings remain in the same metric space as the
 181 original CLIP embeddings, so that Euclidean distance and cosine similarity remain interchangeable
 182 on the hypersphere. Removing this constraint degrades retrieval performance by 4.2 percentage
 183 points (Table 1), confirming that hypersphere consistency is necessary for stable ANN indexing.

184 **Full forward pass.** Let Block_i denote the i -th EGA block consisting of a two-layer MLP with Lay-
 185 erNorm followed by feature gating, and let Refine denote the final linear projection with LayerNorm.
 186 The complete forward pass for an input embedding $\mathbf{z} \in \mathbb{S}^{d-1}$ is:

$$\hat{\mathbf{z}} = \ell_2(\text{Refine}(\text{Block}_2(\text{Block}_1(\mathbf{z})))) . \quad (6)$$

187 The total number of trainable parameters is approximately $4.2M$ for $d = 512$ and hidden dimension
 188 2048, which is negligible compared to the frozen CLIP backbone ($\sim 150M$ parameters).

189 3.3 Local Triplet Training and Gradient Sparsity

190 **Training objective.** Given a mini-batch sampled from $\mathcal{D}_{\text{train}}$, we construct triplets (a, p, n) where a
 191 is an anchor image, p is a positive sampled uniformly from the same class as a , and n is a negative
 192 sampled uniformly from a different class. The adapter is trained with the standard triplet margin loss:

$$\mathcal{L} = \frac{1}{|\mathcal{B}|} \sum_{(a,p,n) \in \mathcal{B}} \max(0, d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m), \quad (7)$$

193 where $d(\cdot, \cdot)$ denotes Euclidean distance, $m > 0$ is the margin hyperparameter, and $\hat{\mathbf{z}} = f_\theta(\mathbf{z})$ is the
 194 transformed embedding. We set $m = 0.2$ throughout all experiments based on the sensitivity analysis
 195 in Table 2.

196 **Gradient sparsity and manifold preservation.** A triplet (a, p, n) contributes a non-zero gradient
 197 to θ if and only if the margin constraint is *violated*:

$$d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m > 0. \quad (8)$$

198 We refer to such triplets as *active*. As the adapter converges on seen classes, increasingly many
 199 triplets satisfy the margin constraint and become inactive, contributing zero gradient. Figure 8 (left)
 200 shows that the active triplet ratio decays from 17.4% at epoch 10 to 3.5% at convergence. At steady
 201 state, **96.5% of triplets produce zero gradient**, meaning the adapter ceases to modify the vast
 202 majority of local neighborhood relationships in the embedding space.

203 This sparsity has a direct consequence for out-of-distribution generalization. Because unseen classes
 204 are absent from $\mathcal{D}_{\text{train}}$, no triplet involving an unseen-class embedding is ever constructed. Provided
 205 that the adapter’s updates remain local—i.e., that gradients from seen-class triplets do not propagate
 206 into unseen-class regions of the embedding space—the CLIP manifold structure for unseen classes
 207 is preserved intact. The gradient sparsity mechanism enforces exactly this locality: once seen-class
 208 neighborhoods satisfy the margin, the corresponding gradients vanish, and only the small fraction of
 209 actively-violated local relationships continue to be updated.

210 **Contrast with global contrastive objectives.** Global objectives such as ICon [1] and SRL [4]
 211 compute losses over the *full* pairwise similarity matrix within each mini-batch. ICon minimizes a
 212 KL divergence between the empirical similarity distribution and the class-membership distribution,
 213 while SRL jointly optimizes batch-wide uniformity and within-class homogeneity. Both objectives
 214 generate dense gradient fields in every training step: every sample pair—regardless of whether
 215 its local geometry already satisfies the retrieval objective—contributes to the loss and receives a
 216 non-zero gradient. This global optimization pressure from seen-class training propagates throughout
 217 the embedding space, distorting the local manifold structure that CLIP had established for unseen
 218 categories.




Figure 1: Label precision across varying K and nprobe settings. EGA consistently outperforms the original latent manifold, demonstrating improved semantic alignment.

Proposition 1 (Gradient Sparsity at Convergence). *Let f_θ be trained with the triplet loss in Eq. (7) with margin $m > 0$ on a dataset $\mathcal{D}_{\text{train}}$ drawn from seen classes $\mathcal{Y}_{\text{seen}}$. Define the active triplet ratio at training step t as*

$$\rho(t) = \Pr_{(a,p,n) \sim \mathcal{B}} \left[d(\hat{\mathbf{z}}_a^{(t)}, \hat{\mathbf{z}}_p^{(t)}) - d(\hat{\mathbf{z}}_a^{(t)}, \hat{\mathbf{z}}_n^{(t)}) + m > 0 \right]. \quad (9)$$

As f_θ converges, $\rho(t) \rightarrow \rho^*$, where

$$\rho^* = \Pr_{(a,p,n)} \left[d(\hat{\mathbf{z}}_a^*, \hat{\mathbf{z}}_p^*) - d(\hat{\mathbf{z}}_a^*, \hat{\mathbf{z}}_n^*) + m > 0 \right] \quad (10)$$

is determined solely by the margin m and the within-class distance distribution of the converged embeddings. Furthermore, $\nabla_\theta \mathcal{L} = \mathbf{0}$ for all triplets (a, p, n) with $d(\hat{\mathbf{z}}_a^*, \hat{\mathbf{z}}_p^*) - d(\hat{\mathbf{z}}_a^*, \hat{\mathbf{z}}_n^*) + m \leq 0$, so the adapter gradient is supported only on the fraction ρ^* of the training distribution. For embeddings whose local geometry already satisfies the margin constraint, the adapter exerts zero update pressure.

Proof Sketch. The loss in Eq. (7) is a sum of hinge terms, each of which has gradient zero whenever $d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m \leq 0$. At convergence, f_θ has minimized the loss over $\mathcal{D}_{\text{train}}$, so the only remaining active triplets are those whose geometry the adapter has not yet been able to separate within the margin. The fraction of such triplets is bounded below by the overlap of the within-class and between-class distance distributions at the converged embedding, which depends on m and the geometry of $\mathcal{Y}_{\text{seen}}$ but not on $\mathcal{Y}_{\text{unseen}}$. Since no triplet involving unseen-class embeddings is ever constructed, the adapter gradient is identically zero over the entire unseen-class region of the embedding space throughout training. See Appendix A for the complete proof. $\square$

The empirical validation of Proposition 1 is provided in Section 4.7, where we track the active triplet ratio throughout training and connect it directly to the OOD retrieval outcomes shown in Figure 8.

4 Evaluation

4.1 Experimental Setup

Experiment Platform All experiments were conducted on the Chameleon Cloud [15] platform using a compute node equipped with 256 AMD EPYC-7763 CPU cores, 512 GiB of RAM, and an NVIDIA A100 GPU of 80 GB HBM2e. The operating system is Ubuntu 24.04 LTS.

4.2 Ground truth of label precision

The metric of label precision is employed to assess the semantic quality of the retrieved results within the approximate nearest neighbor framework. This evaluation focuses on whether the neighbors found in the feature space are categorized under the same physical class as the query. By examining different neighborhood sizes from K=1 to K=10, we can determine the local and global semantic density of the manifold. Our analysis contrasts the original latent manifold with the EGA flattened model to highlight the impact of manifold smoothing on search accuracy.




Figure 2: L2 distance distributions for Top K neighbors versus random background points. The original space exhibits significant overlap, whereas EGA enforces a clear geometric margin.




Figure 3: Recall@K versus nprobe tradeoff. EGA consistently shifts the efficiency frontier upward, enabling high recall at minimal search probes.

Figure 1 demonstrates that EGA flattened achieves a substantial improvement in label precision across all configurations. In the $K=1$ scenario, our method maintains a precision score exceeding 0.7, while the original latent manifold fluctuates between 0.5 and 0.6. This performance gap remains persistent as the search budget nprobe increases from 1 to 10, indicating that the benefits of our approach are not dependent on extensive search efforts. The superior semantic alignment shown in these figures suggests that EGA creates a more organized and discriminative feature space, allowing the indexer to retrieve correct class labels with higher reliability than the baseline.

4.3 Distances in Latent and EGA Spaces

To validate our stance on the systemic disconnect between semantic alignment and geometric proximity, we present empirical evidence across two dimensions: physical distance distributions and retrieval efficiency metrics. We evaluate Euclidean Geodesic Alignment (EGA) against raw foundational model embeddings using CIFAR10 and CIFAR100 datasets.

The results in Figure 2 expose the root cause of retrieval failure. In the original CLIP latent manifold, the distance distributions of true semantic neighbors and random background points exhibit significant overlap. This position paper argues that this overlap constitutes a geometric illusion where noise is indistinguishable from signal, rendering Euclidean indexing inherently unreliable.

4.4 Recalls in Latent and EGA Spaces

The impact of this flattening is quantified in Figure 3. The raw CLIP space achieves a meager 0.51 Recall@1 at nprobe=1 on CIFAR100. To compensate for this disorganized geometry, systems are forced to increase search volume to nprobe=10 to reach a recall of 0.85. This represents a 10x waste of computational resources.

In contrast, EGA elevates the nprobe=1 Recall@1 to 0.68. Notably, EGA at nprobe=1 outperforms the original space at nprobe=5 across all benchmarks. This confirms that fixing the geometry at the




Figure 4: tSNE projection of original CLIP latent manifold (left) versus EGA flattened manifold (right). EGA reveals clearer category clusters and reduced semantic overlap.

source is the only viable path to scalable retrieval. We achieve accuracy parity with brute force multi probe baselines while maintaining the extreme efficiency of single cell lookups. These results verify that our position is a necessary requirement for the next generation of scalable retrieval systems.

4.5 Low dimensional projection visualization

To better understand the structural changes in the feature space, we utilize t SNE to project high dimensional embeddings into a two dimensional plane. Figure 4 compares the spatial distribution of the latent manifold generated by original CLIP against our proposed method. As observed in the left panel, the baseline manifold is characterized by highly dispersed data points and substantial semantic overlap between different categories. For instance, classes such as Class 91 and Class 53 appear entangled in certain regions, which inherently limits the effectiveness of search algorithms.

The right panel of the figure displays the EGA flattened manifold, revealing a significant enhancement in geometric order. Data points corresponding to specific categories, including Class 35 and Class 72, are tightly clustered and separated by well defined boundaries. This visual evidence confirms that our smoothing technique effectively eliminates manifold irregularities and reduces local noise. Such a structured feature distribution is a direct consequence of the manifold flattening process, which facilitates faster convergence in indexing and contributes to the superior recall rates discussed in previous sections.

4.6 Out-of-Distribution Generalization

The generalization capabilities of EGA are evaluated on the CIFAR-10 dataset, providing a rigorous out-of-distribution (OOD) test since all models were trained exclusively on CIFAR-100. Figure 5 illustrates the recall performance across varying nprobe values and K settings. EGA consistently maintains a dominant lead over all baselines, including recent state-of-the-art methods like ICon [1] and SRL [4].

At a minimum search budget of nprobe=1, EGA achieves a Recall@1 of 0.858, significantly outperforming SRL (0.831) and ICon (0.788). As shown in the inset zooms, EGA continues to hold a clear advantage even at nprobe=10, where other models begin to saturate. A key observation is the efficiency gain provided by our method: EGA achieves a recall level at nprobe=5 that matches or exceeds the performance of supervised baselines at nprobe=10. This indicates that the manifold smoothing provided by EGA effectively doubles the search efficiency in unseen data environments by creating a more index-friendly geometric structure.

To further validate the robustness of EGA under fine-grained domain shifts, we extend our evaluation to the FGVC Aircraft dataset. Figure 6 visually corroborates that EGA remains the only method capable of scaling zero-shot retrieval performance on fine-grained domains without destroying the underlying manifold. In this setting, models are trained on 80 known aircraft categories and tested on




Figure 5: Recall@K versus nprobe tradeoff on CIFAR-10. EGA maintains a consistent lead over baselines, demonstrating strong OOD generalization.




Figure 6: Recall versus nprobe tradeoff on 20 unseen FGVC Aircraft classes.

20 strictly unseen categories. For a fair comparison, all baseline methods are restricted to training identical lightweight MLP adapters on top of frozen CLIP features, thereby eliminating the parameter advantage of end-to-end backbone fine-tuning. The empirical results reveal a stark contrast in spatial preservation capabilities. Global contrastive and structural regularization methods, such as ICon and SRL, suffer severe performance degradation when confined to a lightweight adapter bottleneck. Specifically, at $K = 1$ and $nprobe=1$, the recall for ICon and SRL collapses to 0.422 and 0.440 respectively, falling significantly below the original frozen CLIP baseline of 0.773. This phenomenon indicates that global contrastive objectives aggressively distort the latent manifold of unseen fine-grained categories when the representation capacity is constrained. Conversely, EGA exhibits exceptional stability and generalization, elevating the recall to 0.898 under the exact same adapter constraints. By leveraging local geometric anchors rather than forcing global distributions, EGA successfully compresses class-specific clusters while maintaining the structural integrity of the broader semantic space.

To comprehensively evaluate the versatility of EGA across diverse fine-grained domains, we extend our evaluation to the Food-101 dataset. Figure 7 visually corroborates that EGA remains the only method capable of scaling zero-shot retrieval performance on fine-grained domains without destroying the underlying manifold. In this setting, models are trained on 80 known food categories and tested on 21 strictly unseen categories. For a fair comparison, all baseline methods are restricted to training identical lightweight MLP adapters on top of frozen CLIP features, thereby eliminating the parameter advantage of end-to-end backbone fine-tuning. The empirical results reveal a stark contrast in spatial preservation capabilities, particularly given the strong pre-trained representations of common food items. Global contrastive and structural regularization methods, such as ICon and SRL, suffer performance degradation when confined to a lightweight adapter bottleneck. Specifically, at $K = 1$ and $nprobe=1$, the recall for ICon and SRL drops to 0.838 and 0.840 respectively, falling below the original frozen CLIP baseline of 0.842. This phenomenon indicates that global contrastive objectives aggressively distort the already well-structured latent manifold of unseen categories when the representation capacity is constrained. Conversely, EGA exhibits exceptional stability and generalization, elevating the recall to 0.897 under the exact same adapter constraints. By leveraging local geometric anchors rather than forcing global distributions, EGA successfully compresses class-specific clusters while maintaining the structural integrity of the broader semantic space, significantly enhancing retrieval efficiency even on a highly optimized baseline manifold.




Figure 7: Recall versus nprobe tradeoff on 21 unseen Food-101 classes.




Figure 8: Gradient sparsity as the mechanism behind EGA’s OOD generalization.

337 4.7 Mechanistic Analysis: Gradient Sparsity as OOD Protection

338 The superior OOD generalization of EGA raises a fundamental question: why does a triplet-trained
 339 adapter preserve unseen class structure while global contrastive objectives destroy it? We identify
 340 **gradient sparsity** as the key mechanism.

341 A triplet loss with margin m produces a non-zero gradient for a sample triplet (a, p, n) only when the
 342 margin constraint is violated, i.e., when

$$d(a, p) - d(a, n) + m > 0. \quad (11)$$

343 As the adapter converges on seen classes, the fraction of such *active* triplets decays rapidly. Figure 8
 344 (left) tracks this ratio throughout training: starting at 17.4% at epoch 10, it falls to 3.5% by conver-
 345 gence. At steady state, **96.5% of triplets contribute zero gradient**, meaning the adapter updates
 346 only a sparse subset of local neighborhood relationships and leaves the remainder of the embedding
 347 space—including regions occupied by unseen classes—structurally intact.

348 Global contrastive objectives such as ICon [1] and SRL [4] exhibit the opposite behavior. ICon
 349 minimizes a KL divergence over the full pairwise similarity matrix, and SRL jointly optimizes
 350 batch-wide uniformity and homogeneity. Both produce *dense* gradient fields in every training step:
 351 every sample pair contributes to the loss regardless of whether its local geometry already satisfies the
 352 retrieval objective. This global optimization pressure from seen-class training propagates into unseen
 353 regions of the embedding space, distorting the local manifold structure that CLIP had established for
 354 those categories.

355 Figure 8 (right) quantifies the consequence on FGVC Aircraft, where all models are trained on 80 seen
 356 classes and evaluated strictly on 20 unseen classes. ICon and SRL collapse to label precision of 0.423
 357 and 0.441, respectively—**below the frozen CLIP baseline of 0.512**—confirming that dense gradient
 358 updates actively harm zero-shot retrieval. EGA, by contrast, improves to 0.548 while modifying only
 359 3.5% of the embedding geometry at convergence.

360 This analysis points to a key distinction in adapter design for foundation models: **objectives that pro-**
 361 **duce dense gradient fields throughout training systematically degrade the pretrained manifold**

Table 1: Ablation study of EGA components on CIFAR-100.

Variant	Accuracy
Full EGA (Ours)	0.7015
w/o Residual connection	0.0065
w/o Zero-initialization	0.6840
w/o L2-normalization	0.6590

for unseen categories, while objectives that converge to sparse local updates preserve it. The contrast is not merely one of degree—ICoN and SRL fall below the frozen CLIP baseline entirely, suggesting that dense global optimization actively destructs zero-shot retrieval structure rather than simply failing to improve it.

4.8 Ablation Study

We perform a comprehensive ablation study on the CIFAR 100 dataset to verify the effectiveness of our design components. The results are summarized in Table 1. Each architectural choice contributes to the structural integrity of the calibrated manifold.

The residual connection is the most essential part of the design. When the residual connection is removed, the accuracy collapses from 0.7015 to 0.0065. This failure proves that the EGA adapter should refine the existing semantic space rather than replace it. Without this shortcut, the lightweight MLP cannot reconstruct the complex semantic relationships from limited training data.

L2 normalization is also necessary for stable search performance. Its removal leads to a performance loss of approximately 4.2 percent. This confirms that the hypersphere constraint is vital for consistent distance metrics in vector similarity search. Finally, the zero initialization strategy provides a 1.7 percent gain. By starting the adapter as an identity mapping, the model preserves pre trained knowledge at the beginning of the process, which provides a superior starting point for manifold smoothing.

4.9 Hyperparameter Sensitivity Analysis

The internal stability of EGA is evaluated by varying the triplet margin m and the hidden dimension d of the adapter. As presented in Table 2, our method demonstrates a consistent performance plateau across a broad range of settings.

The results for $K = 5$ indicate that EGA maintains high retrieval quality even when the margin m fluctuates. While the $K = 1$ metric shows sensitivity to larger margin values, the overall stability of the manifold is preserved. This trend suggests that a moderate margin setting, specifically between 0.1 and 0.3, is optimal for calibrating fine grained latent spaces. Furthermore, the performance remains nearly identical as the hidden dimension d scales from 512 to 4096. This observation confirms that the geometric correction provided by EGA is structural in nature and does not rely on massive parameter counts.

Table 2: Sensitivity of retrieval performance to different hyperparameter configurations on CIFAR 100.

Triplet Margin m	0.1	0.2	0.3	0.5
Recall@1	0.8420	0.8180	0.8160	0.7400
Recall@5	0.9782	0.9654	0.9610	0.9125
Hidden Dimension d	512	1024	2048	4096
Recall@1	0.8300	0.8340	0.8200	0.8240
Recall@5	0.9695	0.9712	0.9680	0.9688

5 Conclusion

In this paper, we introduced Euclidean Geodesic Alignment (EGA) to bridge the gap between semantic alignment and geometric proximity in foundational representation spaces. We identified

that global contrastive objectives often produce dense gradient fields that distort the latent manifold of unseen categories. To resolve this, EGA utilizes a lightweight residual adapter initialized as an identity mapping and trained with a small margin triplet loss. This localized approach enables the model to refine specific class clusters while preserving the broader semantic integrity established by the pre trained backbone.

Our experimental results across multiple benchmarks demonstrate that EGA significantly improves label precision and retrieval recall, particularly in out of distribution settings. The analysis of gradient sparsity reveals that EGA converges to a regime where only a small fraction of training samples exert update pressure, effectively protecting the manifolds of unseen categories from destructive optimization. This finding suggests that sparse local updates are vital for maintaining the zero shot capabilities of foundational models during adaptation. While our current evaluation focuses on visual representations, future work could explore the application of EGA to more diverse multimodal manifolds. By achieving parity with multi probe search through single cell lookups, EGA proves that correcting the metric geometry at the source is a highly efficient strategy for scalable vector retrieval. Unlike conventional structural regularization methods that rely on global optimization pressures, EGA provides a sparse and targeted solution for robust manifold alignment in vector search systems.

References

- [1] Shaden Naif Alshammari, John R Hershey, Axel Feldmann, William T Freeman, and Mark Hamilton. I-con: A unifying framework for representation learning. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- [2] Qi Chen, Bing Zhao, Haidong Wang, Mingqin Li, Chuanjie Liu, Zengzhong Li, Mao Yang, and Jingdong Wang. Spann: highly-efficient billion-scale approximate nearest neighbor search. In *Proceedings of the 35th International Conference on Neural Information Processing Systems, NIPS ’21*, Red Hook, NY, USA, 2021. Curran Associates Inc.
- [3] Haya Diwan, Jinrui Gou, Cameron Musco, Christopher Musco, and Torsten Suel. Navigable graphs for high-dimensional nearest neighbor search: Constructions and limits. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 59513–59531. Curran Associates, Inc., 2024.
- [4] Zijian Dong, Yilei Wu, Chongyao Chen, Yingtian Zou, Yichi Zhang, and Juan Helen Zhou. Improve representation for imbalanced regression through geometric constraints. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- [5] Sedigheh Eslami and Gerard de Melo. Mitigate the gap: Improving cross-modal alignment in CLIP. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [6] Hao Guo and Youyou Lu. Achieving low-latency graph-based vector search via aligning best-first search algorithm with ssd. In *19th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, pages 171–186, Boston, MA, 2025. USENIX Association.
- [7] Hao Guo and Youyou Lu. Odinann: Direct insert for consistently stable performance in billion-scale graph-based vector search. In *24th USENIX Conference on File and Storage Technologies (FAST)*, Santa Clara, CA, 2026. USENIX Association.
- [8] Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for NLP. In Kamalika Chaudhuri and Ruslan Salakhutdinov, editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 2790–2799. PMLR, 09–15 Jun 2019.
- [9] Edward J Hu, yelong shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- [10] Jie Hu, Li Shen, and Gang Sun. Squeeze-and-excitation networks. In *CVPR*, 2018.

- [11] Piotr Indyk and Haike Xu. Worst-case performance of popular approximate nearest neighbor search implementations: Guarantees and limitations. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 66239–66256. Curran Associates, Inc., 2023.
- [12] Elias Jääsaari, Ville Hyvönen, and Teemu Roos. LoRANN: Low-rank matrix factorization for approximate nearest neighbor search. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [13] Dong Jing, Xiaolong He, Yutian Luo, Nanyi Fei, Guoxing Yang, Wei Wei, Huiwen Zhao, and Zhiwu Lu. FineCLIP: Self-distilled region-based CLIP for better fine-grained understanding. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [14] Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with gpus. *IEEE Transactions on Big Data*, 7(3):535–547, 2021.
- [15] Kate Keahey, Jason Anderson, Zhuo Zhen, Pierre Riteau, Paul Ruth, Dan Stanzione, Mert Cevik, Jacob Collieran, Haryadi S. Gunawi, Cody Hammock, Joe Mambretti, Alexander Barnes, François Halbach, Alex Rocha, and Joe Stubbs. Lessons learned from the chameleon testbed. In *Proceedings of the 2020 USENIX Annual Technical Conference (USENIX ATC '20)*. USENIX Association, July 2020.
- [16] Hyungbin Kim, Incheol Baek, and Yon Dohn Chung. Decoupled contrastive learning for federated learning, 2025.
- [17] Panagiotis Koromilas, Efthymios Georgiou, Giorgos Bouritsas, Theodoros Giannakopoulos, Mihalis A. Nicolaou, and Yannis Panagakis. A principled framework for multi-view contrastive learning, 2025.
- [18] Chiara Lanza, Roberto Pereira, Marco Miozzo, Eduard Angelats, and Paolo Dini. Degradation of feature space in continual learning, 2026.
- [19] Fengming Lin. Vision language models: A survey of 26k papers, 2025.
- [20] Avik Pal, Max van Spengler, Guido Maria D’Amely di Melendugno, Alessandro Flaborea, Fabio Galasso, and Pascal Mettes. Compositional entailment learning for hyperbolic vision-language models. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [21] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In Marina Meila and Tong Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event*, Proceedings of Machine Learning Research, pages 8748–8763. PMLR, 2021.
- [22] Sameera Ramasinghe, Violetta Shevchenko, Gil Avraham, and Ajanthan Thalaiyasingam. Accept the modality gap: An exploration in the hyperbolic space. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 27253–27262, 2024.
- [23] Simone Ricci, Niccolò Biondi, Federico Pernici, Ioannis Patras, and Alberto Del Bimbo. λ -orthogonality regularization for compatible representation learning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2026.
- [24] Daniele Savietto, Declan Campbell, André Panisson, Marco Nurişso, Giovanni Petri, Jonathan D. Cohen, and Alan Perotti. The geometry of representational failures in vision language models, 2026.
- [25] Shivam Rajendra Rai Sharma, Xiaoguang Zhu, Luca Cerny Oliveira, Kartik Patwari, La Rissa Vasquez, David Garcia, Louise Nicole C. Sevilla, Brittany N Dugger, and Chen-Nee Chuah. Benchmarking parameter efficient adaptation of vision language models on pathology. In *NeurIPS 2025 Workshop for Imageomics: Discovering Biological Knowledge from Images Using AI*, 2025.

- [26] Suhas Jayaram Subramanya, Devvrit, Rohan Kadekodi, Ravishankar Krishaswamy, and Harsha Vardhan Simhadri. Diskann: fast accurate billion-point nearest neighbor search on a single node. In *Proceedings of the 33rd International Conference on Neural Information Processing Systems*, 2019.
- [27] Hangke Sui, Yuqing Wang, and Minh N. Do. Unicon: Unified framework for efficient contrastive alignment via kernels. In *The Fourteenth International Conference on Learning Representations*, 2026.
- [28] Liang Wang, Xiang Tao, Qiang Liu, Shu Wu, and Liang Wang. Rethinking graph masked autoencoders through alignment and uniformity. In *Proceedings of the Thirty-Eighth AAAI Conference on Artificial Intelligence and Thirty-Sixth Conference on Innovative Applications of Artificial Intelligence and Fourteenth Symposium on Educational Advances in Artificial Intelligence*, AAAI’24/IAAI’24/EAAI’24. AAAI Press, 2024.
- [29] Mingqi Wu, Qiang Sun, and Archer Y. Yang. PCA++: How uniformity induces robustness to background noise in contrastive learning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2026.
- [30] Baili Xiao, Zhibin Dong, Ke Liang, Suyuan Liu, Siwei Wang, Tianrui Liu, Xingchen Hu, En Zhu, and Xinwang Liu. Easemvc: Efficient dual selection mechanism for deep multi-view clustering. In *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 20716–20726, 2025.
- [31] Xuqian Xue, Yiming Lei, Qi Cai, Hongming Shan, and Junping Zhang. PROTOCOL: Partial optimal transport-enhanced contrastive learning for imbalanced multi-view clustering. In *Forty-second International Conference on Machine Learning*, 2025.
- [32] Ming Yang, Yuzheng Cai, and Weiguo Zheng. Cspg: Crossing sparse proximity graphs for approximate nearest neighbor search. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 103076–103100. Curran Associates, Inc., 2024.
- [33] Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image diffusion models. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 3813–3824, 2023.
- [34] Renrui Zhang, Wei Zhang, Rongyao Fang, Peng Gao, Kunchang Li, Jifeng Dai, Yu Qiao, and Hongsheng Li. Tip-adapter: Training-free adaption of clip for few-shot classification. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXV*, page 493–510, Berlin, Heidelberg, 2022. Springer-Verlag.
- [35] Yingwen Zhang, Meng Wang, Xihua Sheng, Peilin Chen, Junru Li, Li Zhang, and Shiqi Wang. An information-theoretic regularizer for lossy neural image compression. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 15573–15582, October 2025.
- [36] Qingyan Zhao, Ruifang He, Jinpeng Zhang, Chang Liu, and Bo Wang. Representation de-generation problem in prompt-based models for natural language understanding. In Nicoletta Calzolari, Min-Yen Kan, Veronique Hoste, Alessandro Lenci, Sakriani Sakti, and Nianwen Xue, editors, *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 13946–13957, Torino, Italia, May 2024. ELRA and ICCL.
- [37] Xin Zhou, Yongjie Wang, and Zhiqi Shen. Cm³: Calibrating multimodal recommendation, 2025.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The paper presents both theoretical and experimental evidence to support its claims.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: We discuss the limitations of our approach in the conclusion.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Full proofs are provided in the appendix, and all assumptions are clearly stated in the main paper and/or the appendix.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All code and data needed to reproduce the main experimental results are available on the anonymous Github repository.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All the details to reproduce the main experimental results are provided in the anonymous Github repository, including instructions for data access and preparation, and scripts to reproduce all experimental results for the new proposed method and baselines.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: All such details are provided in the anonymous Github repository, and the main paper includes a summary of the experimental setting.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The results are pretty stable so we did not report error bars; we explain this in the main paper.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: These details are provided in the anonymous Github repository, and the main paper includes a summary of the compute resources used for the experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: We believe we have followed the NeurIPS Code of Ethics in all aspects of our research.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.

- 846
- 847
- 848
- 849
- 850
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
 - For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

851

16. Declaration of LLM usage

852

853

854

855

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

856

Answer: [N/A]

857

Justification: [N/A]

858

Guidelines:

- 859
- 860
- 861
- 862
- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
 - Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.

863 A Proof of Proposition 1

864 *Proof.* We prove each claim in turn.

865 **(1) Gradient support.** The triplet loss in Eq. (7) can be written as $\mathcal{L} = \frac{1}{|\mathcal{B}|} \sum_{(a,p,n) \in \mathcal{B}} \ell_{apn}$, where
 866 $\ell_{apn} = \max(0, \delta_{apn} + m)$ and $\delta_{apn} = d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n)$. The subgradient with respect to θ is

$$\nabla_{\theta} \ell_{apn} = \mathbf{1}[\delta_{apn} + m > 0] \cdot \nabla_{\theta} \delta_{apn}, \quad (12)$$

867 which is exactly zero whenever $\delta_{apn} + m \leq 0$. Therefore, $\nabla_{\theta} \mathcal{L}$ receives contributions only from
 868 active triplets satisfying Eq. (8).

869 **(2) Convergence of $\rho(t)$.** Under mild regularity conditions (Lipschitz continuity of f_{θ} and bounded
 870 embedding norms), the sequence $\{\hat{\mathbf{z}}^{(t)}\}$ converges to a stationary point $\hat{\mathbf{z}}^*$ of $\mathcal{L}$. At stationarity,
 871 $\rho(t)$ converges to ρ^* as defined in Eq. (10), since $\rho(t)$ is a continuous functional of the embedding
 872 distribution and the embeddings converge.

873 **(3) Zero gradient over unseen regions.** Let $\mathcal{R}_{\text{unseen}} \subset \mathbb{S}^{d-1}$ denote the region of the embedding
 874 space occupied by unseen-class images under f_{θ} . By construction, all triplets (a, p, n) in $\mathcal{D}_{\text{train}}$ satisfy
 875 $y_a, y_p, y_n \in \mathcal{Y}_{\text{seen}}$, so $\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p, \hat{\mathbf{z}}_n \notin \mathcal{R}_{\text{unseen}}$ for all training triplets. The gradient $\nabla_{\theta} \delta_{apn}$ depends on
 876 f_{θ} only through its evaluation at $\mathbf{z}_a, \mathbf{z}_p, \mathbf{z}_n$, which lie outside $\mathcal{R}_{\text{unseen}}$. Consequently, no training step
 877 produces a gradient that directly modifies f_{θ} in $\mathcal{R}_{\text{unseen}}$, and the CLIP geometry in that region is
 878 preserved throughout training.

879 **(4) Dependence on m .** The residual distance gap δ_{apn}^* at convergence is bounded by the within-class
 880 diameter $\Delta_c = \max_{x, x' \in c} d(\hat{\mathbf{z}}, \hat{\mathbf{z}}')$ for seen class c . For $m \leq \min_c \Delta_c$, the adapter can drive all
 881 within-class distances below the threshold, pushing ρ^* toward zero. For $m > \min_c \Delta_c$, some triplets
 882 remain active at convergence because the within-class spread exceeds the margin, resulting in a
 883 strictly positive ρ^* . This explains the empirically observed increase in ρ^* with larger m reported in
 884 Table 2. $\square$