
EGA: Adapting Frozen Encoders for Vector Search with Bounded Out-of-Distribution Degradation

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Vector search systems built on frozen vision encoders face queries from unseen
2 classes at deployment, yet existing adapter training collapses under this shift:
3 high-capacity adapters with global contrastive losses silently reassign unseen-class
4 samples to wrong seen-class clusters, dropping Label Precision by over 40 points
5 below the frozen baseline. We propose Euclidean Geodesic Alignment (EGA),
6 a residual adapter that couples three simple principles: zero initialization, local
7 triplet loss, and hypersphere projection. These collectively induce a self-limiting
8 dynamic: triplets that already satisfy a small margin stop producing gradients,
9 so the adapter automatically stops updating where the local geometry is already
10 correct. At convergence 96.5% of triplets are gradient-free, leaving unseen-class
11 regions largely untouched while still enabling full-capacity refinement of seen
12 classes. Across five diverse OOD benchmarks, EGA achieves the highest worst-
13 case Label Precision (0.611): the only method to exceed 0.6 on every single
14 benchmark. The design transfers directly to stronger backbones, and we provide
15 an analytical justification linking gradient sparsity to bounded OOD perturbation.

16 1 Introduction

17 A vector similarity search system [16, 30] deployed in production rarely operates on a fixed label
18 space. A retailer indexes product images today and adds a new category line next month; a content
19 platform builds an index over existing tags and continually ingests media from emerging topics; a
20 scientific image database is annotated for known phenotypes and queried for ones discovered later. In
21 such scenarios, the index and the query stream together cover a label space strictly larger than what
22 was available when any model was trained or adapted.

23 The deployment of unseen classes sits awkwardly with how adapters over frozen encoders [25] are
24 trained. The dominant paradigm takes a labeled subset, optimizes the embedding geometry to separate
25 those classes, and ships the resulting adapter as a drop-in replacement for the frozen encoder’s output.
26 The implicit assumption is that geometry tightened for seen classes will, at worst, leave unseen-class
27 geometry unchanged. We show this assumption fails systematically: high-capacity adapters trained
28 with global contrastive objectives reshape the entire embedding distribution, pulling unseen-class
29 samples toward whichever seen-class clusters they happen to resemble in the pre-trained metric. The
30 retrieval system inherits this distortion silently, since approximate nearest neighbor (ANN) indices
31 remain well-behaved geometrically while the neighbors belong to the wrong semantic class.

32 The system designer usually cannot know in advance which out-of-distribution (OOD) will dominate
33 the deployed query stream. The relevant criterion is therefore the *worst-case Label Precision* across
34 plausible OODs, not optimality on any single one. Two design choices in current adapter training
35 degrade this worst-case quantity. First, the loss function: global contrastive objectives such as
36 ICon [1] and SRL [4] compute gradients over the full pairwise similarity matrix in every step, so seen-

class supervision exerts pressure across the entire embedding distribution rather than only on labeled regions. Second, the adapter capacity: a sufficiently expressive adapter has the representational room to comply with that pressure, and zero-initialized residual designs do not by themselves prevent the cumulative perturbation from reaching unseen-class regions. Existing alternatives address one factor at a time. Low-rank adapters such as LoRA [10] cap the second factor, which avoids the worst OOD failures but limits in-distribution (ID) refinement, while keeping the global loss leaves the first factor untouched. We argue that adapters intended to operate under unseen-class queries should be designed against both factors in a coordinated manner.

This work presents Euclidean Geodesic Alignment (EGA), a residual adapter designed against both factors of loss function and adapter capacity, simultaneously. EGA pairs a high-capacity zero-initialized residual architecture with a small-margin triplet loss on the class-aware neighborhood graph of pre-trained embeddings, with ℓ_2 projection onto the unit hypersphere. Despite the technical intricacies, the underlying idea is straightforward, as follows. On the one hand, the triplet loss induces a self-limiting training dynamic: as the adapter converges on seen classes, the fraction of margin-violating triplets decays rapidly. On the other hand, the capacity to refine seen-class neighborhoods is preserved, but the loss stops applying pressure once seen-class neighborhoods are locally separated. As a result, the global reshaping that destroys unseen-class geometry never accumulates.

We evaluate EGA on CIFAR-100 and ImageNet-1K (in-distribution) and five OOD benchmarks (CIFAR-10, FGVC-Aircraft, Food-101, ImageNet-1K held-out classes, and Oxford-IIIT Pet), comparing against the frozen CLIP baseline [25], global contrastive adapters (Icon [1], SRL [4]), and LoRA [10] variants (InfoNCE [32] and Triplet, both at $r = 128$). Across the four primary OOD benchmarks, EGA achieves the highest worst-case Label Precision (0.611), 4.2 percentage points above the next-best method (LoRA+Triplet at 0.569), and is the only method whose Label Precision exceeds 0.6 on every one of them; on Oxford-IIIT Pet, where the frozen baseline already reaches 0.96, EGA matches this label precision while achieving the highest ANN Recall among all evaluated methods. Icon and SRL collapse to 0.470 and 0.375 in their worst case (Aircraft); LoRA+InfoNCE to 0.538 (Aircraft). The benefits transfer to stronger backbones, with EGA improving Label Precision@1 by 4.75 and 8.55 percentage points on DINOv2-large [23] and SigLIP [39].

In summary, this work makes the following contributions:

- We frame adapter tuning over frozen vision-language encoders from a perspective of vector search applications, in which the index and query stream together cover a label space larger than the labeled subset available for adapter training, and we identify worst-case OOD Label Precision as the relevant criterion under deployment uncertainty.
- We propose Euclidean Geodesic Alignment (EGA), a residual adapter design that couples zero-initialized residual architecture, local triplet supervision, and hypersphere projection. This combination induces a self-limiting training dynamic: as seen-class neighborhoods become locally separated, the fraction of active triplets decays rapidly, empirically reaching 96.5% zero-gradient triplets at steady state. We also provide a theoretical analysis of this dynamic, showing that the unseen-class perturbation is bounded by the product of the active triplet ratio and the adapter’s Lipschitz constant.
- Across four OOD benchmarks and three backbones (CLIP ViT-B/32, DINOv2-large, SigLIP), EGA achieves the highest worst-case OOD Label Precision (0.611), exceeding all five baselines by between 4.2 and 23.6 percentage points. EGA is the only method whose Label Precision exceeds 0.6 on every OOD benchmark we evaluate.

2 Background and Related Work

Foundation models such as CLIP [25] produce embedding spaces that encode both seen-class semantics and broader priors over unseen visual concepts. Related work has documented degradation phenomena in adapted representations [20, 37], and recent work documents analogous degradation in contrastive representations, including anisotropic collapse [43, 6], modality gap fracturing [26, 5], and loss of compositional structure [24, 28, 36]. Recent strategies counter these issues with geometric constraints such as orthogonality regularization [27]. *EGA exploits the gradient sparsity of small-margin triplet losses to keep adapter updates concentrated on seen-class neighborhoods, empirically leaving unseen-class regions of the pretrained geometry largely intact.*

Recent representation learning frequently relies on global contrastive objectives to optimize the latent space [31, 19]. A common paradigm jointly enforces alignment for positive pairs and uniformity across the feature distribution [34], applied in graph encoders [33] and multimodal recommendation [46]. While uniformity improves robustness to noise [35], it treats all latent regions equally. Other structural regularizers include unifying frameworks for multi-view alignment [1], geometric constraints for imbalanced regression [4], decoupled objectives [18], information-theoretic density control [42], and region-based distillation [15]. *EGA replaces dense global optimization with sparse local triplet updates, which empirically reduces unseen-class distortion.*

Parameter-efficient adaptation is an approach for transferring large foundation models without full fine-tuning [9, 10], which dominates vision-language research [21]. Lightweight methods, e.g., Visual Prompt Tuning [14] and Tip Adapter [41], specialize the model without modifying the frozen backbone and rival full fine-tuning on many tasks [29], while zero-initialization strategies further constrain the magnitude of pretrained-manifold perturbation [40]. *EGA admits a zero-initialized residual adapter, localized triplet objective, and parameter efficiency with gradient-sparse updates.*

Vector search is the foundational engine for scalable high-dimensional retrieval. The systems community has developed indexing algorithms for billion-scale datasets, including FAISS [16], HNSW [22], DiskANN [30], and SPANN [2]. Recent work refines graph topologies [38, 3], optimizes SSD-based search [7, 8], and accelerates distance computations [13], yet popular ANN implementations still suffer worst-case limitations [12]. *EGA recalibrates the metric geometry before index construction, rendering search backends achieve higher semantic recall without algorithmic changes.*

3 Methodology

3.1 Problem Formulation

Let $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$ denote a frozen vision-language encoder (e.g., CLIP ViT-B/32) that maps an image $x \in \mathcal{X}$ to a unit-normalized embedding $\mathbf{z} = \phi(x) \in \mathbb{S}^{d-1}$. We are given a set of labeled training images $\mathcal{D}_{\text{train}} = \{(x_i, y_i)\}_{i=1}^N$, where $y_i \in \mathcal{Y}_{\text{seen}}$ belongs to a set of *seen* classes. At inference time, the retrieval system must operate over a database that also contains images from a disjoint set of *unseen* classes $\mathcal{Y}_{\text{unseen}}$, with $\mathcal{Y}_{\text{seen}} \cap \mathcal{Y}_{\text{unseen}} = \emptyset$. Our goal is to learn an adapter $f_\theta : \mathbb{S}^{d-1} \rightarrow \mathbb{S}^{d-1}$ such that the transformed embeddings $\hat{\mathbf{z}} = f_\theta(\mathbf{z})$ support more accurate approximate nearest neighbor (ANN) retrieval on *both* seen and unseen classes, without access to unseen-class labels during training.

We evaluate retrieval quality along two complementary dimensions. *Label Precision@K* (LP@K) measures the semantic quality of retrieved results: whether the K nearest neighbors in the transformed space share the same class label as the query:

$$\text{LP@K} = \frac{1}{|\mathcal{Q}|} \sum_{q \in \mathcal{Q}} \frac{1}{K} \sum_{k=1}^K \mathbf{1}[y_{\hat{n}_k(q)} = y_q], \quad (1)$$

where $\mathcal{Q}$ is the query set and $\hat{n}_k(q)$ denotes the k -th nearest neighbor of q in the transformed embedding space. *ANNS Recall@K* (AR@K) measures the geometric fidelity of the ANN index, i.e., the fraction of true K nearest neighbors (under exact search) that are recovered by the approximate index at a given search budget nprobe:

$$\text{AR@K} = \frac{1}{|\mathcal{Q}|} \sum_{q \in \mathcal{Q}} \frac{|\mathcal{N}_K^{\text{exact}}(q) \cap \mathcal{N}_K^{\text{approx}}(q)|}{K}, \quad (2)$$

where $\mathcal{N}_K^{\text{exact}}(q)$ and $\mathcal{N}_K^{\text{approx}}(q)$ denote the exact and approximate top- K neighbor sets, respectively. The two metrics capture distinct failure modes: an adapter may improve LP@K by tightening within-class clusters (semantic improvement) while simultaneously improving AR@K by making the geometry more amenable to IVF-based indexing (geometric improvement).

3.2 Manifold-Preserving Adapter Design

The name Euclidean Geodesic Alignment (EGA) reflects the geometric intuition behind its three design principles. In the pre-trained embedding space, the shortest path (geodesic) between two points on the unit hypersphere $\mathbb{S}^{d-1}$ is the great-circle distance, which corresponds to cosine similarity. EGA

134 respects this intrinsic geometry in three ways: (1) The zero-initialized residual connection ensures
 135 that the adapter begins as an identity map, thus initially preserving all pairwise geodesic distances.
 136 (2) The local triplet loss only perturbs an embedding along the Euclidean chord toward its positive
 137 neighbor when the geodesic separation violates the margin; the perturbation is therefore aligned with
 138 the local geodesic direction. (3) The ℓ_2 projection back onto the hypersphere recovers a valid point
 139 on $\mathbb{S}^{d-1}$, ensuring that the output embedding remains in the same metric space. These principles
 140 allow EGA to perform local geodesic adjustments on the hypersphere without global distortion.

141 **(1) Residual connection with zero initialization.** Each EGA block applies a residual update as:

$$\mathbf{z}' = \mathbf{z} + g_\theta(\mathbf{z}), \quad (3)$$

142 where g_θ is a learned transformation initialized so that $g_\theta(\mathbf{z}) \approx \mathbf{0}$ at the start of training. Concretely,
 143 the final linear layer of g_θ is initialized to zero, making f_θ an identity mapping at initialization. This
 144 guarantees that training begins from the pretrained geometry rather than a random perturbation, and
 145 that the adapter can only *refine* rather than replace the CLIP embedding.

146 **(2) Feature gating.** Inside each residual block, the transformed signal passes through a module:

$$\text{Gate}(\mathbf{h}) = \mathbf{h} \odot \sigma(W_2 \text{GELU}(W_1 \mathbf{h})), \quad (4)$$

147 where $W_1 \in \mathbb{R}^{(d/4) \times d}$, $W_2 \in \mathbb{R}^{d \times (d/4)}$, σ is the Sigmoid function, and $\odot$ denotes elementwise
 148 multiplication. This is analogous to Squeeze-and-Excitation networks [11] to learn a per-dimension
 149 weight that amplifies retrieval-relevant dimensions and suppresses noisy ones. The bottleneck ratio
 150 $1/4$ is chosen to prevent the gating from overfitting on seen-class data, to be validated in Table 4.

151 The output of f_θ is projected back onto the unit hypersphere $\mathbb{S}^{d-1}$ via ℓ_2 :

$$\hat{\mathbf{z}} = \frac{f_\theta(\mathbf{z})}{\|f_\theta(\mathbf{z})\|_2}. \quad (5)$$

152 This constraint ensures that the transformed embeddings remain in the same metric space as the origi-
 153 nal embeddings: Euclidean distance and cosine similarity remain interchangeable on the hypersphere.

154 **(3) Full forward pass.** Let Block_i denote the i -th EGA block consisting of a two-layer MLP
 155 with LayerNorm followed by feature gating, and let Refine denote the final linear projection with
 156 LayerNorm. The complete forward pass for an input embedding $\mathbf{z} \in \mathbb{S}^{d-1}$ is:

$$\hat{\mathbf{z}} = \ell_2(\text{Refine}(\text{Block}_2(\text{Block}_1(\mathbf{z})))) . \quad (6)$$

157 The total number of trainable parameters is approximately $4.2M$ for $d = 512$ and hidden dimension
 158 2048, which is negligible compared to the frozen CLIP backbone ($\sim 88M$ parameters).

159 3.3 Local Triplet Training

160 Given a mini-batch $\mathcal{B}$ sampled from $\mathcal{D}_{\text{train}}$, we construct triplets (a, p, n) where a is an anchor image,
 161 p is a positive sampled uniformly from the same class as a , and n is a negative sampled uniformly
 162 from a different class. The adapter is trained with the standard triplet margin loss:

$$\mathcal{L} = \frac{1}{|\mathcal{B}|} \sum_{(a,p,n) \in \mathcal{B}} \max(0, d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m), \quad (7)$$

163 where $d(\cdot, \cdot)$ denotes distance, $m > 0$ is the margin, and $\hat{\mathbf{z}} = f_\theta(\mathbf{z})$ is the transformed embedding.

164 A triplet (a, p, n) contributes a non-zero gradient to θ if and only if the margin constraint is *violated*:

$$d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m > 0. \quad (8)$$

165 We refer to such triplets as *active*. As the adapter converges on seen classes, increasingly many triplets
 166 satisfy the margin constraint and become inactive, contributing zero gradient. This sparsity has a
 167 direct consequence for out-of-distribution generalization. Because unseen classes are absent from
 168 $\mathcal{D}_{\text{train}}$, no triplet involving an unseen-class embedding is ever constructed. Provided that the adapter’s
 169 updates remain effectively local (i.e., the cumulative perturbation induced by seen-class training
 170 gradients remains small in unseen-class regions of the embedding space), the manifold structure for




Figure 1: Label Precision@ K on CIFAR-100 across $K \in \{1, 3, 5, 10\}$ and $nprobe \in \{1, 5, 10\}$.

171 unseen classes is largely preserved. The gradient sparsity mechanism enforces this locality: once
 172 seen-class neighborhoods satisfy the margin, the corresponding gradients vanish, and only the small
 173 fraction of actively-violated local relationships continue to be updated.

174 The triplet loss in Eq. (7) consists of hinge terms whose subgradient with respect to θ is exactly zero
 175 whenever the margin constraint is satisfied:

$$\nabla_{\theta} \ell_{apn} = \mathbf{1}[d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m > 0] \cdot \nabla_{\theta} \delta_{apn}, \quad (9)$$

176 where $\delta_{apn} = d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n)$. Define the *active triplet ratio* at training step t as

$$\rho(t) = \Pr_{(a,p,n) \sim \mathcal{B}} \left[d(\hat{\mathbf{z}}_a^{(t)}, \hat{\mathbf{z}}_p^{(t)}) - d(\hat{\mathbf{z}}_a^{(t)}, \hat{\mathbf{z}}_n^{(t)}) + m > 0 \right]. \quad (10)$$

177 Under the standard assumption that training converges to a stationary point of $\mathcal{L}$, $\rho(t)$ converges to a
 178 steady-state value $\rho^* \in [0, 1]$ determined by the margin m and the within-class distance distribution of
 179 the converged embeddings. We measure $\rho^* \approx 3.5\%$ at convergence on FGVC-Aircraft (Section 4.3),
 180 meaning that 96.5% of training triplets contribute zero gradient at steady state.

181 Strictly speaking, sparse seen-class gradients do not formally imply that f_{θ} leaves unseen-class
 182 regions unchanged: the adapter is a parametric MLP, not a locally supported transformation, so
 183 any update to θ propagates globally. The OOD measurements in Section 4.2 will demonstrate the
 184 evidence: EGA preserves or improves Label Precision on three unseen-class benchmarks, while ICon
 185 and SRL lack the gradient sparsity and fall below the frozen baseline. A Lipschitz-based bound on
 186 the unseen-class perturbation is discussed in Appendix A.

187 4 Evaluation

188 Experiments were conducted on Chameleon Cloud [17] with 256 AMD EPYC-7763 CPU cores,
 189 512 GiB RAM, and an NVIDIA A100 GPU of 80 GB HBM2e. Datasets include CIFAR-100 and
 190 ImageNet-1K for in-distribution, and CIFAR-10, FGVC-Aircraft, Food-101, ImageNet-1K held-out
 191 classes, and Oxford-IIIT Pet for OOD evaluation; all are available via PyTorch or Kaggle. Backbones
 192 are CLIP ViT-B/32 [25], DINOv2-large [23], and SigLIP [39].

193 We compare EGA against four adapter methods: ICon [1], SRL [4], LoRA + InfoNCE [10, 32],
 194 and LoRA + Triplet [10], all trained on frozen CLIP features under identical splits and protocols.
 195 ICon and SRL use capacity-matched residual adapters (~ 4.2 M parameters) differing from EGA only
 196 in loss, isolating the training objective. LoRA variants use rank $r=128$, matching the parameter
 197 budget; we note that parameter count does not equate functional capacity (EGA permits nonlinear
 198 transformations while LoRA is low-rank), and our comparison controls the parameter budget rather
 199 than claiming architectural equivalence (see Section 4.3). We exclude full fine-tuning (it violates the
 200 frozen-backbone constraint required for index preservation and multi-task serving), as well as linear
 201 probing, Tip-Adapter [41], and prompt tuning [45, 44] (incompatible with our disjoint OOD setting).
 202 More training details and extended baseline results are in Appendices B and C.

203 4.1 In-Distribution Validation

204 Figure 1 reports Label Precision@ K across four neighborhood sizes and three search budgets on
 205 CIFAR-100. EGA raises Label Precision@1 from 0.549 (raw CLIP) to 0.705 at $nprobe = 1$, and




Figure 2: Distance distributions for Top- K neighbors vs. random background points on CIFAR-100.




Figure 3: ANNS Recall@ K versus nprobe on CIFAR-100.

the gap persists across all K and nprobe values, including nprobe = 10, indicating that the adapter tightens within-class neighborhoods rather than benefiting from broader index coverage.

The Label Precision improvement has a direct geometric correlate. Figure 2 contrasts the L_2 distance distributions for true Top- K neighbors against random background points. In the raw CLIP space, the two distributions overlap substantially, which is the condition under which IVF-based ANN indexing degrades. After EGA, the two distributions separate cleanly across all K , and this separation drives the recall-efficiency improvement in Figure 3: raw CLIP achieves 0.667 Recall@1 at nprobe = 1, whereas EGA reaches 0.825 and saturates above 0.99 by nprobe = 5.

Table 1 reports all six methods’ in-distribution performance on both CIFAR-100 and ImageNet-1K. On CIFAR-100, ICon and SRL reach near-perfect Label Precision (0.999 and 0.992) by tightening seen-class clusters under global contrastive objectives. EGA reaches 0.705, an improvement of +16 percentage points over the frozen baseline; the LoRA variants land between EGA and the baseline. On ImageNet-1K (1000 classes), the scale of all methods drops because the test set has approximately seven samples per class for retrieval (vs. ~ 150 on CIFAR-100). The within-2-percentage-point clustering of the four trained adapters (SRL, ICon, EGA, LoRA+Triplet) at this density indicates that ID performance becomes test-set-density-bounded rather than method-bounded as the label space grows.

Table 1: In-distribution retrieval performance at $K = 1$, nprobe = 1 on CIFAR-100 and ImageNet-1K. Standard errors are below 0.01 and omitted.

Method	CIFAR-100		ImageNet-1K	
	LP@1	AR@1	LP@1	AR@1
CLIP (frozen)	0.549	0.667	0.321	0.803
ICOn	0.999	0.992	0.386	0.817
SRL	0.992	0.991	0.421	0.687
LoRA+InfoNCE	0.668	0.879	0.426	0.822
LoRA+Triplet	0.612	0.908	0.359	0.870
EGA	0.705	0.825	0.407	0.817

4.2 Out-of-Distribution Generalization

We now examine performance under strict OOD settings where evaluation categories are disjoint from training categories. Table 2 reports headline numbers at $K = 1$, nprobe = 1 across four OOD

Table 2: OOD performance at $K = 1$, $nprobe = 1$.

Dataset	Metric	CLIP	ICoN	SRL	LoRA+InfoNCE	LoRA+Triplet	EGA
CIFAR-10	AR@1	0.863	<u>0.797</u> $\pm$.003	<u>0.819</u> $\pm$.001	0.869 $\pm$.012	0.861 $\pm$.011	<u>0.833</u> $\pm$.005
	LP@1	0.880	<u>0.560</u> $\pm$.010	<u>0.536</u> $\pm$.008	0.885 $\pm$.006	<u>0.875</u> $\pm$.006	<u>0.810</u> $\pm$.002
Aircraft	AR@1	0.773	0.853 $\pm$.016	0.879 $\pm$.023	0.891 $\pm$.015	0.893 $\pm$.008	0.905 $\pm$.015
	LP@1	0.512	<u>0.470</u> $\pm$.034	<u>0.375</u> $\pm$.032	0.538 $\pm$.020	0.569 $\pm$.011	0.611 $\pm$.020
Food-101	AR@1	0.842	<u>0.821</u> $\pm$.001	<u>0.824</u> $\pm$.015	0.906 $\pm$.003	0.893 $\pm$.008	0.879 $\pm$.011
	LP@1	0.881	<u>0.562</u> $\pm$.011	<u>0.552</u> $\pm$.010	<u>0.883</u> $\pm$.004	<u>0.833</u> $\pm$.007	<u>0.791</u> $\pm$.002
ImageNet-1K	AR@1	0.829	0.849 $\pm$.012	<u>0.747</u> $\pm$.015	0.878 $\pm$.008	0.883 $\pm$.009	0.868 $\pm$.011
	LP@1	0.684	<u>0.620</u> $\pm$.014	<u>0.665</u> $\pm$.012	0.753 $\pm$.009	0.714 $\pm$.010	0.711 $\pm$.008
Worst-case	LP@1	0.512	0.470	0.375	0.538	0.569	0.611

233 benchmarks: CIFAR-10 (trained on CIFAR-100), FGVC-Aircraft (80 seen / 20 unseen classes),
 234 Food-101 (80 seen / 21 unseen classes), and ImageNet-1K held-out classes (800 seen / 200 unseen
 235 classes). For a fair comparison, all methods train identical adapters on top of frozen CLIP features
 236 and differ only in their loss functions. Robustness across search budgets is reported in Appendix C.1.

237 **Worst-case OOD Label Precision.** The bottom row of Table 2 reports the worst-case Label
 238 Precision across the four OOD benchmarks. EGA achieves 0.611, exceeding LoRA+Triplet (0.569)
 239 by 4.2 percentage points, LoRA+InfoNCE (0.538) by 7.3 percentage points (pp), the frozen CLIP
 240 baseline (0.512) by 9.9 pp, ICoN (0.470) by 14.1 pp, and SRL (0.375) by 23.6 pp. EGA is the only
 241 method whose Label Precision exceeds 0.6 on every OOD benchmark we evaluate.

242 The four OOD benchmarks span qualitatively different shifts: CIFAR-10 is a cross-dataset shift;
 243 Aircraft and Food-101 are fine-grained intra-domain shifts where seen and unseen classes share sub-
 244 stantial visual structure; ImageNet held-out classes is a generic large-scale shift where the pretrained
 245 encoder already provides reasonable transfer. The worst-case reflects deployment uncertainty: the
 246 system designer commits to one adapter without knowing in advance which shift will dominate the
 247 query stream. Methods that excel on some shifts but collapse on others (ICoN collapses on Aircraft,
 248 SRL collapses on Aircraft, LoRA variants on Aircraft) cannot guarantee a quality floor.

249 While these four benchmarks span the major categories of distribution shift, the worst-case is driven
 250 by the fine-grained Aircraft split. As a further probe, we evaluate on Oxford-IIIT Pet (37-class
 251 pet breed classification) under the same OOD protocol: the frozen CLIP baseline already reaches
 252 LP@1 = 0.96, and all adapter methods perform comparably; EGA remains the best-performing
 253 method qualitatively, preserving the pretrained geometry without degradation (detailed numbers in
 254 Appendix C.5). This supports the interpretation that EGA’s advantage is specific to hard fine-grained
 255 shifts where the pretrained geometry is genuinely insufficient to separate unseen categories, rather
 256 than to fine-grained datasets in general.

257 **Per-benchmark behavior.** On ANNS Recall, all trained adapters match or exceed the frozen
 258 baseline across Aircraft, Food-101, and ImageNet, confirming that adapter training successfully
 259 reorganizes the geometry for ANN indexing. On Label Precision, however, the global contrastive
 260 methods fall below the frozen baseline on every fine-grained OOD dataset, with degradation exceeding
 261 30 percentage points on Food-101 and CIFAR-10: the retrieved neighbors are geometrically closer
 262 but semantically wrong. EGA improves Label Precision over the frozen baseline on Aircraft and
 263 ImageNet, and its degradations on CIFAR-10 and Food-101 are contained to single-digit percentage
 264 points. The collapse of global contrastive methods is most severe exactly where unseen classes are
 265 visually similar to seen classes, consistent with our hypothesis that dense global gradients pull unseen
 266 samples into incorrect seen-class clusters. On ImageNet, where held-out classes are visually distinct
 267 from training classes, no method collapses, and all approaches, including the frozen baseline, perform
 268 comparably. EGA’s advantage materializes precisely under difficult distribution shifts, which is the
 269 scenario where worst-case guarantees are most valuable.

270 EGA does not dominate every method on every individual benchmark. Methods that restrict adapter
 271 capacity rather than the loss function achieve higher Label Precision on CIFAR-10, Food-101, and
 272 ImageNet, but suffer catastrophic degradation on the fine-grained Aircraft split, which is what




Figure 4: ID/OOD operating points of adapter training in the (ID, OOD) Label Precision plane.

determines their worst-case performance. EGA’s design explicitly prioritizes worst-case robustness; practitioners who can confidently bound the deployed distribution shift may reasonably prefer capacity-restricted alternatives. The trade-off between these two approaches is visualized in Figure 4.

4.3 Two Routes to OOD Stability: Gradient Sparsity and Capacity Restriction

Adapters that avoid OOD collapse must neutralize one of these factors, and the route determines its position in the Label Precision plane.

Route 1: Gradient sparsity (EGA).

A triplet loss with margin m produces a non-zero gradient for a sample triplet (a, p, n) only when the margin constraint is violated:

$$d(a, p) - d(a, n) + m > 0. \quad (11)$$

As the adapter converges on seen classes, the fraction of *active* triplets decays. Figure 5 (left) tracks this ratio from 17.4% at epoch 10 to 3.5% at convergence. At steady state, 96.5% of triplets contribute zero gradient: the loss stops applying pressure once seen-class neighborhoods are locally separated, even though the adapter retains the capacity to refine them further. Sparsity neutralizes the loss-side factor while leaving capacity intact, which permits EGA’s high in-distribution Label Precision in Table 1. The self-limiting dynamic is consistent across datasets: the steady-state active triplet ratio ρ^* is 2.1% on CIFAR-100 and 0.8% on Food-101, compared to 3.5% on FGVC-Aircraft, see Appendix C.4. The higher residual ratio on Aircraft reflects the fine-grained visual similarity among its classes, which leaves a larger fraction of triplets close to the margin boundary.




Figure 5: Gradient sparsity as OOD protection.

Route 2: Capacity restriction (LoRA). The LoRA variants take the other route. They retain a global contrastive loss (InfoNCE) or a triplet loss, but constrain the adapter to a low-rank residual update. At rank $r = 128$, LoRA’s expressive power is limited to rank-128 perturbations of the identity, so the dense gradient field has substantially less room to reshape unseen-class geometry. This explains why LoRA+InfoNCE, despite using the same loss family as ICon, avoids the same OOD collapse: the loss applies pressure everywhere, but the adapter cannot fully comply. The price is paid in-distribution. At capacity-matched scale (~ 4.2 M parameters), LoRA+InfoNCE reaches 0.668 Label Precision on CIFAR-100, 3.7 percentage points below EGA, because the same low-rank constraint that bounds OOD damage also bounds in-distribution refinement.

The two routes occupy different points in Figure 4. EGA, which neutralizes the loss-side factor while keeping capacity, achieves the highest worst-case OOD Label Precision. Capacity-matched LoRA, which neutralizes the capacity-side factor, excels on three of four benchmarks but collapses on Aircraft (worst-case 0.538–0.569). These routes are complementary: removing gradient sparsity from EGA drops Aircraft LP@1 to 0.458 (Appendix C.3), while scaling LoRA rank from 64 to 512 leaves both ID and OOD nearly unchanged. OOD robustness thus requires the combination of gradient

Table 3: Ablation study of EGA

Variant	Accuracy
Full EGA (Ours)	0.705
w/o Residual connection	0.0065
w/o Zero-initialization	0.6840
w/o L2-normalization	0.6590

Table 4: Sensitivity of different hyperparameters

Triplet Margin m	0.1	0.2	0.3	0.5
Recall@1	0.8420	0.8180	0.8160	0.7400
Recall@5	0.9782	0.9654	0.9610	0.9125
Hidden Dimension d	512	1024	2048	4096
Recall@1	0.8300	0.8340	0.8200	0.8240
Recall@5	0.9695	0.9712	0.9680	0.9688

sparsity and sufficient nonlinear capacity. EGA’s in-distribution performance reflects this design priority rather than underfitting: increasing the margin would relax gradient sparsity and degrade OOD, so the reported ID scores represent a deliberate choice to preserve worst-case robustness.

4.4 Ablation and Sensitivity Analysis

Table 3 isolates the contribution of each architectural component. The residual connection is the core element: removing it collapses accuracy from 0.705 to 0.0065, indicating that the adapter cannot reconstruct semantic relationships from target data alone and must operate as a refinement over the pre-trained features. ℓ_2 -normalization keeps the transformed embeddings on the unit hypersphere, and removing it lowers accuracy to 0.6590. Zero-initialization lets the adapter begin as an identity mapping, and removing it lowers accuracy to 0.6840.

Table 4 evaluates stability across hyperparameter variations. EGA maintains consistent quality for margins $m \in \{0.1, 0.2, 0.3\}$, but increasing m to 0.5 degrades Recall@1 to 0.7400: a larger margin forces the adapter to update embeddings that already satisfy local neighborhood constraints, reducing gradient sparsity. Performance is also stable across hidden dimensions from 512 to 4096, indicating that EGA’s behavior is determined by its structural design rather than by specific parameter choices.

4.5 Generalization to More Backbones

To verify that EGA’s benefits are not specific to CLIP ViT-B/32 [25], we further evaluate it on two more recent vision encoders: DINOv2-large [23] and SigLIP [39].

Table 5 shows consistent improvements across all three backbones. The largest gains appear on the relatively weaker CLIP ViT-B/32 (+14.65 pp in Label Precision@1 and +13.15 pp in Recall@1), but the stronger DINOv2-large and SigLIP backbones also see stable improvements (+4.75 pp and +8.55 pp in Label Precision@1 respectively).

Table 5: Retrieval performance on more backbones.

Backbone	Label Precision@1		ANNS Recall@1	
	Raw	+EGA	Raw	+EGA
CLIP ViT-B/32	0.549	0.6955	0.667	0.7985
DINOv2-large	0.8530	0.9005	0.8785	0.9310
SigLIP SO400M	0.7740	0.8595	0.8015	0.9155

5 Conclusion

We studied adapter training for vector search under open-category queries, where the proper metric is worst-case OOD Label Precision rather than in-distribution optimality. Global contrastive adapters silently distort unseen-class geometry; EGA avoids this through a self-limiting triplet loss that vanishes when local class structure is satisfied, automatically preserving unseen regions. Across five OOD benchmarks spanning cross-dataset, fine-grained, and large-scale shifts, EGA achieves the highest worst-case retrieval quality, is the only method above 0.6 on every primary benchmark, and transfers directly to multiple backbones.

Limitations The theoretical bound on unseen-class perturbation requires the empirically observed steady-state active triplet ratio as an input; a closed-form guarantee without this remains open. EGA prioritizes worst-case robustness, so methods that restrict capacity can excel on individual benchmarks, and practitioners who know their deployment shift may prefer those alternatives. Our worst-case analysis covers five benchmarks spanning distinct shift types; the absolute worst-case is determined by the hardest available fine-grained dataset, and additional challenging domains remain a natural extension.

References

- [1] Shaden Naif Alshammari, John R Hershey, Axel Feldmann, William T Freeman, and Mark Hamilton. I-con: A unifying framework for representation learning. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- [2] Qi Chen, Bing Zhao, Haidong Wang, Mingqin Li, Chuanjie Liu, Zengzhong Li, Mao Yang, and Jingdong Wang. Spann: highly-efficient billion-scale approximate nearest neighbor search. In *Proceedings of the 35th International Conference on Neural Information Processing Systems*, NIPS ’21, Red Hook, NY, USA, 2021. Curran Associates Inc.
- [3] Haya Diwan, Jinrui Gou, Cameron Musco, Christopher Musco, and Torsten Suel. Navigable graphs for high-dimensional nearest neighbor search: Constructions and limits. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 59513–59531. Curran Associates, Inc., 2024.
- [4] Zijian Dong, Yilei Wu, Chongyao Chen, Yingtian Zou, Yichi Zhang, and Juan Helen Zhou. Improve representation for imbalanced regression through geometric constraints. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- [5] Sedigheh Eslami and Gerard de Melo. Mitigate the gap: Improving cross-modal alignment in CLIP. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [6] Tianyu Gao, Xingcheng Yao, and Danqi Chen. SimCSE: Simple contrastive learning of sentence embeddings. In Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih, editors, *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 6894–6910, Online and Punta Cana, Dominican Republic, November 2021. Association for Computational Linguistics.
- [7] Hao Guo and Youyou Lu. Achieving low-latency graph-based vector search via aligning best-first search algorithm with ssd. In *19th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, pages 171–186, Boston, MA, 2025. USENIX Association.
- [8] Hao Guo and Youyou Lu. Odinann: Direct insert for consistently stable performance in billion-scale graph-based vector search. In *24th USENIX Conference on File and Storage Technologies (FAST)*, Santa Clara, CA, 2026. USENIX Association.
- [9] Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for NLP. In Kamalika Chaudhuri and Ruslan Salakhutdinov, editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 2790–2799. PMLR, 09–15 Jun 2019.
- [10] Edward J Hu, yelong shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuezhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- [11] Jie Hu, Li Shen, Samuel Albanie, Gang Sun, and Enhua Wu. Squeeze-and-excitation networks. *IEEE Trans. Pattern Anal. Mach. Intell.*, 42(8):2011–2023, August 2020.
- [12] Piotr Indyk and Haike Xu. Worst-case performance of popular approximate nearest neighbor search implementations: Guarantees and limitations. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 66239–66256. Curran Associates, Inc., 2023.
- [13] Elias Jäsaari, Ville Hyvönen, and Teemu Roos. LoRANN: Low-rank matrix factorization for approximate nearest neighbor search. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [14] Menglin Jia, Luming Tang, Bor-Chun Chen, Claire Cardie, Serge Belongie, Bharath Hariharan, and Ser-Nam Lim. Visual prompt tuning. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXIII*, page 709–727, Berlin, Heidelberg, 2022. Springer-Verlag.

- [15] Dong Jing, Xiaolong He, Yutian Luo, Nanyi Fei, Guoxing Yang, Wei Wei, Huiwen Zhao, and Zhiwu Lu. FineCLIP: Self-distilled region-based CLIP for better fine-grained understanding. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [16] Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with gpus. *IEEE Transactions on Big Data*, 7(3):535–547, 2021.
- [17] Kate Keahey, Jason Anderson, Zhuo Zhen, Pierre Riteau, Paul Ruth, Dan Stanzione, Mert Cevik, Jacob Colleran, Haryadi S. Gunawi, Cody Hammock, Joe Mambretti, Alexander Barnes, François Halbach, Alex Rocha, and Joe Stubbs. Lessons learned from the chameleon testbed. In *Proceedings of the 2020 USENIX Annual Technical Conference (USENIX ATC '20)*. USENIX Association, July 2020.
- [18] Hyungbin Kim, Incheol Baek, and Yon Dohn Chung. Decoupled contrastive learning for federated learning, 2025.
- [19] Panagiotis Koromilas, Efthymios Georgiou, Giorgos Bouritsas, Theodoros Giannakopoulos, Mihalis A. Nicolaou, and Yannis Panagakis. A principled framework for multi-view contrastive learning, 2025.
- [20] Chiara Lanza, Roberto Pereira, Marco Miozzo, Eduard Angelats, and Paolo Dini. Degradation of feature space in continual learning, 2026.
- [21] Fengming Lin. Vision language models: A survey of 26k papers, 2025.
- [22] Yu A. Malkov and D. A. Yashunin. Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs. *IEEE Trans. Pattern Anal. Mach. Intell.*, 42(4):824–836, April 2020.
- [23] Maxime Oquab, Timothée Darcet, Théo Moutakanni, Huy V. Vo, Marc Szafraniec, Vasil Khalidov, Pierre Fernandez, Daniel HAZIZA, Francisco Massa, Alaaeldin El-Nouby, Mido Assran, Nicolas Ballas, Wojciech Galuba, Russell Howes, Po-Yao Huang, Shang-Wen Li, Ishan Misra, Michael Rabbat, Vasu Sharma, Gabriel Synnaeve, Hu Xu, Herve Jegou, Julien Mairal, Patrick Labatut, Armand Joulin, and Piotr Bojanowski. DINOv2: Learning robust visual features without supervision. *Transactions on Machine Learning Research*, 2024. Featured Certification.
- [24] Avik Pal, Max van Spengler, Guido Maria D’Amely di Melendugno, Alessandro Flaborea, Fabio Galasso, and Pascal Mettes. Compositional entailment learning for hyperbolic vision-language models. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [25] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In Marina Meila and Tong Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event*, Proceedings of Machine Learning Research, pages 8748–8763. PMLR, 2021.
- [26] Sameera Ramasinghe, Violetta Shevchenko, Gil Avraham, and Ajanthan Thalaiyasingam. Accept the modality gap: An exploration in the hyperbolic space. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 27253–27262, 2024.
- [27] Simone Ricci, Niccolò Biondi, Federico Pernici, Ioannis Patras, and Alberto Del Bimbo. λ -orthogonality regularization for compatible representation learning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2026.
- [28] Daniele Savietto, Declan Campbell, André Panisson, Marco Nurisso, Giovanni Petri, Jonathan D. Cohen, and Alan Perotti. The geometry of representational failures in vision language models, 2026.
- [29] Shivam Rajendra Rai Sharma, Xiaoguang Zhu, Luca Cerny Oliveira, Kartik Patwari, La Rissa Vasquez, David Garcia, Louise Nicole C. Sevilla, Brittany N Dugger, and Chen-Nee Chuah. Benchmarking parameter efficient adaptation of vision language models on pathology. In *NeurIPS 2025 Workshop for Imageomics: Discovering Biological Knowledge from Images Using AI*, 2025.

- [30] Suhas Jayaram Subramanya, Devvrit, Rohan Kadekodi, Ravishankar Krishaswamy, and Harsha Vardhan Simhadri. Diskann: fast accurate billion-point nearest neighbor search on a single node. In *Proceedings of the 33rd International Conference on Neural Information Processing Systems*, 2019.
- [31] Hangke Sui, Yuqing Wang, and Minh N. Do. Unicon: Unified framework for efficient contrastive alignment via kernels. In *The Fourteenth International Conference on Learning Representations*, 2026.
- [32] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. *arXiv preprint arXiv:1807.03748*, 2018.
- [33] Liang Wang, Xiang Tao, Qiang Liu, Shu Wu, and Liang Wang. Rethinking graph masked autoencoders through alignment and uniformity. In *Proceedings of the Thirty-Eighth AAAI Conference on Artificial Intelligence and Thirty-Sixth Conference on Innovative Applications of Artificial Intelligence and Fourteenth Symposium on Educational Advances in Artificial Intelligence, AAAI’24/IAAI’24/EAAI’24*. AAAI Press, 2024.
- [34] Tongzhou Wang and Phillip Isola. Understanding contrastive representation learning through alignment and uniformity on the hypersphere. In *Proceedings of the 37th International Conference on Machine Learning, ICML’20*. JMLR.org, 2020.
- [35] Mingqi Wu, Qiang Sun, and Archer Y. Yang. PCA++: How uniformity induces robustness to background noise in contrastive learning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2026.
- [36] Baili Xiao, Zhibin Dong, Ke Liang, Suyuan Liu, Siwei Wang, Tianrui Liu, Xingchen Hu, En Zhu, and Xinwang Liu. Easemvc: Efficient dual selection mechanism for deep multi-view clustering. In *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 20716–20726, 2025.
- [37] Xuqian Xue, Yiming Lei, Qi Cai, Hongming Shan, and Junping Zhang. PROTOCOL: Partial optimal transport-enhanced contrastive learning for imbalanced multi-view clustering. In *Forty-second International Conference on Machine Learning*, 2025.
- [38] Ming Yang, Yuzheng Cai, and Weiguo Zheng. Cspg: Crossing sparse proximity graphs for approximate nearest neighbor search. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 103076–103100. Curran Associates, Inc., 2024.
- [39] Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language image pre-training. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 11941–11952, 2023.
- [40] Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image diffusion models. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 3813–3824, 2023.
- [41] Renrui Zhang, Wei Zhang, Rongyao Fang, Peng Gao, Kunchang Li, Jifeng Dai, Yu Qiao, and Hongsheng Li. Tip-adaptor: Training-free adaption of clip for few-shot classification. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXV*, page 493–510, Berlin, Heidelberg, 2022. Springer-Verlag.
- [42] Yingwen Zhang, Meng Wang, Xihua Sheng, Peilin Chen, Junru Li, Li Zhang, and Shiqi Wang. An information-theoretic regularizer for lossy neural image compression. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 15573–15582, October 2025.
- [43] Qingyan Zhao, Ruifang He, Jinpeng Zhang, Chang Liu, and Bo Wang. Representation degeneration problem in prompt-based models for natural language understanding. In Nicoletta Calzolari, Min-Yen Kan, Veronique Hoste, Alessandro Lenci, Sakriani Sakti, and Nianwen Xue, editors, *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 13946–13957, Torino, Italia, May 2024. ELRA and ICCL.

- 505 [44] Kaiyang Zhou, Jingkang Yang, Chen Change Loy, and Ziwei Liu. Conditional prompt learning
506 for vision-language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision*
507 *and Pattern Recognition*, pages 16816–16825, 2022.
- 508 [45] Kaiyang Zhou, Jingkang Yang, Chen Change Loy, and Ziwei Liu. Learning to prompt for
509 vision-language models. *Int. J. Comput. Vision*, 130(9):2337–2348, September 2022.
- 510 [46] Xin Zhou, Yongjie Wang, and Zhiqi Shen. Cm³: Calibrating multimodal recommendation,
511 2025.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [\[Yes\]](#)

Justification: The paper presents both theoretical and experimental evidence to support its claims.

Guidelines:

- The answer [\[N/A\]](#) means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [\[No\]](#) or [\[N/A\]](#) answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [\[Yes\]](#)

Justification: We discuss the limitations of our approach in the conclusion.

Guidelines:

- The answer [\[N/A\]](#) means that the paper has no limitation while the answer [\[No\]](#) means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [\[Yes\]](#)

Justification: Full proofs are provided in the appendix, and all assumptions are clearly stated in the main paper and/or the appendix.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All code and data needed to reproduce the main experimental results are available on the anonymous Github repository.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

617 Question: Does the paper provide open access to the data and code, with sufficient instruc-
618 tions to faithfully reproduce the main experimental results, as described in supplemental
619 material?

620 Answer: [Yes]

621 Justification: All the details to reproduce the main experimental results are provided in the
622 anonymous Github repository, including instructions for data access and preparation, and
623 scripts to reproduce all experimental results for the new proposed method and baselines.

624 Guidelines:

- 625 • The answer [N/A] means that paper does not include experiments requiring code.
- 626 • Please see the NeurIPS code and data submission guidelines ([https://neurips.cc/
627 public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 628 • While we encourage the release of code and data, we understand that this might not
629 be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not
630 including code, unless this is central to the contribution (e.g., for a new open-source
631 benchmark).
- 632 • The instructions should contain the exact command and environment needed to run to
633 reproduce the results. See the NeurIPS code and data submission guidelines ([https:
634 //neurips.cc/public/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.
- 635 • The authors should provide instructions on data access and preparation, including how
636 to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- 637 • The authors should provide scripts to reproduce all experimental results for the new
638 proposed method and baselines. If only a subset of experiments are reproducible, they
639 should state which ones are omitted from the script and why.
- 640 • At submission time, to preserve anonymity, the authors should release anonymized
641 versions (if applicable).
- 642 • Providing as much information as possible in supplemental material (appended to the
643 paper) is recommended, but including URLs to data and code is permitted.

644 6. Experimental setting/details

645 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
646 rameters, how they were chosen, type of optimizer) necessary to understand the results?

647 Answer: [Yes]

648 Justification: All such details are provided in the anonymous Github repository, and the
649 main paper includes a summary of the experimental setting.

650 Guidelines:

- 651 • The answer [N/A] means that the paper does not include experiments.
- 652 • The experimental setting should be presented in the core of the paper to a level of detail
653 that is necessary to appreciate the results and make sense of them.
- 654 • The full details can be provided either with the code, in appendix, or as supplemental
655 material.

656 7. Experiment statistical significance

657 Question: Does the paper report error bars suitably and correctly defined or other appropriate
658 information about the statistical significance of the experiments?

659 Answer: [Yes]

660 Justification: The results are pretty stable so we did not report error bars; we explain this in
661 the main paper.

662 Guidelines:

- 663 • The answer [N/A] means that the paper does not include experiments.
- 664 • The authors should answer [Yes] if the results are accompanied by error bars, confidence
665 intervals, or statistical significance tests, at least for the experiments that support the
666 main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: These details are provided in the anonymous Github repository, and the main paper includes a summary of the compute resources used for the experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: We believe we have followed the NeurIPS Code of Ethics in all aspects of our research.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.

- 822 • We recognize that the procedures for this may vary significantly between institutions
823 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
824 guidelines for their institution.
825 • For initial submissions, do not include any information that would break anonymity (if
826 applicable), such as the institution conducting the review.

827 **16. Declaration of LLM usage**

828 Question: Does the paper describe the usage of LLMs if it is an important, original, or
829 non-standard component of the core methods in this research? Note that if the LLM is used
830 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
831 scientific rigor, or originality of the research, declaration is not required.

832 Answer: [N/A]

833 Justification: [N/A]

834 Guidelines:

- 835 • The answer [N/A] means that the core method development in this research does not
836 involve LLMs as any important, original, or non-standard components.
837 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
838 be described.

A Theoretical Analysis of EGA’s OOD Stability

We provide a theoretical analysis to explain why the proposed EGA adapter outperforms other popular adapters in worst-case out-of-distribution (OOD) scenarios. The analysis is structured in four steps, starting from the geometric interpretation of the adapter architecture to deriving a Lipschitz-based bound on the unseen-class perturbation.

1. Manifold Projection and Identity Initialization

Let the pre-trained feature space be defined as a unit hypersphere manifold $\mathcal{M} = \mathbb{S}^{d-1} \subset \mathbb{R}^d$. For any input feature $x \in \mathcal{M}$, the residual adapter learns a continuous perturbation vector field $h_\theta : \mathcal{M} \rightarrow \mathbb{R}^d$. The forward pass of our architecture, which combines a residual connection with L_2 normalization, is geometrically equivalent to a retraction mapping $\Pi_{\mathcal{M}}$ that projects the perturbed point back onto the manifold:

$$f_\theta(x) = \Pi_{\mathcal{M}}(x + h_\theta(x)) = \frac{x + h_\theta(x)}{\|x + h_\theta(x)\|_2}. \quad (12)$$

Under our zero initialization strategy, the initial perturbation field is strictly zero everywhere, meaning $h_{\theta_0}(x) \equiv 0$. Consequently, the initial adapter state represents a strict identity mapping on the manifold, $f_{\theta_0}(x) = x$. This forms the theoretical origin for computing cumulative structural distortions.

2. Path Integral of Unseen Class Distortion

We model the training process as continuous gradient descent where parameters evolve according to $\frac{d\theta_t}{dt} = -\eta \nabla_\theta \mathcal{L}(\theta_t)$, with learning rate η . For any given unseen class sample $x_u \in \mathcal{X}_{unseen}$, its total geometric drift accumulated over the training period T can be formulated using the fundamental theorem of calculus as a path integral along the parameter trajectory:

$$\Delta(x_u) = \|f_{\theta_T}(x_u) - f_{\theta_0}(x_u)\|_2 = \left\| \int_0^T J_{\theta_t}(x_u) \frac{d\theta_t}{dt} dt \right\|_2, \quad (13)$$

where $J_{\theta_t}(x_u) = \nabla_\theta f_{\theta_t}(x_u)$ denotes the Jacobian matrix of the network evaluated at x_u . Applying the triangle inequality for integrals, we establish the fundamental upper bound for geometric distortion:

$$\Delta(x_u) \leq \int_0^T \left\| J_{\theta_t}(x_u) \frac{d\theta_t}{dt} \right\|_2 dt. \quad (14)$$

3. Sparsity Modulated Deformation Bound

In our formulation, the parameter derivative is entirely driven by the active triplet ratio ρ_t . Let N be the batch size. The loss gradient is the average over active triplets. Assuming the gradient norm per individual sample is bounded by G , the batch gradient norm satisfies:

$$\|\nabla_\theta \mathcal{L}(\theta_t)\|_2 = \left\| \frac{1}{N} \sum_{i \in \text{active}} \nabla_\theta \ell_i \right\|_2 \leq \frac{1}{N} \sum_{i \in \text{active}} \|\nabla_\theta \ell_i\|_2 \leq \frac{\rho_t N}{N} G = \rho_t G. \quad (15)$$

This directly bounds the parameter update rate:

$$\left\| \frac{d\theta_t}{dt} \right\|_2 = \eta \|\nabla_\theta \mathcal{L}(\theta_t)\|_2 \leq \eta \rho_t G. \quad (16)$$

To process the integrand from the previous step, we apply the sub-multiplicativity of matrix norms. Assuming the network satisfies L Lipschitz continuity such that the spectral norm of the Jacobian is bounded by $\|J_{\theta_t}(x_u)\|_2 \leq L$, we obtain:

$$\left\| J_{\theta_t}(x_u) \frac{d\theta_t}{dt} \right\|_2 \leq \|J_{\theta_t}(x_u)\|_2 \left\| \frac{d\theta_t}{dt} \right\|_2 \leq L(\eta \rho_t G). \quad (17)$$

Substituting this back into the path integral yields the global perturbation upper bound for any unseen sample:

$$\Delta_{EGA}(x_u) \leq \eta L G \int_0^T \rho_t dt. \quad (18)$$

For global contrastive methods, the equivalent active ratio remains at a constant level near 1, resulting in a perturbation bound that grows linearly with time ηLGT . In contrast, our gradient sparsity mechanism forces ρ_t to decay rapidly to a marginal steady state value. The bounded nature of this integral mathematically locks the maximum possible geometric shift for unseen categories.

4. Geometric Intuition via Subspace Dimensionality.

To explain why the residual active gradients do not coincidentally distort visually similar unseen classes, we analyze the dimensionality of the update space. The parameter update direction $\frac{d\theta_t}{dt}$ resides entirely within a tangent subspace $\mathcal{S}_{active}^{(t)}$ spanned by the gradients of the active triplets. Let P_S denote the orthogonal projection operator onto this active subspace. We can rewrite the update vector as $P_{\mathcal{S}_{active}^{(t)}} \frac{d\theta_t}{dt}$. The true instantaneous perturbation is therefore modulated by the projection of the unseen Jacobian onto this subspace:

$$\left\| J_{\theta_t}(x_u) \frac{d\theta_t}{dt} \right\|_2 = \left\| J_{\theta_t}(x_u) P_{\mathcal{S}_{active}^{(t)}} \frac{d\theta_t}{dt} \right\|_2 \leq \left\| J_{\theta_t}(x_u) P_{\mathcal{S}_{active}^{(t)}} \right\|_2 \left\| \frac{d\theta_t}{dt} \right\|_2. \quad (19)$$

As the active triplet ratio ρ_t drops exponentially, the dimensionality of $\mathcal{S}_{active}^{(t)}$ collapses. Consequently, the spectral norm of the projected Jacobian $\left\| J_{\theta_t}(x_u) P_{\mathcal{S}_{active}^{(t)}} \right\|_2$ probabilistically approaches zero. Under the empirically observed decay of ρ_t , the dimensionality of $\mathcal{S}_{active}^{(t)}$ is small at convergence, and the spectral norm of the projected Jacobian $\left\| J_{\theta_t}(x_u) P_{\mathcal{S}_{active}^{(t)}} \right\|_2$ is correspondingly small in practice. This provides geometric intuition for why the sparse updates are unlikely to systematically corrupt the structure of unseen classes.

B Experimental Setup Details

Backbones models. All experiments in the main paper use frozen CLIP ViT-B/32 [25] image embeddings as the base representation. To verify that EGA’s benefits are not specific to this backbone, we further evaluate EGA on two stronger vision encoders: DINOv2-large [23] and SigLIP [39]. DINOv2-large is a self-supervised vision transformer trained on 142M images, while SigLIP is a supervised vision transformer trained on 400M image-text pairs. Both models produce higher-quality embeddings than CLIP ViT-B/32, as reflected in their stronger raw retrieval performance, and EGA continues to provide meaningful improvements on top of these stronger backbones, as shown in Table 5 in Section 4.4.

Adapter models. We compare EGA against the frozen CLIP baseline and four trained adapters, all sharing the same training data, evaluation protocol, and 75/25 retrieval split:

- **ICon** [1]: a global contrastive objective minimizing the KL divergence between the empirical similarity distribution and the class-membership distribution.
- **SRL** [4]: a structural regularization objective jointly optimizing batch-wide uniformity and within-class homogeneity, also instantiated on the EGAMLP architecture.
- **LoRA + InfoNCE** [10, 32]: a capacity-matched low-rank adapter (rank $r = 128$, $\sim 4.2M$ parameters) with zero-initialized residual structure, trained with supervised InfoNCE.
- **LoRA + Triplet** [10]: identical capacity-matched low-rank architecture as above (rank $r = 128$), but trained with the same small-margin ($m = 0.2$) triplet loss as EGA.

Datasets. We evaluate on six datasets covering both in-distribution and OOD retrieval scenarios. **CIFAR-100** (100 classes, 60K images) serves as a primary in-distribution benchmark with a 75/25 split between database and query sets. **ImageNet-1K** (1000 classes, 35K embedded samples) serves as a larger-scale in-distribution benchmark and additionally provides an unseen-class OOD setting via an 80/20 class split (800 seen / 200 unseen). **CIFAR-10** (10 classes, 60K images) provides a cross-dataset OOD setting: the adapter is trained on CIFAR-100 and evaluated on the disjoint CIFAR-10 classes. **FGVC-Aircraft** (100 fine-grained aircraft variants, 10K images) and **Food-101** (101 food categories, 25K images for the test split we use) provide unseen-class OOD settings: classes are randomly partitioned into 80% seen / 20% unseen with a fixed seed, the adapter is trained on




Figure 6: Label Precision@1 across three OOD benchmarks and three search budgets ($nprobe \in \{1, 5, 10\}$).

the seen split, and retrieval is evaluated strictly on the unseen split. **Oxford-IIIT Pet** (37 pet breeds, $\sim 7K$ images) provides an additional fine-grained OOD setting under the same 80/20 class-disjoint protocol.

Indexing and metrics. For all retrieval evaluations we use FAISS [16] IndexIVFFlat with $nlist=10$ centroids and report results at $nprobe \in \{1, 5, 10\}$. We measure both Label Precision ($LP@K$, Eq. 1) and ANNS Recall ($AR@K$, Eq. 2) at $K \in \{1, 3, 5, 10\}$.

Training. All adapters are trained with AdamW (weight decay 10^{-4}), cosine-annealed learning rate, and batch size 256 (CIFAR-100) or 512 (ImageNet-1K). EGAMLP-based methods (EGA, ICon, SRL) use learning rate 10^{-4} ; LoRA variants use 10^{-3} to compensate for their smaller per-parameter scale. Triplet-based methods sample one positive and one negative per anchor uniformly within each mini-batch; global contrastive methods (ICoN, SRL, LoRA + InfoNCE) operate on the full pairwise similarity matrix. All experiments use three fixed random seed (42, 123, 456) for reproducibility.

Compute platform. All experiments were conducted on the Chameleon Cloud [17] platform using a compute node equipped with 256 AMD EPYC-7763 CPU cores, 512 GiB of RAM, and an NVIDIA A100 GPU of 80 GB HBM2e. The operating system is Ubuntu 24.04 LTS.

C Additional Experimental Results

C.1 Robustness across search budgets.

Figure 6 traces Label Precision@1 across three OOD benchmarks and three search budgets ($nprobe \in \{1, 5, 10\}$). The collapse pattern of ICoN and SRL persists across all 18 measurement points (3 datasets $\times$ 3 budgets $\times$ 2 methods), ruling out the possibility that the degradation is an artifact of any particular benchmark or indexing configuration. EGA’s relative ordering against other methods is also stable across budgets.

C.2 Linear Probing Baseline Analysis

To address potential concerns about additional baselines, we also evaluated Linear Probing as an example. As shown in Table 6 and Table 7, Linear Probing achieves reasonable in-distribution performance (0.5112 LP@1 on CIFAR-100) but completely collapses on out-of-distribution benchmarks (0.0065–0.0133 LP@1). This confirms that Linear Probing is not a meaningful comparison in our cross-dataset OOD setting due to class semantic mismatch.

Table 6: Linear Probing in-distribution performance on CIFAR-100 (75/25 split, mean $\pm$ std over 3 seeds).

Method	LP@1	AR@1
Linear Probing	0.5112 ± 0.0020	0.7533 ± 0.0025

Table 7: Linear Probing OOD performance (mean $\pm$ std over 3 seeds).

Dataset	LP@1	AR@1
Aircraft	0.0133 $\pm$ 0.0038	0.4401 $\pm$ 0.0254
Food-101	0.0077 $\pm$ 0.0019	0.5162 $\pm$ 0.0228
CIFAR-10	0.0065 $\pm$ 0.0005	0.6454 $\pm$ 0.0096

945 C.3 Disentangling Loss and Architecture

946 We conduct two controlled analyses to isolate the contributions of gradient sparsity and architecture
947 capacity to EGA’s OOD robustness.

948 **Architecture without gradient sparsity.** We train EGAMLP (the same residual architecture as
949 EGA, ~ 4.2 M parameters) with the global contrastive InfoNCE loss under the same protocols. Table 8
950 reports OOD Label Precision@1. Under the global objective, the architecture degrades substantially
951 across all shifts: Aircraft drops to 0.458, below even the frozen baseline (0.512), and Food-101 and
952 CIFAR-10 lose 21.4 and 10 points relative to EGA. This confirms that the residual design alone does
953 not prevent unseen-class distortion; gradient sparsity from the triplet loss is essential.

Table 8: EGAMLP trained with InfoNCE on OOD benchmarks ($K=1$, $nprobe=1$).

OOD Benchmark	EGAMLP+InfoNCE	EGA (Triplet)	Frozen CLIP
FGVC-Aircraft	0.458	0.611	0.512
Food-101	0.667	0.791	0.881
CIFAR-10	0.710	0.810	0.880

954 **Gradient sparsity with restricted capacity.** We next examine whether increasing the capacity
955 of a low-rank adapter can recover EGA’s OOD robustness when using the triplet loss. We train
956 LoRA+Triplet at ranks $r \in \{64, 128, 256, 512\}$ and measure both in-distribution (CIFAR-100) and
957 OOD (Aircraft) Label Precision@1. As shown in Table 9, ID performance barely improves with rank
958 (0.6456 at $r=64$ to 0.6480 at $r=512$), remaining far below EGA’s 0.705, while Aircraft LP@1 never
959 exceeds 0.566 and even deteriorates at higher ranks. Thus, the low-rank structure intrinsically caps
960 both ID refinement and OOD protection, regardless of parameter budget.

Table 9: LoRA+Triplet at different ranks ($K=1$, $nprobe=1$).

LoRA Rank	CIFAR-100 LP@1	Aircraft LP@1
64	0.6456	0.5655
128	0.6460	0.5595
256	0.6472	0.5655
512	0.6480	0.5298
EGA (full)	0.7050	0.6110

961 Together, these results establish that the combination of gradient-sparse triplet updates and a high-
962 capacity nonlinear residual architecture is necessary for the OOD robustness observed in EGA, and
963 that neither component alone suffices.

964 C.4 Consistency of gradient sparsity across datasets.

965 Table 10 reports the steady-state active triplet ratio ρ^* for EGA trained on three datasets. The fraction
966 of zero-gradient triplets exceeds 96% in all cases, confirming that the self-limiting dynamic is not
967 specific to Aircraft. The variation in ρ^* aligns with dataset difficulty: fine-grained Aircraft leaves the
968 largest residual active fraction, while Food-101, whose classes are more visually distinct, converges
969 to near-complete sparsity.

Table 10: Steady-state active triplet ratio ρ^* across datasets.

Dataset	ρ^*	Zero-gradient fraction
FGVC-Aircraft	3.5%	96.5%
CIFAR-100	2.1%	97.9%
Food-101	0.8%	99.2%

970 C.5 Additional Fine-Grained Benchmark: Oxford-IIIT Pet

971 We further evaluate EGA and LoRA+Triplet on Oxford-IIIT Pet under the same 80/20 class-disjoint
 972 OOD protocol. Table 11 reports LP@1 and AR@1 at $K=1$, nprobe=1. The frozen CLIP baseline
 973 achieves LP@1 = 0.96 on unseen breeds, indicating that the pretrained encoder already separates
 974 these categories well; all adapter methods perform comparably to the baseline, with EGA retaining
 975 the highest ANNS Recall. This is consistent with our finding that EGA’s advantage materializes on
 976 hard fine-grained shifts where the pretrained geometry is insufficient, while on easier shifts it simply
 977 preserves the already-strong frozen representation.

Table 11: Oxford-IIIT Pet OOD performance ($K=1$, nprobe=1, 80/20 class split).

Method	LP@1	AR@1
Frozen CLIP	0.9595	0.8785
LoRA+Triplet	0.9595	0.9342
EGA	0.9595	0.9544