
When Adapter Training Hurts: Manifold Preservation for Vector Retrieval

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Adapting pre-trained vision-language models for vector retrieval typically relies
2 on global contrastive objectives, which optimize semantic alignment across all
3 training categories simultaneously. We identify a fundamental limitation of this
4 approach: when applied through lightweight adapters, global objectives produce
5 dense gradient fields that systematically distort the latent manifold of unseen
6 categories, causing out-of-distribution (OOD) retrieval performance to collapse
7 below the frozen baseline. Although parameter-efficient low-rank adapters such as
8 LoRA largely preserve performance when trained with global losses like InfoNCE
9 or triplet objectives due to their capacity bottleneck, they do not exploit local
10 sparse updates to the same extent. To address this, we propose Euclidean Geodesic
11 Alignment (EGA), a residual adapter trained with a small-margin triplet loss on
12 the class-aware neighborhood graph of pre-trained embeddings. By operating
13 locally rather than globally, EGA converges to a sparse-gradient state in which
14 96.5% of training triplets produce zero gradient at steady state, leaving the broader
15 embedding space, including regions occupied by unseen categories, structurally
16 intact. Extensive evaluations on CIFAR-10, CIFAR-100, FGVC-Aircraft, and Food-
17 101 demonstrate that EGA consistently outperforms global contrastive methods
18 including ICon and SRL across both in-distribution and OOD settings while also
19 surpassing or matching the LoRA+InfoNCE and LoRA+Triplet baselines. On
20 strictly unseen categories, EGA improves Label Precision over the frozen baseline
21 on Aircraft and CIFAR-10 and bounds the degradation on Food-101, while ICon
22 and SRL fall below the frozen baseline by up to 36%.

23 1 Introduction

24 [Dongfang: TODO] Modern retrieval systems built on foundation models such as CLIP [1] face a
25 fundamental challenge: the pretrained encoder is general-purpose, but each downstream retrieval task
26 requires embeddings tuned to its own category structure. Full fine-tuning is rarely viable due to the
27 size of foundation backbones and the risk of destroying their zero-shot capabilities. The dominant
28 solution is parameter-efficient adaptation through lightweight adapters trained on a labeled subset of
29 the target task while the backbone remains frozen.

30 In practice, adapter training is almost universally driven by global contrastive objectives—methods
31 that compute losses over the full pairwise similarity matrix in each mini-batch and continuously
32 reshape the entire embedding distribution. Such objectives have proven effective for in-distribution
33 classification, where all relevant categories appear during training. However, retrieval systems must
34 also perform well on *unseen* categories that arrive at deployment time but were absent from adapter
35 training. We show that in this out-of-distribution (OOD) regime, the empirical consequence is striking:
36 when paired with a high-capacity adapter architecture, global objectives such as ICon [2] and SRL [3]
37 produce embeddings that are *geometrically* more amenable to IVF-based indexing on unseen-class

queries, yet retrieve neighbors of the wrong semantic class—falling 4–36 percentage points below the frozen baseline in Label Precision on every OOD benchmark we tested. Adapter training under these conditions trades semantic correctness for geometric indexability.

We propose Euclidean Geodesic Alignment (EGA), a manifold-preserving adapter that resolves this tension by replacing dense global optimization with sparse local updates. EGA combines three design choices: (i) a zero-initialized residual structure, (ii) an ℓ_2 normalization constraint, and (iii) a small-margin triplet loss on the class-aware neighborhood graph, each of which restricts the adapter’s influence to local neighborhood relationships that explicitly violate the retrieval objective. As training proceeds, the fraction of triplets producing non-zero gradient decays rapidly: at convergence, 96.5% of triplets contribute zero gradient. The adapter ceases to modify the vast majority of the embedding space, preserving the pretrained semantic structure for unseen categories with bounded degradation under heavy OOD shift.

We evaluate EGA across in-distribution (CIFAR-100) and three OOD benchmarks (CIFAR-10, FGVC-Aircraft, Food-101). On in-distribution data, EGA improves retrieval quality over the frozen baseline. On strictly unseen categories, EGA improves Label Precision over the frozen baseline on Aircraft and CIFAR-10 and bounds the degradation on Food-101, while ICon and SRL fall below the frozen baseline on every dataset—on Food-101, by 32–36 percentage points. Concretely on FGVC-Aircraft, EGA reaches Label Precision 0.548 versus 0.423/0.441 for ICon/SRL, with the frozen baseline at 0.512. We trace this contrast to the gradient sparsity property and argue that it constitutes a general design principle for retrieval adapters built on foundation models.

1.1 Contributions

This paper makes the following contributions:

- We identify a fundamental in-distribution/OOD tradeoff in global contrastive adapter training: while global objectives achieve higher geometric compression on seen categories, they produce dense gradient fields that systematically distort the pretrained manifold of unseen categories, causing retrieval performance to fall below the frozen CLIP baseline on out-of-distribution data.
- We propose EGA, a lightweight residual adapter combining zero initialization, ℓ_2 normalization, and a small-margin triplet loss. By restricting optimization to locally violated neighborhood constraints, EGA refines the pretrained embedding space without destroying its global semantic structure, simultaneously improving retrieval quality on seen categories and preserving manifold integrity for unseen ones.
- We provide a mechanistic explanation grounded in gradient sparsity: EGA’s triplet objective converges to a regime where 96.5% of training triplets contribute zero gradient, leaving unseen class regions structurally unchanged from CLIP. We validate this mechanism empirically and demonstrate through evaluations on three OOD benchmarks (CIFAR-10, FGVC-Aircraft, Food-101) that ICon and SRL fall below the frozen CLIP baseline in Label Precision on every dataset—catastrophically so on Food-101, where they degrade by 32–36 percentage points—while EGA improves Label Precision on Aircraft and CIFAR-10 with bounded degradation on Food-101.

2 Related Work

Manifold Preservation under Adapter Tuning Foundation models such as CLIP [1] produce embedding spaces whose structure encodes both seen class semantics and a broader prior over visual concepts not observed during downstream training. Catastrophic forgetting has been studied extensively in continual and incremental learning settings [4, 5]. Recent work has documented analogous degradation phenomena specific to contrastive representations: anisotropic collapse into narrow conic regions [6, 7], structural fracturing under modality gap [8, 9], and loss of compositional structure under aggressive optimization [10, 11, 12]. To counter these degradations, recent adaptation strategies impose explicit geometric constraints such as orthogonality regularization [13] or restrict optimization to limited subspaces. EGA leverages the inherent sparsity of triplet loss optimization to confine adapter updates to local neighborhoods of seen classes, leaving the unseen class regions of the pretrained manifold structurally unmodified.

Global Contrastive Objectives and Structural Regularization Recent advancements in representation learning often rely on global contrastive objectives to optimize the latent space [14, 15]. A common paradigm involves jointly enforcing alignment for positive pairs and uniformity for the overall feature distribution [16], a strategy applied across various domains including graph encoders [17] and multimodal recommendation systems [18]. While pushing embeddings toward a uniform distribution enhances robustness to background noise [19], it inherently assumes that all latent regions require the same degree of global scattering. To further refine complex distributions, structural regularization techniques introduce auxiliary constraints. For instance, methods like ICon formulate a unifying framework to systematically align multi view representations [2], while geometric constraints are leveraged to improve imbalanced regression tasks [3]. Other approaches attempt to decouple the learning objectives [20] or apply information theoretic regularizers [21] to control the density of the learned representations. Furthermore, region based distillation techniques have been proposed to enhance fine grained understanding [22]. EGA replaces dense global optimization with sparse local triplet updates, preserving out of distribution semantic structure.

Parameter-Efficient Manifold Adaptation Parameter efficient adaptation emerged as a standard to prevent catastrophic forgetting and reduce computational overhead when transferring large foundation models [23, 24]. This paradigm has extensively shifted towards vision language models, dominating recent literature due to the prohibitive cost of full scale fine tuning [25]. Methods such as Visual Prompt Tuning [26] and Tip Adapter demonstrate that lightweight tuning can effectively specialize representations for downstream tasks without altering the frozen backbone [27]. Furthermore, extensive benchmarking confirms that parameter efficient techniques can rival full fine tuning while preserving the generalizable representations of the foundation model [28]. To further protect the pre trained manifold during initial training steps, zero initialization strategies have proven highly effective. For instance, conditional control mechanisms in image generation initialize adapter weights to zero, ensuring the model behaves exactly like the frozen base model before any optimization occurs [29]. EGA integrates a zero initialized residual adapter with a localized geometric objective, achieving both parameter efficiency and targeted manifold smoothing.

Vector Similarity Search and Metric Consistency Vector similarity search is the foundational engine for retrieving high dimensional representations at scale. The systems community has developed highly optimized indexing algorithms to handle billion scale datasets. Foundational works such as FAISS leverage hardware acceleration for fast similarity search [30], and Hierarchical Navigable Small World graphs establish the core topology for efficient approximate retrieval [31]. Meanwhile, hybrid structures like DiskANN and SPANN achieve exceptional speed by efficiently managing memory and disk hierarchies [32, 33]. Recent advancements continue to refine these topologies, exploring crossing sparse proximity graphs [34] and analyzing the theoretical constructions of navigable graphs for high dimensional nearest neighbor search [35]. Furthermore, system level studies optimize retrieval latency and stability by aligning best first search algorithms with solid state drives [36] and enabling direct insertions in large scale indices [37]. To reduce the computational cost of exact distance calculations during the retrieval phase, approximation techniques such as low rank matrix factorization have also been introduced to accelerate inner product evaluations [38]. However, extensive theoretical analyses reveal that popular approximate nearest neighbor implementations still face severe worst case performance limitations [39]. EGA proactively recalibrates the metric manifold prior to index construction, enabling standard search backends to achieve superior semantic recall without requiring complex algorithmic reengineering.

3 Methodology

3.1 Problem Formulation

Let $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$ denote a frozen vision-language encoder (e.g., CLIP ViT-B/32) that maps an image $x \in \mathcal{X}$ to a unit-normalized embedding $\mathbf{z} = \phi(x) \in \mathbb{S}^{d-1}$. We are given a set of labeled training images $\mathcal{D}_{\text{train}} = \{(x_i, y_i)\}_{i=1}^N$, where $y_i \in \mathcal{Y}_{\text{seen}}$ belongs to a set of *seen* classes. At inference time, the retrieval system must operate over a database that also contains images from a disjoint set of *unseen* classes $\mathcal{Y}_{\text{unseen}}$, with $\mathcal{Y}_{\text{seen}} \cap \mathcal{Y}_{\text{unseen}} = \emptyset$.

Our goal is to learn a lightweight adapter $f_\theta : \mathbb{S}^{d-1} \rightarrow \mathbb{S}^{d-1}$ such that the transformed embeddings $\hat{\mathbf{z}} = f_\theta(\mathbf{z})$ support more accurate approximate nearest neighbor (ANN) retrieval on *both* seen and unseen classes, without access to any unseen-class labels during training.

Evaluation metrics. We evaluate retrieval quality along two complementary dimensions.

Label Precision@K measures the semantic quality of retrieved results—whether the K nearest neighbors in the transformed space share the same class label as the query:

$$\text{LP@K} = \frac{1}{|\mathcal{Q}|} \sum_{q \in \mathcal{Q}} \frac{1}{K} \sum_{k=1}^K \mathbf{1}[y_{\hat{n}_k(q)} = y_q], \quad (1)$$

where $\mathcal{Q}$ is the query set and $\hat{n}_k(q)$ denotes the k -th nearest neighbor of q in the transformed embedding space.

ANNS Recall@K measures the geometric fidelity of the ANN index—the fraction of true K nearest neighbors (under exact search) that are recovered by the approximate index at a given search budget n_{probe} :

$$\text{AR@K} = \frac{1}{|\mathcal{Q}|} \sum_{q \in \mathcal{Q}} \frac{|\mathcal{N}_K^{\text{exact}}(q) \cap \mathcal{N}_K^{\text{approx}}(q)|}{K}, \quad (2)$$

where $\mathcal{N}_K^{\text{exact}}(q)$ and $\mathcal{N}_K^{\text{approx}}(q)$ denote the exact and approximate top- K neighbor sets, respectively.

The two metrics capture distinct failure modes: an adapter may improve LP@K by tightening within-class clusters (semantic improvement) while simultaneously improving AR@K by making the geometry more amenable to IVF-based indexing (geometric improvement). As we show in Section 4, global contrastive objectives achieve high AR@K on seen classes but collapse on *both* metrics for unseen classes, whereas EGA improves both simultaneously under the OOD setting.

3.2 Manifold-Preserving Adapter Design

EGA introduces a lightweight adapter f_θ composed of stacked residual blocks with a final refinement layer. The architecture is deliberately constrained by three design principles, each of which serves a distinct role in preserving the pretrained manifold.

Residual connection with zero initialization. Each EGA block applies a residual update of the form

$$\mathbf{z}' = \mathbf{z} + g_\theta(\mathbf{z}), \quad (3)$$

where g_θ is a learned transformation initialized so that $g_\theta(\mathbf{z}) \approx \mathbf{0}$ at the start of training. Concretely, the final linear layer of g_θ is initialized to zero, making f_θ an identity mapping at initialization. This guarantees that training begins from the pretrained geometry rather than a random perturbation, and that the adapter can only *refine* rather than replace the CLIP embedding. As confirmed by our ablation (Table 2), removing the residual connection causes accuracy to collapse from 0.70 to 0.01, because the lightweight MLP cannot reconstruct complex semantic relationships from seen-class data alone.

Feature gating. Inside each residual block, the transformed signal passes through a feature-gating module:

$$\text{Gate}(\mathbf{h}) = \mathbf{h} \odot \sigma(W_2 \text{GELU}(W_1 \mathbf{h})), \quad (4)$$

where $W_1 \in \mathbb{R}^{(d/4) \times d}$, $W_2 \in \mathbb{R}^{d \times (d/4)}$, σ is the sigmoid function, and $\odot$ denotes elementwise multiplication. This structure, analogous to Squeeze-and-Excitation networks [40], learns a per-dimension importance weight that amplifies retrieval-relevant dimensions and suppresses noisy ones. The bottleneck ratio of $1/4$ is chosen to prevent the gating network from overfitting on seen-class data, as validated in Table 3.

ℓ_2 normalization. The final output of f_θ is projected back onto the unit hypersphere $\mathbb{S}^{d-1}$ via ℓ_2 normalization:

$$\hat{\mathbf{z}} = \frac{f_\theta(\mathbf{z})}{\|f_\theta(\mathbf{z})\|_2}. \quad (5)$$

179 This constraint ensures that the transformed embeddings remain in the same metric space as the
 180 original CLIP embeddings, so that Euclidean distance and cosine similarity remain interchangeable
 181 on the hypersphere. Removing this constraint degrades retrieval performance by 4.2 percentage
 182 points (Table 2), confirming that hypersphere consistency is necessary for stable ANN indexing.

183 **Full forward pass.** Let Block_i denote the i -th EGA block consisting of a two-layer MLP with Lay-
 184 erNorm followed by feature gating, and let Refine denote the final linear projection with LayerNorm.
 185 The complete forward pass for an input embedding $\mathbf{z} \in \mathbb{S}^{d-1}$ is:

$$\hat{\mathbf{z}} = \ell_2(\text{Refine}(\text{Block}_2(\text{Block}_1(\mathbf{z})))) . \quad (6)$$

186 The total number of trainable parameters is approximately $4.2M$ for $d = 512$ and hidden dimension
 187 2048, which is negligible compared to the frozen CLIP backbone ($\sim 150M$ parameters).

188 3.3 Local Triplet Training and Gradient Sparsity

189 **Training objective.** Given a mini-batch sampled from $\mathcal{D}_{\text{train}}$, we construct triplets (a, p, n) where a
 190 is an anchor image, p is a positive sampled uniformly from the same class as a , and n is a negative
 191 sampled uniformly from a different class. The adapter is trained with the standard triplet margin loss:

$$\mathcal{L} = \frac{1}{|\mathcal{B}|} \sum_{(a,p,n) \in \mathcal{B}} \max(0, d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m), \quad (7)$$

192 where $d(\cdot, \cdot)$ denotes Euclidean distance, $m > 0$ is the margin hyperparameter, and $\hat{\mathbf{z}} = f_\theta(\mathbf{z})$ is the
 193 transformed embedding. We set $m = 0.2$ throughout all experiments based on the sensitivity analysis
 194 in Table 3.

195 **Gradient sparsity and manifold preservation.** A triplet (a, p, n) contributes a non-zero gradient
 196 to θ if and only if the margin constraint is *violated*:

$$d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m > 0. \quad (8)$$

197 We refer to such triplets as *active*. As the adapter converges on seen classes, increasingly many
 198 triplets satisfy the margin constraint and become inactive, contributing zero gradient. Figure 5 (left)
 199 shows that the active triplet ratio decays from 17.4% at epoch 10 to 3.5% at convergence. At steady
 200 state, **96.5% of triplets produce zero gradient**, meaning the adapter ceases to modify the vast
 201 majority of local neighborhood relationships in the embedding space.

202 This sparsity has a direct consequence for out-of-distribution generalization. Because unseen classes
 203 are absent from $\mathcal{D}_{\text{train}}$, no triplet involving an unseen-class embedding is ever constructed. Provided
 204 that the adapter’s updates remain local—i.e., that gradients from seen-class triplets do not propagate
 205 into unseen-class regions of the embedding space—the CLIP manifold structure for unseen classes
 206 is preserved intact. The gradient sparsity mechanism enforces exactly this locality: once seen-class
 207 neighborhoods satisfy the margin, the corresponding gradients vanish, and only the small fraction of
 208 actively-violated local relationships continue to be updated.

209 **Contrast with global contrastive objectives.** Global objectives such as ICon [2] and SRL [3]
 210 compute losses over the *full* pairwise similarity matrix within each mini-batch. ICon minimizes a
 211 KL divergence between the empirical similarity distribution and the class-membership distribution,
 212 while SRL jointly optimizes batch-wide uniformity and within-class homogeneity. Both objectives
 213 generate dense gradient fields in every training step: every sample pair—regardless of whether
 214 its local geometry already satisfies the retrieval objective—contributes to the loss and receives a
 215 non-zero gradient. This global optimization pressure from seen-class training propagates throughout
 216 the embedding space, distorting the local manifold structure that CLIP had established for unseen
 217 categories.

218 **Proposition 1** (Gradient Sparsity at Convergence). *Let f_θ be trained with the triplet loss in Eq. (7)*
 219 *with margin $m > 0$ on a dataset $\mathcal{D}_{\text{train}}$ drawn from seen classes $\mathcal{Y}_{\text{seen}}$. Define the active triplet ratio*
 220 *at training step t as*

$$\rho(t) = \Pr_{(a,p,n) \sim \mathcal{B}} \left[d(\hat{\mathbf{z}}_a^{(t)}, \hat{\mathbf{z}}_p^{(t)}) - d(\hat{\mathbf{z}}_a^{(t)}, \hat{\mathbf{z}}_n^{(t)}) + m > 0 \right]. \quad (9)$$




Figure 1: Label Precision@ K on CIFAR-100 across $K \in \{1, 3, 5, 10\}$ and $nprobe \in \{1, 5, 10\}$. EGA consistently improves over the raw CLIP latent space.

As f_θ converges, $\rho(t) \rightarrow \rho^*$, where

$$\rho^* = \Pr_{(a,p,n)} [d(\hat{\mathbf{z}}_a^*, \hat{\mathbf{z}}_p^*) - d(\hat{\mathbf{z}}_a^*, \hat{\mathbf{z}}_n^*) + m > 0] \quad (10)$$

is determined solely by the margin m and the within-class distance distribution of the converged embeddings. Furthermore, $\nabla_\theta \mathcal{L} = \mathbf{0}$ for all triplets (a, p, n) with $d(\hat{\mathbf{z}}_a^*, \hat{\mathbf{z}}_p^*) - d(\hat{\mathbf{z}}_a^*, \hat{\mathbf{z}}_n^*) + m \leq 0$, so the adapter gradient is supported only on the fraction ρ^* of the training distribution. For embeddings whose local geometry already satisfies the margin constraint, the adapter exerts zero update pressure.

Proof Sketch. The triplet loss consists of hinge terms whose gradient is zero whenever the margin constraint is satisfied. At convergence, only those triplets whose geometry the adapter has not yet separated within the margin remain active. The fraction of such triplets is determined solely by the margin m and the distance distribution of seen classes, independent of unseen classes. Since no training triplet involves unseen-class embeddings, the adapter exerts zero gradient over the entire unseen-class region of the embedding space. See Appendix A for the complete proof. $\square$

4 Evaluation

4.1 Experimental Setup

Experiments were conducted on the Chameleon Cloud platform using a compute node equipped with 256 AMD EPYC-7763 CPU cores, 512 GiB of RAM, and an NVIDIA A100 GPU of 80 GB HBM2e. Evaluation datasets include CIFAR-100 for in distribution scenarios, CIFAR-10 for cross dataset OOD settings, and FGVC-Aircraft and Food-101 for unseen class OOD evaluations. We compare our method against four baseline models: ICon, SRL, LoRA + InfoNCE, and LoRA + Triplet.

Extended experimental details can be found in Appendix B.

4.2 In-Distribution Validation

Semantic retrieval quality. Figure 1 reports Label Precision@ K across four neighborhood sizes and three search budgets. EGA improves Label Precision@1 from ≈ 0.55 (raw CLIP) to > 0.70 at $nprobe = 1$, and the gap persists across all K and $nprobe$ values. The improvement is not an artifact of larger search budgets: even at $nprobe = 10$, EGA retains a clear advantage, indicating that the adapter genuinely tightens within-class neighborhoods rather than merely benefiting from broader index coverage.

Geometric structure. The Label Precision improvement has a direct geometric correlate. Figure 2 contrasts the L_2 distance distributions for true Top- K neighbors against random background points, before and after EGA. In the raw CLIP space, the two distributions overlap substantially: the distance to a true neighbor is often indistinguishable from the distance to an unrelated point, which is exactly the condition under which IVF-based ANN indexing degrades. After EGA, the two distributions separate cleanly across all K , with neighbors concentrated at small distances and background points pushed toward a distinct mode. This separation is the mechanism by which EGA improves the recall-efficiency frontier shown in Figure 3: the raw CLIP space achieves only 0.67 Recall@1




Figure 2: L_2 distance distributions for Top- K neighbors versus random background points on CIFAR-100. The raw CLIP space exhibits significant overlap; EGA separates the two distributions cleanly.




Figure 3: ANNS Recall@ K versus nprobe on CIFAR-100. EGA shifts the efficiency frontier upward across all K .

255 at nprobe = 1, whereas EGA reaches 0.80 under identical indexing, and saturates above 0.97 by
 256 nprobe = 5.

257 4.3 Out-of-Distribution Generalization

258 We now turn to the central empirical claim of this paper. Under strict OOD settings where evaluation
 259 categories are disjoint from training categories, the two retrieval metrics tell different stories about
 260 adapter quality, and EGA is the only method that performs well on both. Table 1 reports headline
 261 numbers at $K = 1$, nprobe = 1, and Figure 4 traces Label Precision across nprobe $\in \{1, 5, 10\}$
 262 to show that the pattern is robust to search budget. We evaluate on three benchmarks: CIFAR-10
 263 (trained on CIFAR-100), FGVC-Aircraft (80 seen / 20 unseen classes), and Food-101 (80 seen / 21
 264 unseen classes). For a fair comparison, all methods train identical lightweight adapters on top of
 265 frozen CLIP features and differ only in their loss functions.

266 **Two metrics, two stories.** On ANNS Recall, every adapter—including ICon and SRL—improves
 267 or matches the frozen CLIP baseline across all three OOD benchmarks (Table 1). Adapter training
 268 succeeds at *geometric* reorganization: the resulting embedding distribution is more amenable to
 269 IVF-based indexing. On Label Precision, however, the picture inverts. ICon and SRL fall *below*
 270 the frozen CLIP baseline on every dataset, with degradation reaching -32 to -36 percentage points
 271 on Food-101. The retrieved neighbors are geometrically closer in the IVF index, but they belong to
 272 the wrong semantic class. EGA is the only adapter that improves Label Precision on Aircraft (+3.6
 273 pp) and CIFAR-10 (+0.6 pp), and on Food-101 it retains the strongest Label Precision among all
 274 adapters with bounded -8.3 pp degradation.

275 **The pattern is structural, not incidental.** The Label Precision degradation of global contrastive
 276 methods persists across all three datasets and all three search budgets we tested (18 measurement
 277 points in total, Figure 4), ruling out the possibility that it is an artifact of a particular benchmark
 278 or indexing configuration. EGA’s relative advantage over global methods ranges from 4.5 pp on

Table 1: OOD performance at $K = 1$, $nprobe = 1$. Bold = exceeds CLIP; underline = below CLIP. ICon and SRL—trained with global contrastive losses on the full-capacity EGAMLP architecture—fall below the CLIP baseline in Label Precision on every dataset. LoRA, an order-of-magnitude smaller adapter, avoids catastrophic failure regardless of loss; we discuss this capacity-loss interaction in Section 4 and in Limitations.

Dataset	Metric	CLIP	ICon	SRL	LoRA+InfoNCE	LoRA+Triplet	EGA
CIFAR-10	AR@1	0.863	0.913	0.942	0.856	0.866	0.986
	LP@1	0.880	<u>0.841</u>	<u>0.794</u>	0.880	<u>0.870</u>	0.886
Aircraft	AR@1	0.773	0.875	0.863	0.911	0.881	0.898
	LP@1	0.512	<u>0.423</u>	<u>0.441</u>	0.560	0.577	0.548
Food-101	AR@1	0.842	<u>0.811</u>	<u>0.827</u>	0.902	0.899	0.886
	LP@1	0.881	<u>0.557</u>	<u>0.520</u>	<u>0.878</u>	<u>0.833</u>	<u>0.798</u>




Figure 4: Label Precision@1 across three OOD benchmarks and three search budgets.

279 CIFAR-10 to 24 pp on Food-101, and tracks the difficulty of the OOD shift: when seen and unseen
 280 classes are visually similar (as in Food-101’s fine-grained dishes), the dense gradient updates of ICon
 281 and SRL pull unseen samples most aggressively into the wrong seen-class clusters.

282 **Capacity matters too.** We additionally evaluate LoRA [24] as a parameter-efficient baseline
 283 (rank $r = 8$, $\sim 8K$ trainable parameters versus $\sim 4.2M$ for the EGAMLP architecture used by EGA,
 284 ICon, and SRL). Two configurations are included: LoRA + InfoNCE (a global contrastive loss) and
 285 LoRA + Triplet (the same loss as EGA). Both LoRA configurations avoid the catastrophic Label
 286 Precision collapse exhibited by ICon and SRL, even when paired with a global loss—LoRA + In-
 287 foNCE on Food-101 reaches LP 0.878, only 0.3 pp below the frozen baseline. This suggests that the
 288 catastrophic failure of ICon and SRL arises from the *combination* of (i) full-capacity adapter and
 289 (ii) dense global gradients, not from the global loss alone. With an aggressive bottleneck on adapter
 290 capacity, global losses become survivable. Our paper’s main mechanistic claim—that sparse local
 291 gradients preserve unseen-class manifolds—is therefore best read as a sufficient but not necessary
 292 condition for OOD-friendly adaptation. We return to this interaction in Limitations.

293 4.4 Mechanistic Analysis: Gradient Sparsity as OOD Protection

294 The superior OOD generalization of EGA raises a fundamental question: why does a triplet-trained
 295 adapter preserve unseen class structure while global contrastive objectives destroy it? We identify
 296 **gradient sparsity** as the key mechanism.

297 A triplet loss with margin m produces a non-zero gradient for a sample triplet (a, p, n) only when the
 298 margin constraint is violated, i.e., when

$$d(a, p) - d(a, n) + m > 0. \quad (11)$$

299 As the adapter converges on seen classes, the fraction of such *active* triplets decays rapidly. Figure 5
 300 (left) tracks this ratio throughout training: starting at 17.4% at epoch 10, it falls to 3.5% by conver-
 301 gence. At steady state, **96.5% of triplets contribute zero gradient**, meaning the adapter updates
 302 only a sparse subset of local neighborhood relationships and leaves the remainder of the embedding
 303 space—including regions occupied by unseen classes—structurally intact.

Table 2: Ablation study of EGA

Variant	Accuracy
Full EGA (Ours)	0.7015
w/o Residual connection	0.0065
w/o Zero-initialization	0.6840
w/o L2-normalization	0.6590

Table 3: Sensitivity of different hyperparameters

Triplet Margin m	0.1	0.2	0.3	0.5
Recall@1	0.8420	0.8180	0.8160	0.7400
Recall@5	0.9782	0.9654	0.9610	0.9125
Hidden Dimension d	512	1024	2048	4096
Recall@1	0.8300	0.8340	0.8200	0.8240
Recall@5	0.9695	0.9712	0.9680	0.9688

Global contrastive objectives such as ICon [2] and SRL [3] exhibit the opposite behavior. ICon minimizes a KL divergence over the full pairwise similarity matrix, and SRL jointly optimizes batch-wide uniformity and homogeneity. Both produce *dense* gradient fields in every training step: every sample pair contributes to the loss regardless of whether its local geometry already satisfies the retrieval objective. This global optimization pressure shapes the entire embedding distribution based on seen-class similarities, even for regions of the space occupied by unseen-class samples. The result is not loss of geometric structure—in fact, ICon and SRL produce highly compact clusters on unseen data—but rather loss of *semantically correct* structure: unseen samples are pulled toward the clusters of seen classes that they happen to resemble in CLIP’s original metric, rather than toward other unseen samples that share their true category.

Figure 5 (right) and Figure 4 together quantify the consequence: across all three OOD benchmarks, ICon and SRL fall below the frozen CLIP baseline on Label Precision (e.g., on Aircraft, ICon collapses to 0.423 from a baseline of 0.512), confirming that dense gradient updates actively harm zero-shot semantic retrieval. EGA, by contrast, improves on Aircraft and CIFAR-10 while modifying only 3.5% of the embedding geometry at convergence.




Figure 5: Gradient sparsity as OOD protection.

This analysis points to a key distinction in adapter design for foundation models: **when paired with sufficient capacity, dense global objectives reshape the embedding space according to seen-class similarities, projecting unseen samples onto the resulting structure; sparse local objectives leave the pretrained semantic relationships among unseen samples intact.** The contrast is dramatic on the high-capacity EGAMLP architecture: ICon and SRL fall well below the frozen CLIP baseline on Label Precision, while EGA preserves or improves it. As discussed in Section 4, restricting adapter capacity (e.g., via low-rank LoRA) is an alternative way to mitigate the collapse, suggesting that capacity and loss sparsity are complementary mechanisms for OOD robustness.

4.5 Ablation and Sensitivity Analysis

We perform ablation and hyperparameter sensitivity studies on CIFAR-100 to verify the effectiveness and stability of EGA. The results are summarized side-by-side in Table 2 and Table 3.

The residual connection is the most critical component—removing it causes accuracy to collapse from 0.7015 to 0.0065. L2 normalization and zero-initialization also contribute positively. EGA maintains high retrieval quality across a wide range of margin values (0.1–0.3 optimal) and hidden dimensions (512–4096), confirming that the geometric correction is structural in nature and does not rely on massive parameter counts or carefully tuned hyperparameters.

5 Conclusion

We have shown that high-capacity adapters trained with global contrastive objectives—a standard recipe for foundation-model adaptation—can systematically harm OOD retrieval performance, with Label Precision on unseen categories falling well below the frozen baseline. EGA addresses this failure mode by converging to a sparse-gradient regime in which 96.5% of training triplets contribute zero gradient at steady state, preserving the pretrained semantic structure for categories absent from

training. Across three OOD benchmarks (CIFAR-10, FGVC-Aircraft, Food-101), EGA improves Label Precision over the frozen baseline on the easier settings and bounds the degradation on the most challenging one, while ICon and SRL collapse by up to 36 percentage points. These results suggest that gradient sparsity at convergence is a useful design principle for retrieval adapters operating under distribution shift.

Limitations. Our study centers on visual representations from CLIP. Extending EGA to text and multimodal models, exploring unsupervised variants, and disentangling the joint role of adapter capacity and loss sparsity—preliminary results with a low-rank LoRA baseline suggest both factors contribute to OOD robustness—present avenues for future investigation.

References

- [1] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In Marina Meila and Tong Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event*, Proceedings of Machine Learning Research, pages 8748–8763. PMLR, 2021.
- [2] Shaden Naif Alshammari, John R Hershey, Axel Feldmann, William T Freeman, and Mark Hamilton. I-con: A unifying framework for representation learning. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- [3] Zijian Dong, Yilei Wu, Chongyao Chen, Yingtian Zou, Yichi Zhang, and Juan Helen Zhou. Improve representation for imbalanced regression through geometric constraints. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- [4] Chiara Lanza, Roberto Pereira, Marco Miozzo, Eduard Angelats, and Paolo Dini. Degradation of feature space in continual learning, 2026.
- [5] Xuqian Xue, Yiming Lei, Qi Cai, Hongming Shan, and Junping Zhang. PROTOCOL: Partial optimal transport-enhanced contrastive learning for imbalanced multi-view clustering. In *Forty-second International Conference on Machine Learning*, 2025.
- [6] Qingyan Zhao, Ruifang He, Jinpeng Zhang, Chang Liu, and Bo Wang. Representation degeneration problem in prompt-based models for natural language understanding. In Nicoletta Calzolari, Min-Yen Kan, Veronique Hoste, Alessandro Lenci, Sakriani Sakti, and Nianwen Xue, editors, *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 13946–13957, Torino, Italia, May 2024. ELRA and ICCL.
- [7] Tianyu Gao, Xingcheng Yao, and Danqi Chen. SimCSE: Simple contrastive learning of sentence embeddings. In Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih, editors, *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 6894–6910, Online and Punta Cana, Dominican Republic, November 2021. Association for Computational Linguistics.
- [8] Sameera Ramasinghe, Violetta Shevchenko, Gil Avraham, and Ajanthan Thalaisyasingam. Accept the modality gap: An exploration in the hyperbolic space. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 27253–27262, 2024.
- [9] Sedigheh Eslami and Gerard de Melo. Mitigate the gap: Improving cross-modal alignment in CLIP. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [10] Avik Pal, Max van Spengler, Guido Maria D’Amely di Melendugno, Alessandro Flaborea, Fabio Galasso, and Pascal Mettes. Compositional entailment learning for hyperbolic vision-language models. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [11] Daniele Savietto, Declan Campbell, André Panisson, Marco Nurisso, Giovanni Petri, Jonathan D. Cohen, and Alan Perotti. The geometry of representational failures in vision language models, 2026.

- [12] Baili Xiao, Zhibin Dong, Ke Liang, Suyuan Liu, Siwei Wang, Tianrui Liu, Xingchen Hu, En Zhu, and Xinwang Liu. Easemvc: Efficient dual selection mechanism for deep multi-view clustering. In *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 20716–20726, 2025.
- [13] Simone Ricci, Niccolò Biondi, Federico Pernici, Ioannis Patras, and Alberto Del Bimbo. λ -orthogonality regularization for compatible representation learning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2026.
- [14] Hangke Sui, Yuqing Wang, and Minh N. Do. Unicon: Unified framework for efficient contrastive alignment via kernels. In *The Fourteenth International Conference on Learning Representations*, 2026.
- [15] Panagiotis Koromilas, Efthymios Georgiou, Giorgos Bouritsas, Theodoros Giannakopoulos, Mihalalis A. Nicolaou, and Yannis Panagakis. A principled framework for multi-view contrastive learning, 2025.
- [16] Tongzhou Wang and Phillip Isola. Understanding contrastive representation learning through alignment and uniformity on the hypersphere. In *Proceedings of the 37th International Conference on Machine Learning, ICML’20*. JMLR.org, 2020.
- [17] Liang Wang, Xiang Tao, Qiang Liu, Shu Wu, and Liang Wang. Rethinking graph masked autoencoders through alignment and uniformity. In *Proceedings of the Thirty-Eighth AAAI Conference on Artificial Intelligence and Thirty-Sixth Conference on Innovative Applications of Artificial Intelligence and Fourteenth Symposium on Educational Advances in Artificial Intelligence, AAAI’24/IAAI’24/EAAI’24*. AAAI Press, 2024.
- [18] Xin Zhou, Yongjie Wang, and Zhiqi Shen. Cm³: Calibrating multimodal recommendation, 2025.
- [19] Mingqi Wu, Qiang Sun, and Archer Y. Yang. PCA++: How uniformity induces robustness to background noise in contrastive learning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2026.
- [20] Hyungbin Kim, Incheol Baek, and Yon Dohn Chung. Decoupled contrastive learning for federated learning, 2025.
- [21] Yingwen Zhang, Meng Wang, Xihua Sheng, Peilin Chen, Junru Li, Li Zhang, and Shiqi Wang. An information-theoretic regularizer for lossy neural image compression. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 15573–15582, October 2025.
- [22] Dong Jing, Xiaolong He, Yutian Luo, Nanyi Fei, Guoxing Yang, Wei Wei, Huiwen Zhao, and Zhiwu Lu. FineCLIP: Self-distilled region-based CLIP for better fine-grained understanding. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [23] Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for NLP. In Kamalika Chaudhuri and Ruslan Salakhutdinov, editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 2790–2799. PMLR, 09–15 Jun 2019.
- [24] Edward J Hu, yelong shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- [25] Fengming Lin. Vision language models: A survey of 26k papers, 2025.
- [26] Menglin Jia, Luming Tang, Bor-Chun Chen, Claire Cardie, Serge Belongie, Bharath Hariharan, and Ser-Nam Lim. Visual prompt tuning. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXIII*, page 709–727, Berlin, Heidelberg, 2022. Springer-Verlag.

- [27] Renrui Zhang, Wei Zhang, Rongyao Fang, Peng Gao, Kunchang Li, Jifeng Dai, Yu Qiao, and Hongsheng Li. Tip-adapter: Training-free adaption of clip for few-shot classification. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXV*, page 493–510, Berlin, Heidelberg, 2022. Springer-Verlag.
- [28] Shivam Rajendra Rai Sharma, Xiaoguang Zhu, Luca Cerny Oliveira, Kartik Patwari, La Rissa Vasquez, David Garcia, Louise Nicole C. Sevilla, Brittany N Dugger, and Chen-Nee Chuah. Benchmarking parameter efficient adaptation of vision language models on pathology. In *NeurIPS 2025 Workshop for Imageomics: Discovering Biological Knowledge from Images Using AI*, 2025.
- [29] Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image diffusion models. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 3813–3824, 2023.
- [30] Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with gpus. *IEEE Transactions on Big Data*, 7(3):535–547, 2021.
- [31] Yu A. Malkov and D. A. Yashunin. Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs. *IEEE Trans. Pattern Anal. Mach. Intell.*, 42(4):824–836, April 2020.
- [32] Suhas Jayaram Subramanya, Devvrit, Rohan Kadekodi, Ravishankar Krishaswamy, and Harsha Vardhan Simhadri. Diskann: fast accurate billion-point nearest neighbor search on a single node. In *Proceedings of the 33rd International Conference on Neural Information Processing Systems*, 2019.
- [33] Qi Chen, Bing Zhao, Haidong Wang, Mingqin Li, Chuanjie Liu, Zengzhong Li, Mao Yang, and Jingdong Wang. Spann: highly-efficient billion-scale approximate nearest neighbor search. In *Proceedings of the 35th International Conference on Neural Information Processing Systems, NIPS ’21*, Red Hook, NY, USA, 2021. Curran Associates Inc.
- [34] Ming Yang, Yuzheng Cai, and Weiguo Zheng. Cspg: Crossing sparse proximity graphs for approximate nearest neighbor search. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 103076–103100. Curran Associates, Inc., 2024.
- [35] Haya Diwan, Jinrui Gou, Cameron Musco, Christopher Musco, and Torsten Suel. Navigable graphs for high-dimensional nearest neighbor search: Constructions and limits. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 59513–59531. Curran Associates, Inc., 2024.
- [36] Hao Guo and Youyou Lu. Achieving low-latency graph-based vector search via aligning best-first search algorithm with ssd. In *19th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, pages 171–186, Boston, MA, 2025. USENIX Association.
- [37] Hao Guo and Youyou Lu. Odinann: Direct insert for consistently stable performance in billion-scale graph-based vector search. In *24th USENIX Conference on File and Storage Technologies (FAST)*, Santa Clara, CA, 2026. USENIX Association.
- [38] Elias Jääsaari, Ville Hyvönen, and Teemu Roos. LoRANN: Low-rank matrix factorization for approximate nearest neighbor search. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [39] Piotr Indyk and Haike Xu. Worst-case performance of popular approximate nearest neighbor search implementations: Guarantees and limitations. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 66239–66256. Curran Associates, Inc., 2023.
- [40] Jie Hu, Li Shen, Samuel Albanie, Gang Sun, and Enhua Wu. Squeeze-and-excitation networks. *IEEE Trans. Pattern Anal. Mach. Intell.*, 42(8):2011–2023, August 2020.

494 [41] Kate Keahey, Jason Anderson, Zhuo Zhen, Pierre Riteau, Paul Ruth, Dan Stanzione, Mert
495 Cevik, Jacob Collieran, Haryadi S. Gunawi, Cody Hammock, Joe Mambretti, Alexander Barnes,
496 François Halbach, Alex Rocha, and Joe Stubbs. Lessons learned from the chameleon testbed. In
497 *Proceedings of the 2020 USENIX Annual Technical Conference (USENIX ATC '20)*. USENIX
498 Association, July 2020.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The paper presents both theoretical and experimental evidence to support its claims.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: We discuss the limitations of our approach in the conclusion.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Full proofs are provided in the appendix, and all assumptions are clearly stated in the main paper and/or the appendix.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All code and data needed to reproduce the main experimental results are available on the anonymous Github repository.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All the details to reproduce the main experimental results are provided in the anonymous Github repository, including instructions for data access and preparation, and scripts to reproduce all experimental results for the new proposed method and baselines.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: All such details are provided in the anonymous Github repository, and the main paper includes a summary of the experimental setting.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The results are pretty stable so we did not report error bars; we explain this in the main paper.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: These details are provided in the anonymous Github repository, and the main paper includes a summary of the compute resources used for the experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: We believe we have followed the NeurIPS Code of Ethics in all aspects of our research.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.

- 809 • We recognize that the procedures for this may vary significantly between institutions
810 and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the
811 guidelines for their institution.
812 • For initial submissions, do not include any information that would break anonymity (if
813 applicable), such as the institution conducting the review.

814 **16. Declaration of LLM usage**

815 Question: Does the paper describe the usage of LLMs if it is an important, original, or
816 non-standard component of the core methods in this research? Note that if the LLM is used
817 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
818 scientific rigor, or originality of the research, declaration is not required.

819 Answer: [N/A]

820 Justification: [N/A]

821 Guidelines:

- 822 • The answer [N/A] means that the core method development in this research does not
823 involve LLMs as any important, original, or non-standard components.
824 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not
825 be described.

A Proof of Proposition 1

Proof. We prove each claim in turn.

(1) Gradient support. The triplet loss in Eq. (7) can be written as $\mathcal{L} = \frac{1}{|\mathcal{B}|} \sum_{(a,p,n) \in \mathcal{B}} \ell_{apn}$, where $\ell_{apn} = \max(0, \delta_{apn} + m)$ and $\delta_{apn} = d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n)$. The subgradient with respect to θ is

$$\nabla_{\theta} \ell_{apn} = \mathbf{1}[\delta_{apn} + m > 0] \cdot \nabla_{\theta} \delta_{apn}, \quad (12)$$

which is exactly zero whenever $\delta_{apn} + m \leq 0$. Therefore, $\nabla_{\theta} \mathcal{L}$ receives contributions only from active triplets satisfying Eq. (8).

(2) Convergence of $\rho(t)$. Under mild regularity conditions (Lipschitz continuity of f_{θ} and bounded embedding norms), the sequence $\{\hat{\mathbf{z}}^{(t)}\}$ converges to a stationary point $\hat{\mathbf{z}}^*$ of $\mathcal{L}$. At stationarity, $\rho(t)$ converges to ρ^* as defined in Eq. (10), since $\rho(t)$ is a continuous functional of the embedding distribution and the embeddings converge.

(3) Zero gradient over unseen regions. Let $\mathcal{R}_{\text{unseen}} \subset \mathbb{S}^{d-1}$ denote the region of the embedding space occupied by unseen-class images under f_{θ} . By construction, all triplets (a, p, n) in $\mathcal{D}_{\text{train}}$ satisfy $y_a, y_p, y_n \in \mathcal{Y}_{\text{seen}}$, so $\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p, \hat{\mathbf{z}}_n \notin \mathcal{R}_{\text{unseen}}$ for all training triplets. The gradient $\nabla_{\theta} \delta_{apn}$ depends on f_{θ} only through its evaluation at $\mathbf{z}_a, \mathbf{z}_p, \mathbf{z}_n$, which lie outside $\mathcal{R}_{\text{unseen}}$. Consequently, no training step produces a gradient that directly modifies f_{θ} in $\mathcal{R}_{\text{unseen}}$, and the CLIP geometry in that region is preserved throughout training.

(4) Dependence on m . The residual distance gap δ_{apn}^* at convergence is bounded by the within-class diameter $\Delta_c = \max_{x, x' \in c} d(\hat{\mathbf{z}}, \hat{\mathbf{z}}')$ for seen class c . For $m \leq \min_c \Delta_c$, the adapter can drive all within-class distances below the threshold, pushing ρ^* toward zero. For $m > \min_c \Delta_c$, some triplets remain active at convergence because the within-class spread exceeds the margin, resulting in a strictly positive ρ^* . This explains the empirically observed increase in ρ^* with larger m reported in Table 3. $\square$

B More Experimental Details

Datasets. We evaluate on four datasets covering both in-distribution and OOD retrieval scenarios. **CIFAR-100** (100 classes, 60K images) serves as the in-distribution benchmark: the adapter is trained and evaluated on the same class label space, with a 75/25 split between database and query sets. **CIFAR-10** (10 classes, 60K images) provides a cross-dataset OOD setting: the adapter is trained on CIFAR-100 and evaluated on the disjoint CIFAR-10 classes. **FGVC-Aircraft** (100 fine-grained aircraft variants, 10K images) and **Food-101** (101 food categories, 25K images for the test split we use) provide unseen-class OOD settings: classes are randomly partitioned into 80% seen / 20% unseen with a fixed seed, the adapter is trained on the seen split, and retrieval is evaluated strictly on the unseen split. Image embeddings for all datasets are extracted with frozen CLIP ViT-B/32.

Baselines. We compare EGA against the frozen CLIP baseline and four trained adapters, all sharing the same training data, evaluation protocol, and 75/25 retrieval split:

- **Icon** [2]: a global contrastive objective minimizing the KL divergence between the empirical similarity distribution and the class-membership distribution. We instantiate Icon on the same EGAMLP architecture used by our method (~ 4.2 M parameters).
- **SRL** [3]: a structural regularization objective jointly optimizing batch-wide uniformity and within-class homogeneity, also instantiated on the EGAMLP architecture.
- **LoRA + InfoNCE**: a low-rank adapter (rank $r = 8$, $\alpha = 16$, ~ 8 K parameters) with zero-initialized residual structure, trained with supervised InfoNCE—a standard global contrastive loss.
- **LoRA + Triplet**: identical low-rank architecture as above, but trained with the same small-margin ($m = 0.2$) triplet loss as EGA.

The two LoRA configurations isolate the effect of architectural capacity vs. loss function: comparing LoRA + InfoNCE against Icon/SRL holds the loss family roughly fixed while changing capacity, and comparing LoRA + Triplet against EGA holds the loss fixed while changing capacity.

873 **Indexing and metrics.** For all retrieval evaluations we use FAISS [30] IndexIVFFlat with
874 nlist=10 centroids and report results at $nprobe \in \{1, 5, 10\}$. We measure both Label Preci-
875 sion (LP@ K , Eq. 1) and ANNS Recall (AR@ K , Eq. 2) at $K \in \{1, 3, 5, 10\}$.

876 **Training.** All adapters are trained for 150 epochs with AdamW (weight decay 10^{-4}), cosine-
877 annealed learning rate, and batch size 256. EGAMLP-based methods (EGA, ICon, SRL) use learning
878 rate 10^{-4} ; LoRA variants use 10^{-3} to compensate for their smaller parameter count. Triplet-based
879 methods sample one positive and one negative per anchor uniformly within each mini-batch; global
880 contrastive methods (ICoN, SRL, LoRA + InfoNCE) operate on the full pairwise similarity matrix.
881 All experiments use a fixed random seed (42) for reproducibility.

882 **Compute platform.** All experiments were conducted on the Chameleon Cloud [41] platform using
883 a compute node equipped with 256 AMD EPYC-7763 CPU cores, 512 GiB of RAM, and an NVIDIA
884 A100 GPU of 80 GB HBM2e. The operating system is Ubuntu 24.04 LTS.