

---

# EGA: Adapting Frozen Encoders for Vector Search with Bounded OOD Degradation

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

1 Vector similarity search systems built on frozen vision-language encoders routinely  
2 face queries from classes absent at adapter training time, since both the indexed  
3 database and the live query stream contain categories that were never labeled. We  
4 study how lightweight adapters behave under this distribution shift, and find a  
5 consistent in-distribution/out-of-distribution (ID/OOD) tradeoff among existing  
6 training paradigms. Global contrastive objectives such as ICon and SRL, when  
7 paired with high-capacity adapters, achieve near-perfect in-distribution Label Pre-  
8 cision on CIFAR-100 (0.999 and 0.992 at K=1, nprobe=1) but collapse to 0.562  
9 and 0.552 on the unseen classes of Food-101, a drop of 44 percentage points.  
10 Low-capacity LoRA variants avoid this collapse but underperform substantially in-  
11 distribution. We propose Euclidean Geodesic Alignment (EGA), a residual adapter  
12 trained with a small-margin triplet loss on the class-aware neighborhood graph of  
13 pre-trained embeddings. EGA induces gradient sparsity at convergence (96.5%  
14 of triplets produce zero gradient), empirically leaving the unseen-class geometry  
15 largely intact while permitting high-capacity in-distribution refinement. Across  
16 CIFAR-100 (ID) and three OOD benchmarks (CIFAR-10, FGVC-Aircraft, Food-  
17 101), EGA reaches the highest in-distribution Label Precision among adapters that  
18 avoid OOD collapse (0.705 versus 0.6176 for capacity-matched LoRA at r=128),  
19 with bounded OOD degradation (-7 to -9 pp, versus -32 to -33 pp for ICon and  
20 SRL). The benefits transfer to stronger backbones: on DINOv2-large and SigLIP,  
21 EGA improves Label Precision@1 by 4.75 and 8.55 percentage points.

## 22 1 Introduction

23 A vector similarity search system deployed in production rarely operates on a fixed label space. A  
24 retailer indexes product images today and adds a new category line next month; a content platform  
25 builds an index over existing tags and continually ingests media from emerging topics; a scientific  
26 image database is annotated for known phenotypes and queried for ones discovered later. In such  
27 scenarios, the index and the query stream together cover a label space strictly larger than what was  
28 available when any model was trained or adapted.

29 This deployment reality sits awkwardly with how adapters over frozen vision-language encoders are  
30 currently trained. The dominant paradigm takes a labeled subset, optimizes the embedding geometry  
31 to separate those classes, and ships the resulting adapter as a drop-in replacement for the frozen  
32 encoder’s output. The implicit assumption is that geometry tightened for seen classes will, at worst,  
33 leave unseen-class geometry unchanged. We show this assumption fails systematically: high-capacity  
34 adapters trained with global contrastive objectives reshape the entire embedding distribution, pulling  
35 unseen-class samples toward whichever seen-class clusters they happen to resemble in the pre-trained  
36 metric. The retrieval system inherits this distortion silently, since approximate nearest neighbor

indices remain well-behaved geometrically while the neighbors they return belong to the wrong semantic class.

Two design choices in current adapter training are responsible. First, the loss function: global contrastive objectives such as ICon [1] and SRL [4] compute gradients over the full pairwise similarity matrix in every step, so seen-class supervision exerts pressure across the entire embedding distribution rather than only on labeled regions. Second, the adapter capacity: a sufficiently expressive adapter has the representational room to comply with that pressure, and zero-initialized residual designs do not by themselves prevent the cumulative perturbation from reaching unseen-class regions. Existing alternatives address one factor at a time. Low-rank adapters such as LoRA [10] cap the second factor, which avoids the worst out-of-distribution (OOD) failures but limits in-distribution (ID) refinement, while keeping the global loss leaves the first factor untouched. We argue that adapters intended to operate under unseen-class queries should be designed against both factors at once.

The cost of failing on either factor is large. On CIFAR-100, high-capacity global contrastive adapters achieve near-perfect Label Precision: ICon reaches 0.999 and SRL reaches 0.992 at  $K = 1$ ,  $n_{\text{probe}} = 1$ . On the strictly unseen classes of Food-101, the same adapters collapse to 0.562 and 0.552, which is 44 percentage points below their in-distribution scores and between 32 and 33 percentage points below the frozen CLIP baseline. LoRA variants avoid this collapse but pay for it in-distribution, with Label Precision bounded between 0.61 and 0.67, well below what a high-capacity adapter can achieve when the loss does not over-extend into unseen-class regions.

We propose Euclidean Geodesic Alignment (EGA), a residual adapter designed against both factors simultaneously. EGA pairs a high-capacity zero-initialized residual architecture with a small-margin triplet loss on the class-aware neighborhood graph of pre-trained embeddings, with  $\ell_2$  projection onto the unit hypersphere. The triplet loss induces a self-limiting training dynamic: as the adapter converges on seen classes, the fraction of margin-violating triplets decays rapidly, reaching 96.5% zero-gradient triplets at steady state. The capacity to refine seen-class neighborhoods is preserved, but the loss stops applying pressure once seen-class neighborhoods are locally separated, so the global reshaping that destroys unseen-class geometry under ICon and SRL never accumulates.

We evaluate EGA on CIFAR-100 (ID) and three OOD benchmarks (CIFAR-10, FGVC-Aircraft, Food-101), comparing against the frozen CLIP baseline, global contrastive adapters (Icon, SRL), and LoRA configurations with ranks  $r = 64$  and  $r = 128$  that are capacity-matched to EGA. The six methods occupy distinct operating points in the (ID, OOD) Label Precision plane. EGA reaches the highest in-distribution Label Precision among adapters that avoid OOD collapse, achieving 0.705 on CIFAR-100 versus 0.6176 for capacity-matched LoRA at  $r = 128$ , while bounding OOD degradation between  $-7$  and  $-9$  percentage points on CIFAR-10 and Food-101 (compared to between  $-32$  and  $-33$  for ICon and SRL) and improving by 9.9 percentage points on FGVC-Aircraft. The benefits transfer to stronger backbones, with EGA improving Label Precision@1 by 4.75 and 8.55 percentage points on DINOv2-large and SigLIP. EGA does not dominate every method on every axis: on CIFAR-10 and Food-101, capacity-matched LoRA achieves higher OOD Label Precision, and deployments dominated by OOD queries on these distributions may prefer LoRA at the cost of weaker in-distribution refinement.

In summary, this work makes the following contributions:

- We frame adapter tuning over frozen vision-language encoders from a vector search deployment perspective, in which the index and query stream together cover a label space larger than the labeled subset available for adapter training, and we characterize the failure modes of existing adapter paradigms under this lens.
- We propose Euclidean Geodesic Alignment (EGA), a residual adapter that combines a high-capacity zero-initialized architecture with a small-margin triplet objective. The triplet loss converges to a sparse-gradient regime (96.5% zero-gradient triplets at steady state), which empirically preserves unseen-class geometry while permitting in-distribution refinement.
- Across CIFAR-100 and three OOD benchmarks, and across three backbones (CLIP ViT-B/32, DINOv2-large, SigLIP), EGA occupies a distinct operating point in the (ID, OOD) Label Precision plane, achieving the highest in-distribution Label Precision among adapters that avoid OOD collapse, while capacity-matched LoRA variants preserve OOD precision better on two of three benchmarks at the cost of 8.7 percentage points in-distribution.

## 2 Background and Related Work

Foundation models such as CLIP [25] produce embedding spaces that encode both seen-class semantics and broader priors over unseen visual concepts. Related work has documented degradation phenomena in adapted representations [20, 37], and recent work documents analogous degradation in contrastive representations, including anisotropic collapse [44, 6], modality gap fracturing [26, 5], and loss of compositional structure [24, 28, 36]. Recent strategies counter these issues with geometric constraints such as orthogonality regularization [27]. *EGA exploits the gradient sparsity of small-margin triplet losses to keep adapter updates concentrated on seen-class neighborhoods, empirically leaving unseen-class regions of the pretrained geometry largely intact.*

Recent representation learning frequently relies on global contrastive objectives to optimize the latent space [31, 19]. A common paradigm jointly enforces alignment for positive pairs and uniformity across the feature distribution [34], applied in graph encoders [33] and multimodal recommendation [47]. While uniformity improves robustness to noise [35], it treats all latent regions equally. Other structural regularizers include unifying frameworks for multi-view alignment [1], geometric constraints for imbalanced regression [4], decoupled objectives [18], information-theoretic density control [43], and region-based distillation [15]. *EGA replaces dense global optimization with sparse local triplet updates, which empirically reduces unseen-class distortion.*

Parameter-efficient adaptation is an approach for transferring large foundation models without full fine-tuning [9, 10], which dominates vision-language research [21]. Lightweight methods, e.g., Visual Prompt Tuning [14] and Tip Adapter [41], specialize the model without modifying the frozen backbone and rival full fine-tuning on many tasks [29], while zero-initialization strategies further constrain the magnitude of pretrained-manifold perturbation [40]. *EGA admits a zero-initialized residual adapter, localized triplet objective, and parameter efficiency with gradient-sparse updates.*

Vector search is the foundational engine for scalable high-dimensional retrieval. The systems community has developed indexing algorithms for billion-scale datasets, including FAISS [16], HNSW [22], DiskANN [30], and SPANN [2]. Recent work refines graph topologies [38, 3], optimizes SSD-based search [7, 8], and accelerates distance computations [13], yet popular ANN implementations still suffer worst-case limitations [12]. *EGA recalibrates the metric geometry before index construction, rendering search backends achieve higher semantic recall without algorithmic changes.*

## 3 Methodology

### 3.1 Problem Formulation

Let  $\phi : \mathcal{X} \rightarrow \mathbb{R}^d$  denote a frozen vision-language encoder (e.g., CLIP ViT-B/32) that maps an image  $x \in \mathcal{X}$  to a unit-normalized embedding  $\mathbf{z} = \phi(x) \in \mathbb{S}^{d-1}$ . We are given a set of labeled training images  $\mathcal{D}_{\text{train}} = \{(x_i, y_i)\}_{i=1}^N$ , where  $y_i \in \mathcal{Y}_{\text{seen}}$  belongs to a set of *seen* classes. At inference time, the retrieval system must operate over a database that also contains images from a disjoint set of *unseen* classes  $\mathcal{Y}_{\text{unseen}}$ , with  $\mathcal{Y}_{\text{seen}} \cap \mathcal{Y}_{\text{unseen}} = \emptyset$ . Our goal is to learn an adapter  $f_\theta : \mathbb{S}^{d-1} \rightarrow \mathbb{S}^{d-1}$  such that the transformed embeddings  $\hat{\mathbf{z}} = f_\theta(\mathbf{z})$  support more accurate approximate nearest neighbor (ANN) retrieval on *both* seen and unseen classes, without access to unseen-class labels during training.

We evaluate retrieval quality along two complementary dimensions. *Label Precision@K* (LP@K) measures the semantic quality of retrieved results: whether the  $K$  nearest neighbors in the transformed space share the same class label as the query:

$$\text{LP@K} = \frac{1}{|\mathcal{Q}|} \sum_{q \in \mathcal{Q}} \frac{1}{K} \sum_{k=1}^K \mathbf{1}[y_{\hat{n}_k(q)} = y_q], \quad (1)$$

where  $\mathcal{Q}$  is the query set and  $\hat{n}_k(q)$  denotes the  $k$ -th nearest neighbor of  $q$  in the transformed embedding space. *ANNS Recall@K* (AR@K) measures the geometric fidelity of the ANN index, i.e., the fraction of true  $K$  nearest neighbors (under exact search) that are recovered by the approximate index at a given search budget nprobe:

$$\text{AR@K} = \frac{1}{|\mathcal{Q}|} \sum_{q \in \mathcal{Q}} \frac{|\mathcal{N}_K^{\text{exact}}(q) \cap \mathcal{N}_K^{\text{approx}}(q)|}{K}, \quad (2)$$

where  $\mathcal{N}_K^{\text{exact}}(q)$  and  $\mathcal{N}_K^{\text{approx}}(q)$  denote the exact and approximate top- $K$  neighbor sets, respectively. The two metrics capture distinct failure modes: an adapter may improve LP@ $K$  by tightening within-class clusters (semantic improvement) while simultaneously improving AR@ $K$  by making the geometry more amenable to IVF-based indexing (geometric improvement).

### 3.2 Manifold-Preserving Adapter Design

EGA introduces a lightweight adapter  $f_\theta$  composed of stacked residual blocks with a final refinement layer. The architecture is constrained by three design principles, each of which contributes to keeping the adapter’s updates anchored to the pretrained geometry.

**(1) Residual connection with zero initialization** Each EGA block applies a residual update as:

$$\mathbf{z}' = \mathbf{z} + g_\theta(\mathbf{z}), \quad (3)$$

where  $g_\theta$  is a learned transformation initialized so that  $g_\theta(\mathbf{z}) \approx \mathbf{0}$  at the start of training. Concretely, the final linear layer of  $g_\theta$  is initialized to zero, making  $f_\theta$  an identity mapping at initialization. This guarantees that training begins from the pretrained geometry rather than a random perturbation, and that the adapter can only *refine* rather than replace the CLIP embedding. As confirmed by our ablation (Table 4), removing the residual connection causes accuracy to collapse from 0.70 to 0.01, because the lightweight MLP cannot reconstruct complex semantic relationships from seen-class data alone.

**(2) Feature gating** Inside each residual block, the transformed signal passes through a module:

$$\text{Gate}(\mathbf{h}) = \mathbf{h} \odot \sigma(W_2 \text{GELU}(W_1 \mathbf{h})), \quad (4)$$

where  $W_1 \in \mathbb{R}^{(d/4) \times d}$ ,  $W_2 \in \mathbb{R}^{d \times (d/4)}$ ,  $\sigma$  is the Sigmoid function, and  $\odot$  denotes elementwise multiplication. This is analogous to Squeeze-and-Excitation networks [11] to learn a per-dimension weight that amplifies retrieval-relevant dimensions and suppresses noisy ones. The bottleneck ratio 1/4 is chosen to prevent the gating from overfitting on seen-class data, as validated in Table 5.

The output of  $f_\theta$  is projected back onto the unit hypersphere  $\mathbb{S}^{d-1}$  via  $\ell_2$ :

$$\hat{\mathbf{z}} = \frac{f_\theta(\mathbf{z})}{\|f_\theta(\mathbf{z})\|_2}. \quad (5)$$

This constraint ensures that the transformed embeddings remain in the same metric space as the original CLIP embeddings, so that Euclidean distance and cosine similarity remain interchangeable on the hypersphere. Removing this constraint degrades retrieval performance by 4.2 percentage points (Table 4), confirming that hypersphere consistency is necessary for stable ANN indexing.

**(3) Full forward pass** Let  $\text{Block}_i$  denote the  $i$ -th EGA block consisting of a two-layer MLP with LayerNorm followed by feature gating, and let  $\text{Refine}$  denote the final linear projection with LayerNorm. The complete forward pass for an input embedding  $\mathbf{z} \in \mathbb{S}^{d-1}$  is:

$$\hat{\mathbf{z}} = \ell_2(\text{Refine}(\text{Block}_2(\text{Block}_1(\mathbf{z})))) . \quad (6)$$

The total number of trainable parameters is approximately  $4.2M$  for  $d = 512$  and hidden dimension 2048, which is negligible compared to the frozen CLIP backbone ( $\sim 150M$  parameters).

### 3.3 Local Triplet Training

Given a mini-batch  $\mathcal{B}$  sampled from  $\mathcal{D}_{\text{train}}$ , we construct triplets  $(a, p, n)$  where  $a$  is an anchor image,  $p$  is a positive sampled uniformly from the same class as  $a$ , and  $n$  is a negative sampled uniformly from a different class. The adapter is trained with the standard triplet margin loss:

$$\mathcal{L} = \frac{1}{|\mathcal{B}|} \sum_{(a,p,n) \in \mathcal{B}} \max(0, d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m), \quad (7)$$

where  $d(\cdot, \cdot)$  denotes distance,  $m > 0$  is the margin, and  $\hat{\mathbf{z}} = f_\theta(\mathbf{z})$  is the transformed embedding.

A triplet  $(a, p, n)$  contributes a non-zero gradient to  $\theta$  if and only if the margin constraint is *violated*:

$$d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m > 0. \quad (8)$$

We refer to such triplets as *active*. As the adapter converges on seen classes, increasingly many triplets satisfy the margin constraint and become inactive, contributing zero gradient. This sparsity has a direct consequence for out-of-distribution generalization. Because unseen classes are absent from  $\mathcal{D}_{\text{train}}$ , no triplet involving an unseen-class embedding is ever constructed. Provided that the adapter’s updates remain effectively local (i.e., the cumulative perturbation induced by seen-class training gradients remains small in unseen-class regions of the embedding space), the manifold structure for unseen classes is largely preserved. The gradient sparsity mechanism enforces this locality: once seen-class neighborhoods satisfy the margin, the corresponding gradients vanish, and only the small fraction of actively-violated local relationships continue to be updated.

We formalize these intuitions in the following observation, which characterizes the steady-state behavior of triplet-trained adapters and its implications for manifold preservation.

**Gradient sparsity at convergence.** The triplet loss in Eq. (7) consists of hinge terms whose subgradient with respect to  $\theta$  is exactly zero whenever the margin constraint is satisfied:

$$\nabla_{\theta} \ell_{apn} = \mathbf{1}[d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n) + m > 0] \cdot \nabla_{\theta} \delta_{apn}, \quad (9)$$

where  $\delta_{apn} = d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n)$ . Define the *active triplet ratio* at training step  $t$  as

$$\rho(t) = \Pr_{(a,p,n) \sim \mathcal{B}} \left[ d(\hat{\mathbf{z}}_a^{(t)}, \hat{\mathbf{z}}_p^{(t)}) - d(\hat{\mathbf{z}}_a^{(t)}, \hat{\mathbf{z}}_n^{(t)}) + m > 0 \right]. \quad (10)$$

Under the standard assumption that training converges to a stationary point of  $\mathcal{L}$ ,  $\rho(t)$  converges to a steady-state value  $\rho^* \in [0, 1]$  determined by the margin  $m$  and the within-class distance distribution of the converged embeddings. We measure  $\rho^* \approx 3.5\%$  at convergence on FGVC-Aircraft (Section 4.3), meaning that 96.5% of training triplets contribute zero gradient at steady state.

**From sparsity to OOD locality: an empirical claim.** Strictly speaking, sparse seen-class gradients do not formally imply that  $f_{\theta}$  leaves unseen-class regions unchanged: the adapter is a parametric MLP, not a locally supported transformation, so any update to  $\theta$  propagates globally. We therefore make the weaker, empirical claim that the residual zero-initialized architecture combined with gradient sparsity keeps the cumulative perturbation in unseen-class regions small enough to preserve their pretrained semantic structure for retrieval. The OOD measurements in Section 4.2 are the evidence: EGA preserves or improves Label Precision on three unseen-class benchmarks, while ICon and SRL lack the gradient sparsity and fall below the frozen baseline. A formal Lipschitz-based bound on the unseen-class perturbation is left for future work; we discuss technical conditions in Appendix A.

## 4 Evaluation

Experiments were conducted on a server with 256 AMD EPYC-7763 CPU cores, 512 GiB RAM, and an NVIDIA A100 GPU of 80 GB HBM2e. Evaluation datasets include CIFAR-100 for in-distribution scenarios, CIFAR-10 for cross-dataset OOD settings, and FGVC-Aircraft and Food-101 for unseen-class OOD evaluations. All of these datasets can be directly downloaded from the Pytorch library. The backbone models include CLIP ViT-B/32 [25], DINOv2-large [23], and SigLIP [39]. We compare our method against four adapter models: ICon [1], SRL [4], LoRA + InfoNCE [10, 32], and LoRA + Triplet [10]; We compare our method against four adapter models: ICon [1], SRL [4], LoRA + InfoNCE [10, 32], and LoRA + Triplet [10]. We omit Linear Probing, Tip-Adapter [42], and prompt tuning methods (CoOp/CoCoOp [46, 45]) because they are not compatible with our cross-dataset OOD setting. Extended experimental details can be found in Appendix B.

### 4.1 In-Distribution Validation

Figure 1 reports Label Precision@ $K$  across four neighborhood sizes and three search budgets. EGA raises Label Precision@1 from 0.549 (raw CLIP) to 0.705 at  $n_{\text{probe}} = 1$ , and the gap persists across all  $K$  and  $n_{\text{probe}}$  values, including  $n_{\text{probe}} = 10$ , indicating that the adapter tightens within-class neighborhoods rather than benefiting from broader index coverage.

The Label Precision improvement has a direct geometric correlate. Figure 2 contrasts the  $L_2$  distance distributions for true Top- $K$  neighbors against random background points. In the raw CLIP space, the two distributions overlap substantially, which is the condition under which IVF-based ANN indexing




Figure 1: Label Precision@ $K$  on CIFAR-100 across  $K \in \{1, 3, 5, 10\}$  and  $nprobe \in \{1, 5, 10\}$ .




Figure 2: Distance distributions for Top- $K$  neighbors vs. random background points on CIFAR-100.

Table 1: In-distribution retrieval performance on CIFAR-100 at  $K = 1$ ,  $nprobe = 1$ . Standard errors are below 0.01 and omitted for clarity.

| Metric | CLIP  | Icon         | SRL          | LoRA+InfoNCE | LoRA+Triplet | EGA          |
|--------|-------|--------------|--------------|--------------|--------------|--------------|
| LP@1   | 0.549 | <b>0.999</b> | <b>0.992</b> | <b>0.668</b> | <b>0.612</b> | <b>0.705</b> |
| AR@1   | 0.667 | <b>0.992</b> | <b>0.991</b> | <b>0.879</b> | <b>0.908</b> | <b>0.825</b> |

degrades. After EGA, the two distributions separate cleanly across all  $K$ , and this separation drives the recall-efficiency improvement in Figure 3: raw CLIP achieves 0.667 Recall@1 at  $nprobe = 1$ , whereas EGA reaches 0.825 and saturates above 0.99 by  $nprobe = 5$ .

Table 1 reports all six methods’ in-distribution performance, since the contrast with their out-of-distribution behavior in Section 4.2 is central to our analysis. Icon and SRL reach near-perfect Label Precision (0.999 and 0.992) by maximally tightening seen-class clusters under their global contrastive objectives. EGA reaches 0.705, an improvement of +16 pp over the frozen baseline, and the LoRA variants land between EGA and the baseline. As Section 4.2 will show, the same property that lets Icon and SRL dominate this metric is what causes their catastrophic out-of-distribution collapse.

## 4.2 Out-of-Distribution Generalization

Under strict OOD settings where evaluation categories are disjoint from training categories, the two retrieval metrics tell different stories. Table 2 reports headline numbers at  $K = 1$ ,  $nprobe = 1$ . We evaluate on three benchmarks: CIFAR-10 (trained on CIFAR-100), FGVC-Aircraft (80 seen / 20 unseen classes), and Food-101 (80 seen / 21 unseen classes). For a fair comparison, all methods train identical lightweight adapters on top of frozen CLIP features and differ only in their loss functions. Robustness across search budgets ( $nprobe \in \{1, 5, 10\}$ ) is reported in Appendix C.

On ANNS Recall, the adapters either match or exceed the frozen CLIP baseline on Aircraft and Food-101, indicating that adapter training succeeds at geometric reorganization. On Label Precision, the picture inverts. Icon and SRL fall *below* the frozen baseline on every dataset, with degradation




Figure 3: ANNS Recall@ $K$  versus nprobe on CIFAR-100.

Table 2: OOD performance at  $K = 1$ , nprobe = 1, mean  $\pm$  std over 3 random seeds (42, 123, 456).

| Dataset  | Metric | CLIP  | Icon                    | SRL                     | LoRA+InfoNCE            | LoRA+Triplet            | EGA                     |
|----------|--------|-------|-------------------------|-------------------------|-------------------------|-------------------------|-------------------------|
| CIFAR-10 | AR@1   | 0.863 | <u>0.797</u> $\pm$ .003 | <u>0.819</u> $\pm$ .001 | <b>0.869</b> $\pm$ .012 | 0.861 $\pm$ .011        | <u>0.833</u> $\pm$ .005 |
|          | LP@1   | 0.880 | <u>0.560</u> $\pm$ .010 | <u>0.536</u> $\pm$ .008 | <b>0.885</b> $\pm$ .006 | 0.875 $\pm$ .006        | 0.810 $\pm$ .002        |
| Aircraft | AR@1   | 0.773 | <b>0.853</b> $\pm$ .016 | <b>0.879</b> $\pm$ .023 | <b>0.891</b> $\pm$ .015 | <b>0.893</b> $\pm$ .008 | <b>0.905</b> $\pm$ .015 |
|          | LP@1   | 0.512 | <u>0.470</u> $\pm$ .034 | <u>0.375</u> $\pm$ .032 | <b>0.538</b> $\pm$ .020 | <b>0.569</b> $\pm$ .011 | <b>0.611</b> $\pm$ .020 |
| Food-101 | AR@1   | 0.842 | <u>0.821</u> $\pm$ .001 | <u>0.824</u> $\pm$ .015 | <b>0.906</b> $\pm$ .003 | <b>0.893</b> $\pm$ .008 | <b>0.879</b> $\pm$ .011 |
|          | LP@1   | 0.881 | <u>0.562</u> $\pm$ .011 | <u>0.552</u> $\pm$ .010 | <u>0.883</u> $\pm$ .004 | <u>0.833</u> $\pm$ .007 | <u>0.791</u> $\pm$ .002 |

Table 3: Performance of EGA versus capacity-matched LoRA ( $K = 1$ , nprobe=1).

| Dataset        | Metric | EGA          | InfoNCE<br>r=64 | InfoNCE<br>r=128 | Triplet<br>r=64 | Triplet<br>r=128 |
|----------------|--------|--------------|-----------------|------------------|-----------------|------------------|
| CIFAR-100 (ID) | LP@1   | <b>0.705</b> | 0.6256          | 0.6176           | 0.5568          | 0.5472           |
|                | AR@1   | <b>0.825</b> | 0.879           | 0.879            | 0.908           | 0.908            |
| Aircraft (OOD) | LP@1   | <b>0.611</b> | 0.5417          | 0.5298           | 0.5833          | 0.5417           |
|                | AR@1   | <b>0.905</b> | 0.6190          | 0.6131           | 0.6667          | 0.6726           |
| Food-101 (OOD) | LP@1   | 0.791        | <b>0.8667</b>   | <b>0.8599</b>    | <b>0.8599</b>   | 0.8568           |
|                | AR@1   | <b>0.879</b> | 0.6504          | 0.6702           | 0.6756          | 0.6740           |
| CIFAR-10 (OOD) | LP@1   | 0.810        | <b>0.8788</b>   | <b>0.8728</b>    | <b>0.8768</b>   | 0.8756           |
|                | AR@1   | <b>0.833</b> | 0.6008          | 0.5940           | 0.6216          | 0.6096           |

reaching  $-32$  to  $-33$  percentage points on Food-101 and  $-32$  to  $-34$  on CIFAR-10: the retrieved neighbors are geometrically closer in the IVF index but belong to the wrong semantic class. EGA improves Label Precision over the frozen baseline on Aircraft (+9.9 pp), with bounded degradation on CIFAR-10 ( $-7.0$  pp) and Food-101 ( $-9.0$  pp). EGA’s advantage over global methods ranges from 14 pp on Aircraft to 23 pp on Food-101 and 25 pp on CIFAR-10, and tracks the difficulty of the OOD shift: when seen and unseen classes are visually similar, as in Food-101’s fine-grained dishes, the dense gradient updates of Icon and SRL pull unseen samples into the wrong seen-class clusters.

Cross-referencing Table 1 and Table 2 reveals the structural tradeoff. Icon and SRL achieve near-perfect Label Precision on CIFAR-100 (0.999 and 0.992) but collapse to 0.562 and 0.552 on Food-101. The frozen baseline and LoRA variants sit at the opposite extreme, with modest in-distribution precision (0.55 to 0.67) but stable performance under distribution shift. EGA sits between these extremes: 0.705 in-distribution and 0.791 on Food-101, 0.611 on Aircraft, and 0.810 on CIFAR-10. The six methods occupy distinct operating points in the (ID, OOD) Label Precision plane, shown in Figure 4; the dashed line connects the points that are not strictly dominated. Icon and SRL achieve the highest in-distribution precision but collapse on OOD, while LoRA variants achieve competitive OOD precision on Food-101 and CIFAR-10 but trail EGA in-distribution by 8.7 percentage points.

Table 3 compares EGA against capacity-matched LoRA ( $r = 64$  and  $r = 128$ ). On CIFAR-100, EGA achieves 0.705 LP@1, exceeding the best capacity-matched LoRA by 7.9 percentage points. Across OOD benchmarks, EGA matches or exceeds LoRA on Aircraft and remains within 7–9 percentage points on Food-101 and CIFAR-10. LoRA’s in-distribution performance slightly decreases with increasing rank, which we attribute to the low-rank constraint acting as an implicit regularizer.




Figure 4: ID/OOD operating points of adapter training paradigms in the (ID, OOD) Label Precision plane. The dashed line connects the non-dominated points.

### 4.3 Two Routes to OOD Stability: Gradient Sparsity and Capacity Restriction

The introduction identified two factors through which adapter training can distort unseen-class geometry: a loss that exerts pressure across the full embedding distribution, and an adapter capacity expressive enough to comply with that pressure. Adapters that avoid OOD collapse must neutralize one of these factors, and the route determines its position in the Label Precision plane.

#### Route 1: Gradient sparsity (EGA).

A triplet loss with margin  $m$  produces a non-zero gradient for a sample triplet  $(a, p, n)$  only when the margin constraint is violated:

$$d(a, p) - d(a, n) + m > 0. \quad (11)$$

As the adapter converges on seen classes, the fraction of *active* triplets decays. Figure 5 (left) tracks this ratio from 17.4% at epoch 10 to 3.5% at convergence. At steady state, 96.5% of triplets contribute zero gradient: the loss stops applying pressure once seen-class neighborhoods are locally separated, even though the adapter retains the capacity to refine them further. Sparsity neutralizes the loss-side factor while leaving capacity intact, which permits EGA’s high in-distribution Label Precision in Table 1.




Figure 5: Gradient sparsity as OOD protection.

**Route 2: Capacity restriction (LoRA).** The LoRA variants take the other route. They retain a global contrastive loss (InfoNCE) or a triplet loss, but constrain the adapter to a low-rank residual update. Even at  $r = 128$ , LoRA’s expressive power is limited to rank-128 perturbations of the identity, so the dense gradient field has substantially less room to reshape unseen-class geometry. This explains why LoRA+InfoNCE, despite using the same loss family as ICon, avoids the same OOD collapse: the loss applies pressure everywhere, but the adapter cannot fully comply. The price is paid in-distribution. At capacity-matched scale ( $r = 128$ ,  $\sim 4.2$ M parameters), LoRA+InfoNCE reaches 0.6176 Label Precision on CIFAR-100, 8.7 percentage points below EGA, because the same low-rank constraint that bounds OOD damage also bounds in-distribution refinement.

The two routes occupy different points in Figure 4. EGA, which neutralizes the loss-side factor while keeping capacity, achieves the highest in-distribution Label Precision among adapters that avoid OOD collapse. Capacity-matched LoRA, which neutralizes the capacity-side factor while keeping a global loss, achieves higher OOD Label Precision on CIFAR-10 and Food-101 at the cost of in-distribution refinement. Whether either route subsumes the other under stronger conditions, for example by combining gradient sparsity with low-rank constraints, is a question we leave open.

### 4.4 Ablation and Sensitivity Analysis

Table 4 isolates the contribution of each architectural component. The residual connection is the core element: removing it collapses accuracy from 0.7015 to 0.0065, indicating that the adapter cannot



Table 4: Ablation study of EGA

| Variant                 | Accuracy |
|-------------------------|----------|
| Full EGA (Ours)         | 0.705    |
| w/o Residual connection | 0.0065   |
| w/o Zero-initialization | 0.6840   |
| w/o L2-normalization    | 0.6590   |

Table 5: Sensitivity of different hyperparameters

| Triplet Margin $m$   | 0.1    | 0.2    | 0.3    | 0.5    |
|----------------------|--------|--------|--------|--------|
| Recall@1             | 0.8420 | 0.8180 | 0.8160 | 0.7400 |
| Recall@5             | 0.9782 | 0.9654 | 0.9610 | 0.9125 |
| Hidden Dimension $d$ | 512    | 1024   | 2048   | 4096   |
| Recall@1             | 0.8300 | 0.8340 | 0.8200 | 0.8240 |
| Recall@5             | 0.9695 | 0.9712 | 0.9680 | 0.9688 |

reconstruct semantic relationships from target data alone and must operate as a refinement over the pre-trained features.  $\ell_2$ -normalization keeps the transformed embeddings on the unit hypersphere, and removing it lowers accuracy to 0.6590. Zero-initialization lets the adapter begin as an identity mapping, and removing it lowers accuracy to 0.6840.

Table 5 evaluates stability across hyperparameter variations. EGA maintains consistent quality for margins  $m \in \{0.1, 0.2, 0.3\}$ , but increasing  $m$  to 0.5 degrades Recall@1 to 0.7400: a larger margin forces the adapter to update embeddings that already satisfy local neighborhood constraints, reducing gradient sparsity. Performance is also stable across hidden dimensions from 512 to 4096, indicating that EGA’s behavior is determined by its structural design rather than by specific parameter choices.

#### 4.5 Generalization to More Backbones

To verify that EGA’s benefits are not specific to CLIP ViT-B/32, we further evaluate it on two stronger vision encoders: DINOv2-large [23] and SigLIP [39].

Table 6 shows consistent improvements across all three backbones. The largest gains appear on the relatively weaker CLIP ViT-B/32 (+13.95 pp in Label Precision@1 and +13.20 pp in Recall@1), but the stronger DINOv2-large and SigLIP backbones also see stable improvements (+4.75 pp and +8.55 pp in Label Precision@1 respectively). EGA’s design therefore transfers to modern high-capacity vision encoders rather than relying on idiosyncrasies of the CLIP representation.

Table 6: Retrieval performance on more backbones.

| Backbone      | Label Precision@1 |               | ANNS Recall@1 |               |
|---------------|-------------------|---------------|---------------|---------------|
|               | Raw               | +EGA          | Raw           | +EGA          |
| CLIP ViT-B/32 | 0.5560            | <b>0.6955</b> | 0.6665        | <b>0.7985</b> |
| DINOv2-large  | 0.8530            | <b>0.9005</b> | 0.8785        | <b>0.9310</b> |
| SigLIP SO400M | 0.7740            | <b>0.8595</b> | 0.8015        | <b>0.9155</b> |

## 5 Conclusion

We approached adapter training from a vector search deployment perspective, where the index and query stream together cover a label space larger than the labeled subset available at adapter training time. Under this lens, different paradigms occupy distinct operating points in the (ID, OOD) Label Precision plane, and the choice among them reflects which slice of the deployed query distribution one is willing to serve poorly. EGA reaches 0.705 in-distribution Label Precision on CIFAR-100, with OOD degradation bounded at  $-7$  to  $-9$  pp on CIFAR-10 and Food-101 and a  $+9.9$  pp improvement on FGVC-Aircraft. ICon and SRL (0.999, 0.992 in-distribution) fall 32 to 33 pp below the frozen baseline on Food-101; capacity-matched LoRA ( $r = 128$ ) preserves OOD precision better on two of three benchmarks but trails EGA by 8.7 pp in-distribution. The benefits transfer to DINOv2-large and SigLIP (Label Precision@1 gains of 4.75 and 8.55 pp). We associate this behavior with gradient sparsity at convergence (96.5% zero-gradient triplets) and suggest it as a diagnostic for adapter designs intended to operate under unseen-class queries.

**Limitations.** First, our gradient-sparsity analysis (Section 3.3) is empirical rather than formal, and we do not provide a Lipschitz bound on the unseen-class perturbation. Second, EGA does not dominate every competitor on every axis: on CIFAR-10 and Food-101, capacity-matched LoRA achieves higher OOD Label Precision, and deployments dominated by OOD queries on these distributions may prefer LoRA. Third, our comparisons vary the loss family and architecture together in some cases, which prevents a clean disentanglement of the two factors. Controlled sweeps under matched conditions are the natural next step.

## References

- [1] Shaden Naif Alshammari, John R Hershey, Axel Feldmann, William T Freeman, and Mark Hamilton. I-con: A unifying framework for representation learning. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025.
- [2] Qi Chen, Bing Zhao, Haidong Wang, Mingqin Li, Chuanjie Liu, Zengzhong Li, Mao Yang, and Jingdong Wang. Spann: highly-efficient billion-scale approximate nearest neighbor search. In *Proceedings of the 35th International Conference on Neural Information Processing Systems, NIPS '21*, Red Hook, NY, USA, 2021. Curran Associates Inc.
- [3] Haya Diwan, Jinrui Gou, Cameron Musco, Christopher Musco, and Torsten Suel. Navigable graphs for high-dimensional nearest neighbor search: Constructions and limits. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 59513–59531. Curran Associates, Inc., 2024.
- [4] Zijian Dong, Yilei Wu, Chongyao Chen, Yingtian Zou, Yichi Zhang, and Juan Helen Zhou. Improve representation for imbalanced regression through geometric constraints. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- [5] Sedigheh Eslami and Gerard de Melo. Mitigate the gap: Improving cross-modal alignment in CLIP. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [6] Tianyu Gao, Xingcheng Yao, and Danqi Chen. SimCSE: Simple contrastive learning of sentence embeddings. In Marie-Francine Moens, Xuanjing Huang, Lucia Specia, and Scott Wen-tau Yih, editors, *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 6894–6910, Online and Punta Cana, Dominican Republic, November 2021. Association for Computational Linguistics.
- [7] Hao Guo and Youyou Lu. Achieving low-latency graph-based vector search via aligning best-first search algorithm with ssd. In *19th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, pages 171–186, Boston, MA, 2025. USENIX Association.
- [8] Hao Guo and Youyou Lu. Odinann: Direct insert for consistently stable performance in billion-scale graph-based vector search. In *24th USENIX Conference on File and Storage Technologies (FAST)*, Santa Clara, CA, 2026. USENIX Association.
- [9] Neil Houlsby, Andrei Giurgiu, Stanislaw Jastrzebski, Bruna Morrone, Quentin De Laroussilhe, Andrea Gesmundo, Mona Attariyan, and Sylvain Gelly. Parameter-efficient transfer learning for NLP. In Kamalika Chaudhuri and Ruslan Salakhutdinov, editors, *Proceedings of the 36th International Conference on Machine Learning*, volume 97 of *Proceedings of Machine Learning Research*, pages 2790–2799. PMLR, 09–15 Jun 2019.
- [10] Edward J Hu, yelong shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In *International Conference on Learning Representations*, 2022.
- [11] Jie Hu, Li Shen, Samuel Albanie, Gang Sun, and Enhua Wu. Squeeze-and-excitation networks. *IEEE Trans. Pattern Anal. Mach. Intell.*, 42(8):2011–2023, August 2020.
- [12] Piotr Indyk and Haike Xu. Worst-case performance of popular approximate nearest neighbor search implementations: Guarantees and limitations. In A. Oh, T. Naumann, A. Globerson, K. Saenko, M. Hardt, and S. Levine, editors, *Advances in Neural Information Processing Systems*, volume 36, pages 66239–66256. Curran Associates, Inc., 2023.
- [13] Elias Jääsaari, Ville Hyvönen, and Teemu Roos. LoRANN: Low-rank matrix factorization for approximate nearest neighbor search. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [14] Menglin Jia, Luming Tang, Bor-Chun Chen, Claire Cardie, Serge Belongie, Bharath Hariharan, and Ser-Nam Lim. Visual prompt tuning. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXIII*, page 709–727, Berlin, Heidelberg, 2022. Springer-Verlag.

- [15] Dong Jing, Xiaolong He, Yutian Luo, Nanyi Fei, Guoxing Yang, Wei Wei, Huiwen Zhao, and Zhiwu Lu. FineCLIP: Self-distilled region-based CLIP for better fine-grained understanding. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.
- [16] Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with gpus. *IEEE Transactions on Big Data*, 7(3):535–547, 2021.
- [17] Kate Keahey, Jason Anderson, Zhuo Zhen, Pierre Riteau, Paul Ruth, Dan Stanzione, Mert Cevik, Jacob Colleran, Haryadi S. Gunawi, Cody Hammock, Joe Mambretti, Alexander Barnes, François Halbach, Alex Rocha, and Joe Stubbs. Lessons learned from the chameleon testbed. In *Proceedings of the 2020 USENIX Annual Technical Conference (USENIX ATC '20)*. USENIX Association, July 2020.
- [18] Hyungbin Kim, Incheol Baek, and Yon Dohn Chung. Decoupled contrastive learning for federated learning, 2025.
- [19] Panagiotis Koromilas, Efthymios Georgiou, Giorgos Bouritsas, Theodoros Giannakopoulos, Mihalis A. Nicolaou, and Yannis Panagakis. A principled framework for multi-view contrastive learning, 2025.
- [20] Chiara Lanza, Roberto Pereira, Marco Miozzo, Eduard Angelats, and Paolo Dini. Degradation of feature space in continual learning, 2026.
- [21] Fengming Lin. Vision language models: A survey of 26k papers, 2025.
- [22] Yu A. Malkov and D. A. Yashunin. Efficient and robust approximate nearest neighbor search using hierarchical navigable small world graphs. *IEEE Trans. Pattern Anal. Mach. Intell.*, 42(4):824–836, April 2020.
- [23] Maxime Oquab, Timothée Darcet, Théo Moutakanni, Huy V. Vo, Marc Szafraniec, Vasil Khalidov, Pierre Fernandez, Daniel HAZIZA, Francisco Massa, Alaaeldin El-Nouby, Mido Assran, Nicolas Ballas, Wojciech Galuba, Russell Howes, Po-Yao Huang, Shang-Wen Li, Ishan Misra, Michael Rabbat, Vasu Sharma, Gabriel Synnaeve, Hu Xu, Herve Jegou, Julien Mairal, Patrick Labatut, Armand Joulin, and Piotr Bojanowski. DINOv2: Learning robust visual features without supervision. *Transactions on Machine Learning Research*, 2024. Featured Certification.
- [24] Avik Pal, Max van Spengler, Guido Maria D’Amely di Melendugno, Alessandro Flaborea, Fabio Galasso, and Pascal Mettes. Compositional entailment learning for hyperbolic vision-language models. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [25] Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal, Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, Gretchen Krueger, and Ilya Sutskever. Learning transferable visual models from natural language supervision. In Marina Meila and Tong Zhang, editors, *Proceedings of the 38th International Conference on Machine Learning, ICML 2021, 18-24 July 2021, Virtual Event*, Proceedings of Machine Learning Research, pages 8748–8763. PMLR, 2021.
- [26] Sameera Ramasinghe, Violetta Shevchenko, Gil Avraham, and Ajanthan Thalaiyasingam. Accept the modality gap: An exploration in the hyperbolic space. In *2024 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 27253–27262, 2024.
- [27] Simone Ricci, Niccolò Biondi, Federico Pernici, Ioannis Patras, and Alberto Del Bimbo.  $\lambda$ -orthogonality regularization for compatible representation learning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2026.
- [28] Daniele Savietto, Declan Campbell, André Panisson, Marco Nurisso, Giovanni Petri, Jonathan D. Cohen, and Alan Perotti. The geometry of representational failures in vision language models, 2026.
- [29] Shivam Rajendra Rai Sharma, Xiaoguang Zhu, Luca Cerny Oliveira, Kartik Patwari, La Rissa Vasquez, David Garcia, Louise Nicole C. Sevilla, Brittany N Dugger, and Chen-Nee Chuah. Benchmarking parameter efficient adaptation of vision language models on pathology. In *NeurIPS 2025 Workshop for Imageomics: Discovering Biological Knowledge from Images Using AI*, 2025.

- [30] Suhas Jayaram Subramanya, Devvrit, Rohan Kadekodi, Ravishankar Krishaswamy, and Harsha Vardhan Simhadri. Diskann: fast accurate billion-point nearest neighbor search on a single node. In *Proceedings of the 33rd International Conference on Neural Information Processing Systems*, 2019.
- [31] Hangke Sui, Yuqing Wang, and Minh N. Do. Unicon: Unified framework for efficient contrastive alignment via kernels. In *The Fourteenth International Conference on Learning Representations*, 2026.
- [32] Aaron van den Oord, Yazhe Li, and Oriol Vinyals. Representation learning with contrastive predictive coding. *arXiv preprint arXiv:1807.03748*, 2018.
- [33] Liang Wang, Xiang Tao, Qiang Liu, Shu Wu, and Liang Wang. Rethinking graph masked autoencoders through alignment and uniformity. In *Proceedings of the Thirty-Eighth AAAI Conference on Artificial Intelligence and Thirty-Sixth Conference on Innovative Applications of Artificial Intelligence and Fourteenth Symposium on Educational Advances in Artificial Intelligence, AAAI’24/IAAI’24/EAAI’24*. AAAI Press, 2024.
- [34] Tongzhou Wang and Phillip Isola. Understanding contrastive representation learning through alignment and uniformity on the hypersphere. In *Proceedings of the 37th International Conference on Machine Learning, ICML’20*. JMLR.org, 2020.
- [35] Mingqi Wu, Qiang Sun, and Archer Y. Yang. PCA++: How uniformity induces robustness to background noise in contrastive learning. In *The Thirty-ninth Annual Conference on Neural Information Processing Systems*, 2026.
- [36] Baili Xiao, Zhibin Dong, Ke Liang, Suyuan Liu, Siwei Wang, Tianrui Liu, Xingchen Hu, En Zhu, and Xinwang Liu. Easemvc: Efficient dual selection mechanism for deep multi-view clustering. In *2025 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 20716–20726, 2025.
- [37] Xuqian Xue, Yiming Lei, Qi Cai, Hongming Shan, and Junping Zhang. PROTOCOL: Partial optimal transport-enhanced contrastive learning for imbalanced multi-view clustering. In *Forty-second International Conference on Machine Learning*, 2025.
- [38] Ming Yang, Yuzheng Cai, and Weiguo Zheng. Cspg: Crossing sparse proximity graphs for approximate nearest neighbor search. In A. Globerson, L. Mackey, D. Belgrave, A. Fan, U. Paquet, J. Tomczak, and C. Zhang, editors, *Advances in Neural Information Processing Systems*, volume 37, pages 103076–103100. Curran Associates, Inc., 2024.
- [39] Xiaohua Zhai, Basil Mustafa, Alexander Kolesnikov, and Lucas Beyer. Sigmoid loss for language image pre-training. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 11941–11952, 2023.
- [40] Lvmin Zhang, Anyi Rao, and Maneesh Agrawala. Adding conditional control to text-to-image diffusion models. In *2023 IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 3813–3824, 2023.
- [41] Renrui Zhang, Wei Zhang, Rongyao Fang, Peng Gao, Kunchang Li, Jifeng Dai, Yu Qiao, and Hongsheng Li. Tip-adapter: Training-free adaption of clip for few-shot classification. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXV*, page 493–510, Berlin, Heidelberg, 2022. Springer-Verlag.
- [42] Renrui Zhang, Wei Zhang, Rongyao Fang, Peng Gao, Kunchang Li, Jifeng Dai, Yu Qiao, and Hongsheng Li. Tip-adapter: Training-free adaption of clip for few-shot classification. In *Computer Vision – ECCV 2022: 17th European Conference, Tel Aviv, Israel, October 23–27, 2022, Proceedings, Part XXXV*, page 493–510, Berlin, Heidelberg, 2022. Springer-Verlag.
- [43] Yingwen Zhang, Meng Wang, Xihua Sheng, Peilin Chen, Junru Li, Li Zhang, and Shiqi Wang. An information-theoretic regularizer for lossy neural image compression. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pages 15573–15582, October 2025.

- 487 [44] Qingyan Zhao, Ruifang He, Jinpeng Zhang, Chang Liu, and Bo Wang. Representation de-  
488 generation problem in prompt-based models for natural language understanding. In Nicoletta  
489 Calzolari, Min-Yen Kan, Veronique Hoste, Alessandro Lenci, Sakriani Sakti, and Nianwen Xue,  
490 editors, *Proceedings of the 2024 Joint International Conference on Computational Linguistics,  
491 Language Resources and Evaluation (LREC-COLING 2024)*, pages 13946–13957, Torino, Italia,  
492 May 2024. ELRA and ICCL.
- 493 [45] Kaiyang Zhou, Jingkan Yang, Chen Change Loy, and Ziwei Liu. Conditional prompt learning  
494 for vision-language models. In *Proceedings of the IEEE/CVF Conference on Computer Vision  
495 and Pattern Recognition*, pages 16816–16825, 2022.
- 496 [46] Kaiyang Zhou, Jingkan Yang, Chen Change Loy, and Ziwei Liu. Learning to prompt for  
497 vision-language models. *Int. J. Comput. Vision*, 130(9):2337–2348, September 2022.
- 498 [47] Xin Zhou, Yongjie Wang, and Zhiqi Shen. Cm<sup>3</sup>: Calibrating multimodal recommendation,  
499 2025.

## NeurIPS Paper Checklist

### 1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The paper presents both theoretical and experimental evidence to support its claims.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

### 2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: We discuss the limitations of our approach in the conclusion.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Full proofs are provided in the appendix, and all assumptions are clearly stated in the main paper and/or the appendix.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

#### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: All code and data needed to reproduce the main experimental results are available on the anonymous Github repository.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

#### 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: All the details to reproduce the main experimental results are provided in the anonymous Github repository, including instructions for data access and preparation, and scripts to reproduce all experimental results for the new proposed method and baselines.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: All such details are provided in the anonymous Github repository, and the main paper includes a summary of the experimental setting.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The results are pretty stable so we did not report error bars; we explain this in the main paper.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.



- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: These details are provided in the anonymous Github repository, and the main paper includes a summary of the compute resources used for the experiments.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: We believe we have followed the NeurIPS Code of Ethics in all aspects of our research.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.

- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

### 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

### 14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: [N/A]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.

- 810           • We recognize that the procedures for this may vary significantly between institutions  
811           and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the  
812           guidelines for their institution.  
813           • For initial submissions, do not include any information that would break anonymity (if  
814           applicable), such as the institution conducting the review.

815   **16. Declaration of LLM usage**

816   Question: Does the paper describe the usage of LLMs if it is an important, original, or  
817   non-standard component of the core methods in this research? Note that if the LLM is used  
818   only for writing, editing, or formatting purposes and does *not* impact the core methodology,  
819   scientific rigor, or originality of the research, declaration is not required.

820   Answer: [N/A]

821   Justification: [N/A]

822   Guidelines:

- 823           • The answer [N/A] means that the core method development in this research does not  
824           involve LLMs as any important, original, or non-standard components.  
825           • Please refer to our LLM policy in the NeurIPS handbook for what should or should not  
826           be described.

## A Technical Conditions for the Sparsity-Locality Claim

This appendix elaborates on the technical conditions underlying the empirical claim made in Section 3.3.

**(1) Gradient support.** The triplet loss in Eq. (7) can be written as  $\mathcal{L} = \frac{1}{|\mathcal{B}|} \sum_{(a,p,n) \in \mathcal{B}} \ell_{apn}$ , where  $\ell_{apn} = \max(0, \delta_{apn} + m)$  and  $\delta_{apn} = d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_p) - d(\hat{\mathbf{z}}_a, \hat{\mathbf{z}}_n)$ . The subgradient with respect to  $\theta$  is

$$\nabla_{\theta} \ell_{apn} = \mathbf{1}[\delta_{apn} + m > 0] \cdot \nabla_{\theta} \delta_{apn}, \quad (12)$$

which is exactly zero whenever  $\delta_{apn} + m \leq 0$ . Therefore,  $\nabla_{\theta} \mathcal{L}$  receives contributions only from active triplets satisfying Eq. (8).

**(2) Convergence of  $\rho(t)$ .** Under mild regularity conditions (Lipschitz continuity of  $f_{\theta}$  and bounded embedding norms), the sequence  $\{\hat{\mathbf{z}}^{(t)}\}$  converges to a stationary point  $\hat{\mathbf{z}}^*$  of  $\mathcal{L}$ . At stationarity,  $\rho(t)$  converges to a steady-state value  $\rho^*$  determined by the within-class distance distribution of the converged embeddings, since  $\rho(t)$  is a continuous functional of the embedding distribution and the embeddings converge.

**(3) Empirical locality.** Strictly speaking, parameter updates  $\nabla_{\theta} \delta_{apn}$  computed on seen-class triplets affect  $f_{\theta}$  globally (including its behavior on unseen-class inputs) since  $f_{\theta}$  is a parametric MLP and is not locally supported on the input manifold. We therefore do not claim formal zero-gradient over unseen regions. Instead, we observe an empirical locality property: when (i) the residual structure with zero initialization keeps the cumulative perturbation magnitude bounded by the sum of training gradients, and (ii) gradient sparsity at convergence ( $\rho^* \approx 3.5\%$  on FGVC-Aircraft) attenuates this sum, the resulting perturbation in unseen-class regions is small enough to preserve the pretrained semantic structure for retrieval purposes. This observation is empirically supported by the OOD Label Precision results in Section 4.2 and Figure 6, where EGA preserves or improves Label Precision on unseen categories while ICon and SRL, which lack the gradient sparsity property, collapse below the frozen baseline. A formal Lipschitz-based bound on this perturbation is left for future work.

**(4) Dependence on  $m$ .** The residual distance gap  $\delta_{apn}^*$  at convergence is bounded by the within-class diameter  $\Delta_c = \max_{x,x' \in c} d(\hat{\mathbf{z}}, \hat{\mathbf{z}}')$  for seen class  $c$ . For  $m \leq \min_c \Delta_c$ , the adapter can drive all within-class distances below the threshold, pushing  $\rho^*$  toward zero. For  $m > \min_c \Delta_c$ , some triplets remain active at convergence because the within-class spread exceeds the margin, resulting in a strictly positive  $\rho^*$ . This explains the empirically observed increase in  $\rho^*$  with larger  $m$  reported in Table 5.

## B Experimental Setup Details

**Backbones models.** All experiments in the main paper use frozen CLIP ViT-B/32 [25] image embeddings as the base representation. To verify that EGA’s benefits are not specific to this backbone, we further evaluate EGA on two stronger vision encoders: DINOv2-large [23] and SigLIP [39]. DINOv2-large is a self-supervised vision transformer trained on 142M images, while SigLIP is a supervised vision transformer trained on 400M image-text pairs. Both models produce higher-quality embeddings than CLIP ViT-B/32, as reflected in their stronger raw retrieval performance, and EGA continues to provide meaningful improvements on top of these stronger backbones, as shown in Table 6 in Section 4.4.

**Adapter models.** We compare EGA against the frozen CLIP baseline and four trained adapters, all sharing the same training data, evaluation protocol, and 75/25 retrieval split:

- **ICon** [1]: a global contrastive objective minimizing the KL divergence between the empirical similarity distribution and the class-membership distribution. We instantiate ICon on the same EGAMLP architecture used by our method ( $\sim 4.2\text{M}$  parameters).
- **SRL** [4]: a structural regularization objective jointly optimizing batch-wide uniformity and within-class homogeneity, also instantiated on the EGAMLP architecture.

- **LoRA + InfoNCE** [10, 32]: a capacity-matched low-rank adapter (rank  $r = 64$  and  $r = 128$ ,  $\sim 2.6\text{M}$ – $4.2\text{M}$  parameters) with zero-initialized residual structure, trained with supervised InfoNCE, which is a standard global contrastive loss.
- **LoRA + Triplet** [10]: identical capacity-matched low-rank architecture as above (rank  $r = 64$  and  $r = 128$ ), but trained with the same small-margin ( $m = 0.2$ ) triplet loss as EGA.

The two LoRA configurations isolate the effect of architectural capacity vs. loss function: comparing LoRA + InfoNCE against ICon/SRL holds the loss family roughly fixed while changing capacity, and comparing LoRA + Triplet against EGA holds the loss fixed while changing capacity.

**Datasets.** We evaluate on four datasets covering both in-distribution and OOD retrieval scenarios. **CIFAR-100** (100 classes, 60K images) serves as the in-distribution benchmark: the adapter is trained and evaluated on the same class label space, with a 75/25 split between database and query sets. **CIFAR-10** (10 classes, 60K images) provides a cross-dataset OOD setting: the adapter is trained on CIFAR-100 and evaluated on the disjoint CIFAR-10 classes. **FGVC-Aircraft** (100 fine-grained aircraft variants, 10K images) and **Food-101** (101 food categories, 25K images for the test split we use) provide unseen-class OOD settings: classes are randomly partitioned into 80% seen / 20% unseen with a fixed seed, the adapter is trained on the seen split, and retrieval is evaluated strictly on the unseen split. Image embeddings for all datasets are extracted with frozen CLIP ViT-B/32.

**Indexing and metrics.** For all retrieval evaluations we use FAISS [16] IndexIVFFlat with  $nlist=10$  centroids and report results at  $nprobe \in \{1, 5, 10\}$ . We measure both Label Precision ( $LP@K$ , Eq. 1) and ANNS Recall ( $AR@K$ , Eq. 2) at  $K \in \{1, 3, 5, 10\}$ .

**Training.** All adapters are trained for 150 epochs with AdamW (weight decay  $10^{-4}$ ), cosine-annealed learning rate, and batch size 256. EGAMLP-based methods (EGA, ICon, SRL) use learning rate  $10^{-4}$ ; LoRA variants use  $10^{-3}$  to compensate for their smaller parameter count. Triplet-based methods sample one positive and one negative per anchor uniformly within each mini-batch; global contrastive methods (ICoN, SRL, LoRA + InfoNCE) operate on the full pairwise similarity matrix. All experiments use a fixed random seed (42) for reproducibility.

**Compute platform.** All experiments were conducted on the Chameleon Cloud [17] platform using a compute node equipped with 256 AMD EPYC-7763 CPU cores, 512 GiB of RAM, and an NVIDIA A100 GPU of 80 GB HBM2e. The operating system is Ubuntu 24.04 LTS.

## C Additional Experimental Results

**Robustness across search budgets.** Figure 6 traces Label Precision@1 across three OOD benchmarks and three search budgets ( $nprobe \in \{1, 5, 10\}$ ). The collapse pattern of ICon and SRL persists across all 18 measurement points (3 datasets  $\times$  3 budgets  $\times$  2 methods), ruling out the possibility that the degradation is an artifact of any particular benchmark or indexing configuration. EGA’s relative ordering against other methods is also stable across budgets.




Figure 6: Label Precision@1 across three OOD benchmarks and three search budgets ( $nprobe \in \{1, 5, 10\}$ ).